

Stay, React, or Conform? Examining Social Influence Effects on Aesthetic Judgments

Omar Lizardo & Sara Skiles

Revised Manuscript — September 16, 2026

Abstract

In this study, we address the question, foundational to sociological studies of taste, of the extent to which others’ aesthetic judgments causally influence individual evaluations of cultural works, and how this influence depends on the social distance between self and other. While past research has documented observational associations between cultural tastes and social networks, the underlying causal mechanisms—whether individuals actively conform, react, or remain anchored in their initial judgments—have rarely been submitted to experimental test. Using a national survey experiment ($N = 2,275$), we manipulate exposure to others’ aesthetic evaluations (likes vs. dislikes) and the status characteristics of the evaluating group across repeated trials. Using a difference-in-differences framework and inverse probability weighting, we find substantial evidence of *valence asymmetry*: while negative evaluations from others broadly suppress positive exposure drift across groups, positive evaluations facilitate upward shifts, conditional on status congruence. Furthermore, analyzing discrete behavioral choices shows that evaluative stability (“staying”) is not driven by high status alone, but by *structural status consistency*: both high-consistent and low-consistent individuals display high evaluative inertia, whereas status-inconsistent individuals exhibit evaluative fluidity.

1 Introduction

Taste expressions and consumption have long been shown to be a function of individuals’ desire to differentiate themselves from others (Simmel, 1957; DiMaggio, 1987). Bourdieu (1984) suggests that taste differences evoke a visceral reaction and that individuals engage in symbolic boundary work and status distinction by rejecting others’ tastes. In this way, expressing *distastes* serves to communicate social distance from others (Bryson, 1996; Lamont and Molnár, 2002). Accordingly, in contexts where individuals discover an unexpected similarity between their taste and that of members of an outgroup, they abandon these tastes to maintain social boundaries (Berger and Heath, 2008; Simmel, 1957).

However, other people’s tastes do not have to be solely a source of distinction and separation; they can also be a basis for building solidarity and discovering shared perspectives, especially among people considered similar in some socially relevant way. In this way, expressing shared likes and dislikes can signal affiliation with others and serve as a mechanism for creating bounded groups (DiMaggio, 1987; Lizardo, 2006). Accordingly, contexts in which individuals find a *similarity* of tastes with members of their ingroup can foster community and solidarity (Wohl, 2015; Vanzella-Yang and Abrutyn, 2022).

Despite the rich sociological literature linking cultural tastes and network ties (Erickson, 1996; Lizardo, 2006; Lewis and Kaufman, 2018), this body of work has almost exclusively relied on static, cross-sectional, or observational survey data. However, it is well known that observing similar people having similar tastes and dissimilar people having dissimilar tastes in cross-sectional data inevitably conflates prior social sorting, shared environmental exposure, and active social influence (Centola, 2011; Aral et al., 2009; Shalizi and Thomas, 2011). Even when dynamic observational data are available, threats to causal inference from unobserved confounders remain. Some work has addressed these issues by moving to more elegant experimental frameworks (Salganik et al., 2006) but in these settings, the social feedback comes from a relatively anonymous mass. As such, the question of *whether learning that a given cultural object is regarded positively or negatively by ingroup or outgroup others causally affects our transient evaluation of that object* in the ways predicted by the aforementioned theories remains open.

1.1 Study Description

To answer this question, we designed an experimental study to disentangle three distinct social influence mechanisms shaping aesthetic taste for a cultural object. In our study, participants were asked to tell us how they felt about a painting—Whistler’s *Nocturne*—on a scale ranging from “dislike it very much” to “like it very much.” After we obtained this evaluation, in the experimental part of the survey, we told a subset of participants that our survey had so far shown that the group of participants that was most likely to either like (for one subset of participants) or dislike (for another subset) was either people with a PhD and working in higher education, people holding a bachelor’s degree in business and working in financial investment, or people with a high school education and working in food service. They then spent between 10 and 20 minutes answering various sociodemographic questions and completing survey modules. Participants were then shown the same painting one last time and asked to provide us with their final take on it, using the same evaluative categories

as in the beginning.

1.2 Hypotheses

1.2.1 Mere Exposure and Positive Drift

For participants in the Control Condition who are simply asked to rate the painting twice without being exposed to the opinions of others, repeated exposure to a cultural object across trials is expected to produce an upward drift in evaluation even in the absence of external social information, a phenomenon in psychology known as the *mere exposure effect* (Zajonc, 1968; Reber et al., 1998).¹ So our baseline hypothesis is simply that we should observe such positive drift, with participants who are asked to judge the same cultural object twice being more likely to produce a more positive judgment (on average) the second time around.

1.2.2 Social Influence and Generalized Alignment

External social cues then either facilitate or suppress this exposure drift. In the case of generalized influence without status cues, exposure to positive evaluations from others is expected to enhance this positive exposure drift, whereas exposure to negative evaluations from others is expected to suppress it (Salganik et al., 2006). Accordingly, we should expect (*Hypothesis 1*) that individuals will *adjust* their judgments to align with the aggregate evaluations of others, shifting positively when provided information that others like the object and negatively when provided information that others dislike it.

1.2.3 Cultural Goodwill and Asymmetric Alignment

Following Bourdieu’s (1984) idea of “cultural goodwill,” we should expect (*Hypothesis 2*) lower-status respondents to be more likely to *align* their judgments with those of high-status others, such that they will tend to move towards more positive evaluations when provided information that high-status others liked the cultural object, and towards negative evaluations when provided information that high-status others disliked the painting.

1.2.4 Status-Based Distancing and Asymmetric Disalignment

An alternative hypothesis, based on the idea of *reactance* and distancing from the judgments of others of different status, particularly when it comes to high-status individuals distancing themselves from the judgments of low-status others (Bourdieu, 1984). Accordingly, we should expect (*hypothesis 3*) that high-status individuals will shift away from the judgments of low-status individuals, moving towards dislike when provided with information that low-status others liked the cultural object, and towards liking when provided with information that low-status others disliked the painting.

¹For an application of the mere exposure effect to historical dynamics in the aesthetic evaluation of art see Cutting (2003).

1.2.5 Status Inconsistency and Aesthetic Malleability

Finally, with regard to status crystallization and status inconsistency (Hope, 1975), we should expect (*Hypothesis 4*) individuals with consistent status profiles (e.g., Middle Class with College or Working Class without College) to exhibit a higher propensity to maintain their initial evaluation (“staying”) than status-inconsistent individuals, whose judgments should be more malleable given information on the aesthetic opinions of others.

2 Data and Methods

2.1 Sample Recruitment and Survey Administration

Data for this study were collected between February 28 and March 6, 2012, via an online experimental survey administered by Survey Sampling International (SSI), a leading research firm specializing in online panel recruitment and survey sampling. A total of 3,782 panel members accessed the survey instrument. Of these, 97 were excluded due to lack of informed consent ($N = 85$) or age under 18 ($N = 12$), 1,394 were screened out because demographic quota targets had already been met, and 16 dropped out during early screens, resulting in an analytic sample of $N = 2,275$ completed responses (with $N = 2,253$ providing complete initial and post-treatment evaluations). Participants received monetary compensation (\$3.92) via SSI, funded by an NSF Doctoral Dissertation Improvement Grant (Award SES-1203426) and the University of Notre Dame Institute for Scholarship in Liberal Arts.

The sampling frame employed dynamic quota matching calibrated to U.S. Census benchmarks for gender, age, race/ethnicity, and geographic region to achieve a diverse and quasi-representative national sample of U.S. adults. To ensure adequate statistical power for examining status-differentiated cultural dynamics—particularly within high-cultural-capital occupational cells—the sample incorporated a deliberate oversample of college graduates (40.5% holding a bachelor’s degree or higher in the sample versus approximately 28% nationally in 2012). Table 1 presents the full demographic breakdown of the sample, compared with U.S. population parameters, separated by experimental treatment condition.²

The randomization mechanism built into the *Qualtrics* survey engine ensures that internal validity is maximized: respondents were randomly assigned to treatment and control conditions, creating balanced experimental groups across all observed baseline characteristics (Table 1). At the same time, the demographic diversity across gender, age, race, and geographic region provides substantial ecological validity and heterogeneity compared to conventional undergraduate convenience samples.³

²Because the survey used a quota-matched, opt-in online consumer panel rather than a strict probability-based sample with known inclusion probabilities (such as the GSS or ANES), the reported standard errors, p -values, and confidence intervals should be interpreted as reflecting *model-based and experimental design uncertainty within the sample*, rather than formal design-based population parameter estimates.

³Because cell sample sizes vary across experimental conditions and observational status consistency subgroups, we also report a comprehensive Minimum Detectable Effect (MDE) power analysis in Appendix A (Table A.1 and Figure A.1), explicitly delineating the statistical precision and empirical detection boundaries of the experimental design.

Table 1: Demographic Characteristics of Survey Sample Compared to U.S. Population Benchmarks

Demographic Variable	Sample N	Sample %	U.S. Pop %	Treatment	No Alter	Taste Only
Gender						
Women	1,184	52.1%	51.5%	51.7%	52.2%	53.9%
Men	1,091	47.9%	48.5%	48.3%	47.8%	46.1%
Race / Ethnicity						
White	1,463	64.3%	65.1%	64.7%	63.2%	63.1%
Hispanic / Latino	365	16.1%	15.8%	16.2%	14.8%	15.9%
Black / African American	279	12.3%	12.3%	12.1%	13.7%	12.1%
Asian / Pacific Islander	103	4.5%	4.5%	4.4%	5.5%	4.6%
Other / Multiracial	65	2.8%	2.3%	2.5%	2.7%	4.3%
Age Group						
18–24	272	12.0%	13.1%	11.9%	11.0%	12.7%
25–34	410	18.0%	17.5%	18.2%	16.5%	17.9%
35–44	402	17.7%	17.5%	17.9%	13.7%	18.7%
45–54	441	19.4%	19.2%	19.4%	20.9%	18.4%
55–64	372	16.3%	15.6%	16.5%	15.4%	16.1%
65+	377	16.6%	17.2%	16.1%	22.5%	16.2%
Geographic Region						
Northeast	417	18.3%	17.9%	18.8%	16.5%	17.0%
South	816	35.9%	37.2%	36.6%	39.6%	30.5%
Midwest	505	22.2%	21.6%	21.2%	26.4%	24.8%
West	536	23.6%	23.3%	23.4%	17.6%	27.7%
Education (Oversampled)						
Less than High School	89	3.9%	13.3%	3.8%	3.3%	4.6%
High School Diploma	474	20.8%	30.4%	20.3%	23.6%	21.9%
Some College / AA	791	34.8%	28.5%	36.0%	31.3%	30.6%
Bachelor’s Degree	625	27.5%	18.1%	27.3%	27.5%	28.5%
Graduate / Professional	296	13.0%	9.6%	12.6%	14.3%	14.5%
Total N	2,275	100.0%		1,746	182	347

Note: U.S. population benchmarks are drawn from the 2010/2012 U.S. Census Bureau Current Population Survey (CPS). The analytic sample oversamples college-educated respondents to populate higher-status occupational cells.

2.2 Experimental Design and Stimulus

All respondents viewed an image of the same artwork: *Nocturne: Blue and Gold – Old Battersea Bridge* by James McNeill Whistler (1872), displayed in Figure 1. The painting was selected based on an online pilot study ($N = 100$) testing eight artworks; the Whistler painting was chosen because it generated the highest initial pre-treatment variance in aesthetic evaluations across all demographic subgroups.

The experiment used an incomplete factorial design with within-subject pre- and post-trials. In Trial 1 (pre-treatment baseline), respondents viewed the artwork accompanied by



Figure 1: Experimental Stimulus: *Nocturne: Blue and Gold – Old Battersea Bridge*, by James McNeill Whistler (1872)

the verbatim prompt:

“Please look at this painting for a few seconds.”

“Give us your first impression of this painting – how much do you like it?”

Responses were recorded using a 7-point Likert scale (ranging from 1 = “I like it very much” to 7 = “I dislike it very much”, along with an “I can’t see the image” screening option; reverse-coded so that 7 = “like very much” and 1 = “dislike very much”). Next, during an intervening survey block, respondents completed demographic questions and a musical taste battery before treatment respondents read a brief vignette describing how a group of other people evaluated the artwork. Finally, in Trial 2 (post-treatment final impression), respondents were shown the painting again and asked:

“Please look at this painting for a few seconds again.”

“Give us your final impression of this painting – how much do you like it?”

on the identical rating scale. (The complete verbatim wording, layout, and sequencing across all survey screens are transcribed in Appendix B.)

To prevent terminological confusion, we distinguish between the *pre-treatment baseline measurement* (Trial 1, which captures initial, uninfluenced aesthetic preference) and the *Control Condition* (the experimental group that receives no information about others between trials, allowing us to observe pure temporal re-evaluation in the absence of treatment). We refer to changes over repeated trials in this unexposed group as *exposure drift*.

2.3 Mitigation of Experimenter Demand Characteristics and Respondent Reactivity

A common methodological concern in within-subject experimental designs is whether respondents become aware that an evaluative shift is being measured, inducing experimenter demand compliance or artificial consistency. We implemented four procedural and design safeguards to mitigate demand awareness and assess respondent reactivity.

First, Trial 1 and Trial 2 were separated by an extensive cognitive and temporal buffer comprising 15 to 20 minutes of substantive survey batteries. Between aesthetic evaluations, respondents completed detailed modules on educational history, parental degree attainment, employment status, 22-category Standard Occupational Classification (SOC) coding, annual household income, and a comprehensive 20-genre musical taste battery evaluating preference ratings, listening frequency, and listener demographic stereotypes. This substantial intervening distance prevented respondents from mechanically recalling or feeling compelled to repeat their Trial 1 ratings.

Second, the experimental feedback screens were framed naturalistically as observational descriptive summaries from prior survey waves rather than prescriptive social norms: respondents were told, “*Our survey so far indicates that the group of people most likely to say that they [LIKE / DISLIKE] the painting you saw have completed...*” Immediately following this prompt, respondents completed attribution and subjective similarity probes concerning the evaluating group, maintaining the perception that the survey was investigating social perception rather than testing their personal susceptibility to persuasion. When respondents subsequently reached the end of the survey and were re-shown the painting, they were simply asked for their “final impression” using the identical rating scale, without being reminded of their initial Trial 1 score and without any instructional prompt linking their second rating to the opinions of others or suggesting that conformity was expected.

Third, in the following analyses, all treatment shifts are compared against the unexposed Control Condition. Because control respondents completed the identical intervening survey battery without receiving any cue from others, comparing treatment conditions against this baseline cleanly separates genuine informational influence from any residual re-test reactivity or temporal fatigue.

Finally, the empirical distribution of discrete behavioral choices provides powerful evidence against demand compliance. Over 74% of respondents across conditions exhibited complete evaluative stability (“staying”), and a substantial proportion displayed active oppositional reactance by shifting in the direction opposite to the cue. Had respondents been compliance-driven, conformity rates would have dominated across conditions, rather than emerging selectively in status-inconsistent configurations.

2.4 Variables and Measurement

Aesthetic evaluation was measured on a reversed 1–7 scale. In continuous models, the outcome is the repeated evaluation across trials. In behavioral models, change between trials is categorized into three mutually exclusive choices: *Stay* ($\Delta = 0$), *Conform* (moving in the direction of others’ evaluations), and *React* (moving in the opposite direction).

The treatment conditions crossed others’ evaluation (Like vs. Dislike) with others’ social status. High-status others were differentiated into Cultural Specialists (possessing multiple advanced degrees including a PhD and working in higher education) and Economic Specialists (holding a bachelor’s degree in business and working in financial investment). Low-status others were described as possessing a high school education and working in food service. In addition, two Taste-Only conditions provided other like-or-dislike evaluations without any status indicators, and a pure Control Condition provided no information about others between trials ($N = 182$).

We collapse the Economic Specialist and Cultural Specialist conditions into a single High-Status condition (*+Status*) for the primary analysis. This consolidation is justified on both statistical and substantive grounds. First, crossing respondents’ own status consistency quadrants by all nine experimental conditions would expand the design to 36 cells, reducing cell sizes in smaller subgroups to $N \approx 20\text{--}35$ (and $N = 23$ in control), severely inflating the Minimum Detectable Effect threshold ($d > 0.45$). Second, formal linear hypothesis tests confirm that responses to Economic Specialists and Cultural Specialists are statistically indistinguishable for both positive cues ($F = 0.29, p = 0.591$ in the pooled sample; $F = 0.92, p = 0.337$ among college graduates) and negative cues ($F = 0.06, p = 0.803$ in the pooled sample; $F = 0.04, p = 0.846$ among college graduates). Collapsing them into a single high-status category effectively doubles the statistical power for high-status cues ($N \approx 575$ per arm) without obscuring any empirical heterogeneity.

To capture respondents’ structural social location, we examine objective educational attainment (College Degree vs. No College) and subjective class self-identification (Middle Class vs. Working Class). Methodologically, we analyze these dimensions stepwise: we first evaluate the separate moderating effects of educational attainment alone and subjective class identification alone (Section 3.2, Figure 3, and Table 2) before examining their intersection in a four-quadrant typology of status consistency: Working Class without College, Working Class with College, Middle Class without College, and Middle Class with College. In the base operationalization, “Middle Class” includes a small minority of respondents who self-identified as “Lower Class” and “Middle Class” includes an even smaller proportion of respondents who self-identify as “Upper Class.” To verify that our substantive conclusions are not sensitive to boundaries of subjective class, we fit models using an alternative categorization that restricts the sample only to self-identified working- and middle-class respondents (excluding lower- and upper-class extremes) as an additional sensitivity analysis in Appendix F.

A key methodological challenge in experimental design examining status-contingent influences stems from the hybrid nature of the design. While the experimental cues (others’ evaluation valence and others’ status) were randomly assigned, respondents’ own social locations—captured by the crossing of educational attainment and subjective class identification—represent pre-existing, observational attributes (Cruces et al., 2013). Consequently, individuals do not have an equal probability of belonging to each status group: for example, Middle Class individuals with a college degree are systematically older, more likely to have college-educated parents, and differ in racial/ethnic composition compared to non-college Working Class individuals. If unaddressed, comparisons across status groups (and their interactions with the randomized treatments) risk confounding the effect of social class with these correlated background characteristics.

To solve this problem and strengthen causal inference, we implement Inverse Probability Weighting (IPW) using generalized propensity scores (Rosenbaum and Rubin, 1983). The logic of IPW proceeds in two steps. First, for each respondent, we estimate a multinomial logistic regression modeling the probability of belonging to their observed status consistency group $S_i \in \{1, 2, 3, 4\}$ conditional on exogenous baseline covariates $\mathbf{X}_i$ (age, gender, race/ethnicity, and parental educational capital): $e_s(\mathbf{X}_i) = P(S_i = s \mid \mathbf{X}_i)$. Second, each observation is weighted by the inverse of its estimated propensity score to target the Average Treatment Effect (ATE) across the pooled sample: $w_i = 1/e_s(\mathbf{X}_i) = 1/P(S_i = s \mid \mathbf{X}_i)$.

Intuitively, IPW creates a synthetic “pseudo-population” by up-weighting underrepresented demographic profiles within each status group (e.g., working-class respondents whose parents completed college) and down-weighting overrepresented profiles. This weighting scheme breaks the empirical association between the baseline covariates $\mathbf{X}_i$ and status group membership. Under the assumption of conditional exchangeability (no unmeasured confounders given $\mathbf{X}_i$) and positivity, IPW ensures that status groups are demographically comparable. As a result, subsequent differences in evaluative shifts, behavioral inertia, or reactance across status categories can be interpreted as the direct consequence of structural status location and status consistency, rather than demographic artifacts.

Figure 2 displays the standardized mean differences (SMD) for all baseline covariates across the status categories before and after Inverse Probability Weighting (IPW), generated using `cobalt` (Greifer, 2022). Prior to weighting, status groups exhibited substantial demographic disparities: for instance, college-educated middle-class respondents were older and had higher rates of parental college attainment (raw SMDs ranging from 0.35 to 0.50). Following IPW estimation, the weighted pseudo-population achieved balance, with all absolute standardized mean differences compressed well below the conventional conservative threshold of 0.10 standard deviations. This confirms that subsequent model estimates isolate the effect of social class and status consistency, net of observed demographic sorting.




Figure 2: Covariate Balance Across Observational Status Groups Before and After IPW Weighting. The dashed lines indicate the standard ± 0.10 threshold for acceptable balance.

2.5 Difference-in-Differences Analytic Strategy

To strictly isolate information-induced shifts from exposure drift, we formulate our analysis within a Difference-in-Differences (DiD) framework. Rather than examining simple pre-to-post changes within each condition in isolation—which conflates the effect of information from others with the natural exposure drift that occurs upon viewing an artwork a second time—the DiD strategy compares every treatment condition against the influence-unexposed Control Condition.

Accordingly, we specify a linear mixed-effects model with random intercepts for respondents of the form:

$$\text{Taste}_{it} = \beta_0 + \beta_1 \text{Trial}_{it} + \sum_{k=2}^K \beta_k \text{Cond}_{ik} + \sum_{k=2}^K \delta_k (\text{Trial}_{it} \times \text{Cond}_{ik}) + u_i + \epsilon_{it} \quad (1)$$

where i indexes individual respondents ($i = 1, \dots, N$), $t \in \{1, 2\}$ indexes measurement trials (Trial 1 pre-treatment baseline vs. Trial 2 post-treatment evaluation), $k \in \{1, \dots, K\}$ indexes the experimental conditions with the unexposed Control Condition ($k = 1$) serving as the omitted reference category, $u_i \sim \mathcal{N}(0, \sigma_u^2)$ is the individual-specific random intercept, and $\epsilon_{it} \sim \mathcal{N}(0, \sigma_\epsilon^2)$ is the trial-level residual error.

The intercept (β_0) represents the expected baseline aesthetic evaluation at Trial 1 (pre-treatment) for respondents assigned to the unexposed Control Condition: $\beta_0 = E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]$. The Trial main effect (β_1) captures the counterfactual *exposure drift* from Trial 1 to Trial 2 in the absence of external information from others: $\beta_1 = E[\text{Taste} \mid \text{Trial} = 2, \text{Control}] - E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]$. When repeated viewing induces an intrinsic appreciation of the artwork (the mere exposure effect), β_1 should be positive and statistically significant.

The condition main effects (β_k) reflect pre-treatment differences between condition k and the Control Condition at Trial 1: $\beta_k = E[\text{Taste} \mid \text{Trial} = 1, \text{Cond} = k] - E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]$. Because respondents were randomly assigned to experimental conditions before viewing the painting, these coefficients assess baseline balance and are expected to be close to zero. The interaction coefficients (δ_k) represent the *causal Difference-in-Differences treatment effects*. Each δ_k measures whether the change in evaluations from Trial 1 to Trial 2 in treatment condition k differed significantly from the change observed in the unexposed Control Condition:

$$\begin{aligned} \delta_k = & (E[\text{Taste} \mid \text{Trial} = 2, k] - E[\text{Taste} \mid \text{Trial} = 1, k]) \\ & - (E[\text{Taste} \mid \text{Trial} = 2, \text{Control}] - E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]) \end{aligned} \quad (2)$$

Substantively, δ_k answers the counterfactual question: *Did exposure to evaluation k cause respondents' aesthetic judgments to change significantly more or less than they would have changed under mere repeated exposure alone?* A positive estimate ($\delta_k > 0$) indicates that feedback from others facilitated aesthetic appreciation beyond natural exposure drift, whereas a negative estimate ($\delta_k < 0$) indicates that feedback from others suppressed or reversed the positive exposure drift that would have naturally occurred.

Estimating the simple pooled DiD model in Equation 1 without differentiating by social location yields mostly muted or null effects (reported in Appendix Table A.2 and Appendix

Figure A.2). As our introductory considerations suggest, however, this pooled invariance is an artifact of aggregation hiding underlying heterogeneity in the causal effect (Brand and Thomas, 2013). In essence, because aesthetic influence operates relationally based on the status congruence between ego and alter, pooling across socially dissimilar respondents washes out opposing directional adjustments. The substantive dynamics of social influence emerge when we examine models interacting experimental conditions with education, subjective social class, and status consistency.

Instead, to test our focal theoretical hypotheses regarding how social influence is moderated by respondents' structural social location (educational attainment, subjective social class, and the four status consistency quadrants), we extend Equation 1 to a fully moderated Difference-in-Differences specification:

$$\begin{aligned}
\text{Taste}_{it} = & \beta_0 + \beta_1 \text{Trial}_{it} + \sum_{k=2}^K \beta_k \text{Cond}_{ik} + \sum_{m=2}^M \gamma_m \text{Status}_{im} \\
& + \sum_{k=2}^K \delta_k (\text{Trial}_{it} \times \text{Cond}_{ik}) + \sum_{m=2}^M \theta_m (\text{Trial}_{it} \times \text{Status}_{im}) \\
& + \sum_{k=2}^K \sum_{m=2}^M \lambda_{km} (\text{Cond}_{ik} \times \text{Status}_{im}) \\
& + \sum_{k=2}^K \sum_{m=2}^M \psi_{km} (\text{Trial}_{it} \times \text{Cond}_{ik} \times \text{Status}_{im}) + u_i + \epsilon_{it}
\end{aligned} \tag{3}$$

Where Status_{im} represents indicators for respondent i 's structural social location across M categories ($m = 1, \dots, M$, with $m = 1$ denoting the reference group). In this moderated framework, θ_m captures group-specific differences in counterfactual baseline exposure drift in the Control Condition ($\text{Trial}_{it} \times \text{Status}_{im}$), δ_k represents the DiD treatment effect for the reference status group, and ψ_{km} represents the three-way interaction coefficient ($\text{Trial}_{it} \times \text{Cond}_{ik} \times \text{Status}_{im}$) testing whether the causal DiD treatment effect of condition k diverges significantly across status groups. The status-specific causal DiD treatment effect for subgroup m is given by the linear combination $\delta_{km} = \delta_k + \psi_{km}$.⁴

⁴By the Frisch-Waugh-Lovell theorem, estimating Equation 1 (or Equation 3) via linear mixed models with individual random intercepts (u_i) is mathematically equivalent to regressing individual change scores ($\Delta \text{Taste}_i = \text{Taste}_{i2} - \text{Taste}_{i1}$) on the treatment and status interaction terms with the Control Condition as the omitted reference group:

$$\Delta \text{Taste}_i = \beta_1 + \sum_{m=2}^M \theta_m \text{Status}_{im} + \sum_{k=2}^K \delta_k \text{Cond}_{ik} + \sum_{k=2}^K \sum_{m=2}^M \psi_{km} (\text{Cond}_{ik} \times \text{Status}_{im}) + \epsilon_i \tag{4}$$

where the intercept and status main effects estimate the control exposure drift within each status group, and the interaction terms estimate the status-specific DiD treatment effects.

3 Results

3.1 Counterfactual Exposure Drift in the Control Condition

Evaluating the unexposed Control Condition allows us to observe the counterfactual drift in aesthetic judgments over repeated trials in the complete absence of cues from others. Because educational attainment is an observational attribute rather than an experimentally assigned condition, comparisons between college-educated and non-college respondents in the Control Condition may be subject to demographic confounding. To address this, we estimate a repeated-measures linear mixed model within the Control Condition ($N = 177$ respondents, 354 observations) incorporating Inverse Probability Weighting (IPW) to balance baseline age, gender, race/ethnicity, and parental education, using a slight variation of Equation 1:

$$\text{Taste}_{it} = \alpha_0 + \beta_1 \text{Trial}_{it} + \alpha_1 \text{College}_i + \beta_2 (\text{Trial}_{it} \times \text{College}_i) + u_i + \epsilon_{it} \quad (5)$$

With respondents without a college degree serving as the reference category.

The model reveals that the counterfactual exposure drift across trials for respondents without a college degree is statistically indistinguishable from zero ($\beta_1 = +0.032, SE = 0.058, t = 0.55, p = 0.582$). In contrast, the interaction between Trial and College degree is positive and statistically significant ($\beta_2 = +0.247, SE = 0.082, t = 3.00, p = 0.0027^{**}$). Combining these parameters demonstrates that college-educated respondents exhibit substantial upward exposure drift (Total Drift = $+0.279, SE = 0.063, t = 4.43, p < 0.001$). Thus, after balancing baseline demographic covariates, repeated exposure elicits an intrinsic aesthetic appreciation specifically among respondents with institutionalized cultural capital, whereas non-college respondents maintain static evaluations.

3.2 Primary Main Models: Social Influence Moderated by Education and Subjective Class

To examine how social influence operates across distinct dimensions of class and status difference, we first inspect how educational attainment alone (College Degree vs. No College) and subjective social class alone (Working Class vs. Middle Class) moderate the experimental treatments, using a model based on Equation 3. Table 2 presents the complete set of Difference-in-Differences estimates across all cue conditions for both dimensions, and Figure 3 visualizes these treatment effects with 95% confidence intervals compared against the unexposed Control Condition.

Inspecting the separate models reveals that education is the primary structural fault line governing responsiveness to external evaluations (Panel A of Figure 3 and Table 2). Non-college respondents tend to either maintain or positively shift their evaluations. In particular, non-college degree respondents align with an unmarked majority in terms of likes (DiD = $+0.228, SE = 0.100, p = 0.022^*$) and are marginally likely ($p < 0.10$), consistent with the cultural goodwill thesis, to align positively when told that high-status people liked the painting. Puzzlingly, and not predicted by any theory, people without a college degree tend to react to the painting when informed that low-status people *dislike* it by increasing their evaluation of it (DiD = $+0.198, SE = 0.088, p = 0.025^*$).




Figure 3: Difference-in-Differences Treatment Effects Moderated Separately by Educational Attainment (Panel A) and Subjective Social Class (Panel B). Horizontal bars depict point estimates with 95% confidence intervals, compared with the unexposed Control Condition. Colors indicate statistical significance: red ($p < 0.05$), orange ($p < 0.10$), and grey ($p \geq 0.10$).

In sharp contrast, college-educated respondents—whose positive baseline exposure drift was established in Section 3.2—experience substantial evaluative suppression from external negative cues: High-Status Dislike ($\text{DiD} = -0.275, SE = 0.094, p = 0.003^{**}$), Generalized Dislike ($\text{DiD} = -0.302, SE = 0.118, p = 0.010^*$), and Low-Status Dislike ($\text{DiD} = -0.226, SE = 0.101, p = 0.026^*$). In addition, college graduates exhibit Bourdieusian taste abandonment when exposed to Low-Status Likes ($\text{DiD} = -0.248, SE = 0.104, p = 0.017^*$).

By comparison, examining subjective social class in isolation (Panel B of Figure 3) produces a more diffuse picture. Self-identified working-class respondents show no statistically reliable DiD adjustments across any condition, masking substantial underlying heterogeneity. Middle-class self-identifiers display significant negative shifts under High-Status Dislike ($\text{DiD} = -0.182, SE = 0.086, p = 0.035^*$) and Low-Status Like ($\text{DiD} = -0.202, SE = 0.095, p = 0.034^*$), but these effects are attenuated relative to the education-stratified mod-

Table 2: Difference-in-Differences Estimates Moderated Separately by Education and Subjective Class

Treatment Condition	Panel A: Education Alone		Panel B: Subjective Class Alone	
	No College	College Degree	Working Class	Middle Class
High-Status Like	+0.129 [†] (0.078)	-0.079 (0.093)	+0.085 (0.084)	-0.008 (0.086)
Low-Status Like	-0.040 (0.087)	-0.248* (0.104)	-0.057 (0.093)	-0.202* (0.095)
High-Status Dislike	-0.011 (0.078)	-0.275** (0.094)	-0.065 (0.083)	-0.182* (0.086)
Low-Status Dislike	+0.198* (0.088)	-0.226* (0.101)	+0.107 (0.093)	-0.076 (0.095)
Generalized Like	+0.228* (0.100)	-0.082 (0.112)	+0.106 (0.106)	+0.079 (0.106)
Generalized Dislike	+0.000 (0.097)	-0.302* (0.118)	-0.168 (0.104)	-0.084 (0.107)
Control Condition Drift	-0.000 (0.069)	+0.257** (0.081)	+0.065 (0.072)	+0.155* (0.076)
Sample Size (<i>N</i>)	1,340	913	1,099	1,154

[†] $p < 0.10$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Note: Estimates represent Difference-in-Differences treatment contrasts relative to the unexposed Control Condition within each demographic subpopulation. Panel A estimates moderation by educational attainment alone (No College Degree vs. Bachelor’s Degree or Higher). Panel B estimates moderation by subjective social class identification alone (Working Class vs. Middle Class).

els.

Crucially, this stepwise comparison demonstrates why analyzing education and subjective class separately is insufficient. Because objective credential attainment and subjective class identification cross-cut one another—nearly 28% of working-class self-identifiers hold college degrees, while over 28% of middle-class identifiers do not—univariate stratification conflates status-consistent actors with status-inconsistent actors. As we show below, the apparent lack of responsiveness among “working-class” respondents is an artifact of aggregating non-college workers (who exhibit normative inertia) with college-educated workers (who exhibit sharp reactivity against high-status cues). Examining their four-quadrant intersection is therefore essential for isolating the distinct mechanisms of status consistency and cross-status influence.

3.3 Status Consistency and Relational Influence: Hypotheses 1–3

Table 3 presents the complete 24-cell matrix of Difference-in-Differences treatment effect estimates across all experimental conditions and all four status consistency quadrants, estimated using Inverse Probability Weighting (IPW) and benchmarked against the unexposed Control Condition. Figure A.3 in the Appendix displays the corresponding point estimates and 95% confidence intervals.

Evaluating our status-differentiated hypotheses through the DiD framework across these quadrants reveals which theoretical expectations are empirically supported versus unsupported.

First, same-status alignment (*Hypothesis 1*) is not supported. Hypothesis 1 predicted that individuals would align their judgments with the positive and negative evaluations of

Table 3: Difference-in-Differences Treatment Effects Across Status Consistency Quadrants (IPW-Weighted)

Treatment Condition	Working Class Identification		Middle Class Identification	
	No College	College Degree	No College	College Degree
High-Status Like	+0.130 (0.115)	-0.024 (0.121)	+0.038 (0.133)	-0.118 (0.126)
Low-Status Like	-0.110 (0.132)	+0.028 (0.133)	-0.111 (0.146)	-0.296* (0.144)
High-Status Dislike	+0.006 (0.115)	-0.394*** (0.119)	-0.209 (0.132)	-0.195 (0.128)
Low-Status Dislike	+0.185 (0.132)	-0.171 (0.131)	+0.089 (0.149)	-0.227† (0.136)
Generalized Like	+0.138 (0.158)	-0.125 (0.145)	+0.338† (0.173)	-0.055 (0.152)
Generalized Dislike	-0.120 (0.140)	-0.396* (0.154)	+0.134 (0.168)	-0.264† (0.151)
Control Exposure Drift	-0.013 (0.099)	+0.311** (0.104)	+0.098 (0.120)	+0.241* (0.111)
Analytic Sample (N)	764	335	576	577

Note: † $p < 0.10$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Standard errors in parentheses. Estimates represent Difference-in-Differences treatment contrasts relative to the unexposed Control Condition within each status consistency quadrant, estimated using Inverse Probability Weighting (IPW) balancing baseline age, gender, race/ethnicity, and parental education.

others sharing their status. When evaluated relative to the control condition counterfactual, neither high-status consistent respondents ($\text{DiD} = -0.118, SE = 0.126, p = 0.348$) nor low-status consistent respondents ($\text{DiD} = -0.110, SE = 0.132, p = 0.404$) show a statistically significant DiD shift when exposed to same-status likes. Similarly, low-status respondents exposed to same-status dislikes do not exhibit a statistically significant shift ($\text{DiD} = +0.185, SE = 0.132, p = 0.163$). Because none of these estimates are statistically significant ($p \geq 0.10$), the data do not support Hypothesis 1.

Second, symbolic power and upward alignment (*Hypothesis 2*) is not supported for low-status consistent individuals. Hypothesis 2 predicted that low-status respondents would align with high-status evaluations. Among low-status consistent respondents (Working Class/No College), neither exposure to High-Status Likes ($\text{DiD} = +0.130, SE = 0.115, p = 0.256$) nor High-Status Dislikes ($\text{DiD} = +0.006, SE = 0.115, p = 0.959$) produces a statistically significant change. Thus, Hypothesis 2 is not supported for low-status consistent individuals. However, for status-inconsistent workers who hold a college degree (Working Class/College), exposure to High-Status Dislikes produces a statistically significant negative DiD shift of $\text{DiD} = -0.394 (SE = 0.119, t = -3.30, p = 0.001^{**})$, alongside a significant negative shift under Generalized Dislikes ($\text{DiD} = -0.396, SE = 0.154, p = 0.010^{*}$). This indicates that sensitivity to high-status negative evaluation is concentrated among individuals who possess institutional cultural capital but identify with the working class.

Third, cross-status reactance and taste abandonment (*Hypothesis 3*) is strongly supported. Netting out the counterfactual exposure drift provides direct empirical support for Hypothesis 3. For high-status consistent respondents (Middle Class/College), exposure to Low-Status Likes (working-class others liking the painting) produces a statistically significant negative DiD divergence of $\text{DiD} = -0.296 (SE = 0.144, t = -2.06, p = 0.039^{*})$ relative to the control condition. Because college-educated middle-class respondents naturally drift upward by +0.200 under mere exposure, discovering that working-class others liked the art-

work actively reverses their appreciation. High-status consistent respondents also exhibit marginal negative shifts under Generalized Dislikes ($\text{DiD} = -0.264, SE = 0.151, p = 0.081^\dagger$) and Low-Status Dislikes ($\text{DiD} = -0.227, SE = 0.136, p = 0.097^\dagger$). This confirms that high-status individuals distance their judgments from low-status endorsements, providing experimental evidence of Bourdieusian symbolic distinction.

Finally, regarding status-inconsistent positive influence, among Middle Class individuals without a college degree, exposure to Generalized Likes produces a marginally significant positive DiD shift of $\text{DiD} = +0.338 (SE = 0.173, t = 1.96, p = 0.050^\dagger)$. All other conditions for this group remain statistically indistinguishable from zero ($p \geq 0.11$).

3.4 Discrete Behavioral Choices: Inertia, Conformity, and Reactance

While continuous change models identify average shifts across conditions, they mask non-linear behavioral choices where opposing individual movements cancel each other out in the aggregate. For instance, an average shift close to zero could reflect complete evaluative rigidity, or it could reflect equal proportions of respondents conforming and reacting in opposite directions. To examine these discrete underlying choices, we categorize responses into three mutually exclusive behaviors: *Stay* ($\Delta = 0$, maintaining one’s initial evaluation), *Conform* (moving in the direction of others’ evaluation), and *React* (moving in the opposite direction of others’ evaluation) and use this as the outcome variable in a multinomial logit model with the four status consistency categories as predictors.

Figure 4 presents the Average Marginal Effects (AME) on the probability of each behavioral choice from an IPW-weighted multinomial logistic regression, compared against the reference group of Working Class respondents without a college degree. In the baseline reference group, the predicted probabilities are 74.4% for staying, 13.6% for conforming, and 12.0% for reacting.

The marginal effects provide compelling empirical confirmation of *Hypothesis 4* and offer new insights into the social dynamics of status inconsistency. First, in terms of evaluative stability (Panel A: Stay), structural status consistency acts as the primary anchor of aesthetic confidence and inertia. Middle-class individuals with a college degree are significantly more likely to maintain their initial judgment than non-college workers ($\text{AME} = +0.046, SE = 0.013, z = 3.64, p < 0.001$), with an overall stay rate of 79.0%. In sharp contrast, status-inconsistent individuals without a college degree exhibit significantly lower stability ($\text{AME} = -0.054, SE = 0.013, z = -4.05, p < 0.001$), with their stay rate falling to 65.5%. College-educated workers also display attenuated stability ($\text{AME} = -0.022, p = 0.105$, stay rate 68.2%). This shows that evaluative inertia is not a simple linear function of high status; rather, both high-consistent (Middle/College) and low-consistent (Working/No College) individuals occupy unambiguous positions within the social structure that foster normative certainty, whereas status inconsistency breeds evaluative fluidity.

Second, in terms of conformity (Panel B), status inconsistency emerges as the primary driver of susceptibility to external social influence. Both status-inconsistent groups display large, highly significant increases in the probability of conforming to others’ judgments. Working Class individuals with a college degree are +5.4% more likely to conform relative




Figure 4: Average Marginal Effects on Behavioral Choices (Stay, Conform, React) Across Status-Consistency Groups. Horizontal bars indicate differences in choice probability relative to the baseline reference category (Working Class, No College), with 95% confidence intervals. Colors indicate statistical significance: red ($p < 0.05$) and grey ($p \geq 0.10$).

to the reference group ($AME = +0.054, SE = 0.011, z = 4.86, p < 0.001$), and Middle Class individuals without college are +4.1% more likely to conform ($AME = +0.041, SE = 0.011, z = 3.82, p < 0.001$). In stark contrast, high-status consistent individuals are significantly less likely to conform ($AME = -0.021, SE = 0.010, z = -2.11, p = 0.035$). This shows that when actors occupy contradictory locations in social space—possessing formal educational credentials while identifying as working class, or claiming middle-class status without institutional degrees—they experience evaluative cross-pressures and cognitive ambiguity that heighten their inclination to align with external social signals.

Finally, regarding oppositional reactance (Panel C), behavioral divergence is strongly structured by educational capital. Both college-educated groups are significantly less likely to move away from others' judgments compared to non-college workers: Working Class respondents with a college degree are -3.2% less likely to react ($AME = -0.032, SE = 0.009, z = -3.54, p < 0.001$), and Middle Class college graduates are -2.5% less likely to react ($AME = -0.025, SE = 0.009, z = -2.70, p = 0.007$). Meanwhile, middle-class individuals without a college degree do not differ significantly from their working-class peers ($AME = +0.013, p = 0.191$). This shows that oppositional reactance to external cues is concentrated among individuals without a college education, whereas institutionalized cultural capital inhibits oppositional defiance.

To confirm that these behavioral dynamics are driven specifically by status-laden social cues rather than by generic feedback about others' opinions, we re-estimated the multinomial model excluding the two Taste-Only conditions ($N = 1,870$). As displayed in Appendix Figure A.4, the results are virtually identical: structural status consistency continues to anchor evaluative inertia (Middle Class with College: +4.8%, $p < 0.001$), status inconsistency heightens conformity (Working Class with College: +6.0%, $p < 0.001$; Middle Class without College: +4.3%, $p < 0.001$), and oppositional reactance remains concentrated among non-

college workers (-3.0% and -2.1% for college-educated groups, $p < 0.05$). Netting out generic popularity cues therefore confirms that status consistency uniquely structures discrete behavioral responses to social influence.

4 Discussion

4.1 Summary of Key Results

In this study, we investigated how access to others’ aesthetic evaluations causally influences individual taste judgments in real time. Using a national survey experiment with a Difference-in-Differences design and inverse probability weighting, we resolved several ambiguities present in observational literature. Table 4 provides a systematic overview of our focal hypotheses, the empirical DiD point estimates, and the resulting empirical verdicts.

As summarized in Table 4, the DiD framework clarifies which classic sociological intuitions withstand experimental scrutiny. First, we established that repeated evaluation produces a positive exposure drift, but this drift is contingent on educational capital: college-educated individuals experience an intrinsic upward appreciation upon second exposure ($\beta_1 = +0.257, p = 0.016$), whereas non-college individuals do not ($\beta_1 = 0.000, p = 1.00$).

Second, our findings show strong empirical evidence for *valence asymmetry*. In the pooled sample, negative evaluations from others and low-status endorsements produce statistically reliable suppressive shifts relative to the control condition ($p < 0.05$ to $p < 0.10$), actively halting exposure drift. In contrast, positive evaluations from others failed to produce statistically significant facilitative shifts in the aggregate ($p \geq 0.10$), indicating that cultural distastes and out-group endorsements exert a far more dependable veto on aesthetic appreciation than positive cues from others exert on facilitation.

Third, the DiD framework reveals that two widely cited mechanisms—same-status alignment (*Hypothesis 1*) and upward cultural goodwill by low-status actors (*Hypothesis 2*)—are not supported once counterfactual exposure drift is accounted for. For low-status consistent individuals, exposure to same-status likes (DiD = $-0.110, p = 0.404$) and high-status likes (DiD = $+0.130, p = 0.256$) remain statistically indistinguishable from zero. Instead, the primary status-driven dynamic is cross-status reactance (*Hypothesis 3*): high-status consistent respondents exposed to low-status likes exhibit a statistically significant negative divergence (DiD = $-0.296, p = 0.039$), providing clear experimental evidence of Bourdieusian symbolic distinction and taste abandonment.

Finally, discrete behavioral modeling provides robust confirmation of *Hypothesis 4*. Aesthetic stability (“staying”) is governed by structural status consistency rather than high status alone: both high-consistent and low-consistent individuals exhibit the highest inertia, whereas status-inconsistent individuals display heightened fluidity and conformity.

4.2 Limitations and Suggestions for Future Work

Several limitations of this study suggest directions for future research. In terms of operationalization and measurement, our experiment relied on a single cultural stimulus (Whistler’s *Nocturne*), selected for its high baseline variance. Future studies should examine whether so-

Table 4: Summary of Hypotheses, Difference-in-Differences Estimates, and Empirical Verdicts

Hypothesis / Mechanism	Theoretical Expectation	Empirical DiD Estimate / Statistic	Empirical Verdict
<i>Mere Exposure Drift (Control)</i>	Upward drift in evaluation across repeated trials without external social cues	$\beta_1 = +0.107$ ($p = 0.041^*$)	Supported ($p < 0.05$)
<i>Valence Asymmetry</i>	Negative cues from others halt exposure drift; positive cues facilitate drift	High-Status Dislike: DiD = -0.121 ($p = 0.043^*$); Generalized Dislike: DiD = -0.125 ($p = 0.095^\dagger$); Positive cues: null ($p \geq 0.198$)	Supported ($p < 0.05$)
<i>Hypothesis 1: Same-Status Alignment</i>	Individuals align judgments with aggregate evaluations of others sharing their status	High-status like: DiD = -0.118 ($p = 0.348$); Low-status like: DiD = -0.110 ($p = 0.404$)	Not Supported ($p \geq 0.10$)
<i>Hypothesis 2: Symbolic Power</i>	Lower-status individuals align judgments with high-status others	Low-status consistent like: DiD = $+0.130$ ($p = 0.256$); Working/College dislike: DiD = -0.394 ($p = 0.001^{**}$)	Not Supported for Low-Status Consistent ($p \geq 0.10$)
<i>Hypothesis 3: Cross-Status Reactance</i>	Individuals distance judgments from outgroup evaluations	High-status exposed to Low-Status Like: DiD = -0.296 ($p = 0.039^*$); Low-Status Dislike: DiD = -0.227 ($p = 0.097^\dagger$)	Supported ($p < 0.05$)
<i>Hypothesis 4: Status Consistency Inertia</i>	Consistent status profiles exhibit higher stability (“staying”) than inconsistent profiles	Stay rates: 79.0% (Middle/College) & 74.4% (Working/No College) vs. 65.5% & 68.2% for inconsistent groups	Supported ($p < 0.001$)

Note: Difference-in-Differences estimates control treatment shifts against the unexposed Control Condition. Status-subgroup estimates and behavioral models incorporate Inverse Probability Weighting (IPW) on age, gender, race/ethnicity, and parental education.

cial influence dynamics operate similarly across abstract, figurative, and commercial cultural forms.

Regarding threats to causal inference and contextual confounding, while our design randomized information about others, respondents’ status locations were observational. Although inverse probability weighting balanced observed demographic covariates, unmeasured psychological dispositions could contribute to status-group differences. Future research using laboratory status induction could provide complementary evidence.

Third, our study examined the unidirectional influence of aggregate fictional others on individuals. In everyday life, cultural influence operates bidirectionally within evolving social networks. Investigating how reciprocal feedback loops shape aesthetic consensus represents an important next step.

Fourth, regarding temporal granularity, our experiment measured short-term evaluative shifts during a survey session. Longitudinal studies tracking whether experimentally induced taste shifts persist over weeks or months would help determine the durability of these social influence effects.

Finally, regarding institutional scope conditions, our sample was recruited via a quota-matched online panel of U.S. adults. While providing greater demographic diversity than student subject pools, cross-national replications are needed to assess how national cultural institutions shape status distinction and conformity.

4.3 Implications: Status Consistency and Evaluative Fluidity

Our findings carry broad implications for sociological theories of taste and stratification. Classic cultural sociology has often framed aesthetic boundary work in vertical terms: high-status actors use exclusive tastes to mark distinction, while lower-status actors display cultural goodwill. Our results qualify this model by showing that cultural influence is bounded by valence asymmetry and structural consistency.

Most importantly, our analysis highlights status consistency as a neglected driver of aesthetic orientations. Status-consistent individuals—both high and low—display high evaluative stability, anchored in unambiguous structural locations. In contrast, status-inconsistent individuals experience normative cross-pressures that produce evaluative fluidity. By examining how structural consistency interacts with external social signals, cultural sociology can move beyond static homophily toward dynamic, interactional models of aesthetic judgment.

References

- Aral, S., Muchnik, L., and Sundararajan, A. (2009). Distinguishing influence-based contagion from homophily-driven diffusion in dynamic networks. *Proceedings of the National Academy of Sciences of the United States of America*, 106(51):21544–21549.
- Berger, J. and Heath, C. (2008). Who drives divergence? identity signaling, outgroup dissimilarity, and the abandonment of cultural tastes. *Journal of Personality and Social Psychology*, 95(3):593–607.
- Bloom, H. S. (1995). Minimum detectable effects: A simple way to measure the statistical precision of field experiments. *Evaluation Review*, 19(5):547–556.
- Bourdieu, P. (1984). *Distinction: A social critique of the judgement of taste*. Harvard university press.
- Brand, J. E. and Thomas, J. S. (2013). Causal effect heterogeneity. In *Handbook of causal analysis for social research*, pages 189–213. Springer.

- Bryson, B. (1996). “anything but heavy metal”: Symbolic exclusion and musical dislikes. *American Sociological Review*, 61(5):884–899.
- Centola, D. (2011). An experimental study of homophily in the adoption of health behavior. *Science*, 334(6060):1269–1272.
- Cohen, J. (1988). *Statistical Power Analysis for the Behavioral Sciences*. Lawrence Erlbaum Associates, Hillsdale, NJ, 2nd edition.
- Cruces, G., Perez-Truglia, R., and Tetaz, M. (2013). Biases in perceptions of sorting: How home-ownership and inequality shape beliefs about position. *Journal of Public Economics*, 98:100–112.
- Cutting, J. E. (2003). Gustave caillebotte, french impressionism, and mere exposure. *Psychonomic Bulletin & Review*, 10(2):319–343.
- DiMaggio, P. (1987). Classification in art. *American Sociological Review*, pages 440–455.
- Erickson, B. H. (1996). Culture, class, and connections. *American Journal of Sociology*, 102(1):217–251.
- Gelman, A. and Carlin, J. (2014). Beyond power calculations: Assessing type s (sign) and type m (magnitude) errors. *Perspectives on Psychological Science*, 9(6):641–651.
- Greifer, N. (2022). *cobalt: Covariate Balance Tables and Plots*. R package version 4.4.1.
- Hoenig, J. M. and Heisey, D. M. (2001). The abuse of power: The pervasive fallacy of analyzing power for data analytic results. *The American Statistician*, 55(1):19–24.
- Hope, K. (1975). Models of status inconsistency and social mobility effects. *American Sociological Review*, pages 322–343.
- Lamont, M. and Molnár, V. (2002). The study of boundaries in the social sciences. *Annual review of sociology*, 28(1):167–195.
- Lewis, K. and Kaufman, J. (2018). The conversion of cultural tastes into social network ties. *The American journal of sociology*, 123(6):1684–1742. doi: 10.1086/697525.
- Lizardo, O. (2006). How cultural tastes shape personal networks. *American Sociological Review*, 71(5):778–807.
- Reber, R., Winkielman, P., and Schwarz, N. (1998). Effects of perceptual fluency on affective judgments. *Psychological science*, 9(1):45–48.
- Rosenbaum, P. R. and Rubin, D. B. (1983). The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41–55.
- Salganik, M. J., Dodds, P. S., and Watts, D. J. (2006). Experimental study of inequality and unpredictability in an artificial cultural market. *Science*, 311(5762):854–856.

- Shalizi, C. R. and Thomas, A. C. (2011). Homophily and contagion are generically confounded in observational social network studies. *Sociological methods & research*, 40(2):211–239.
- Simmel, G. (1957). Fashion. *American journal of sociology*, 62(6):541–558.
- Vanzella-Yang, A. and Abrutyn, S. (2022). Tastes, emotions, and social cohesion: Toward a cultural theory of social exchange. *Journal for the Theory of Social Behaviour*, 52(2):315–335.
- Wohl, H. (2015). Community sense: The cohesive power of aesthetic judgment. *Sociological Theory*, 33(4):299–326.
- Zajonc, R. B. (1968). Attitudinal effects of mere exposure. *Journal of Personality and Social Psychology*, 9(2p2):1–27.

A Statistical Power and Minimum Detectable Effect (MDE) Analysis

To rigorously evaluate the statistical precision and empirical detection boundaries of the study, we conducted a formal ex-post Minimum Detectable Effect (MDE) analysis (Bloom, 1995). Methodological scholarship has established that retrospective or “post-hoc” power calculations based on observed effect sizes and sample estimates are fundamentally flawed and uninformative, functioning merely as a tautological transformation of the observed p -value (Gelman and Carlin, 2014; Hoenig and Heisey, 2001). By contrast, an ex-post MDE analysis answers a principled design question: given the achieved cell sample sizes and the empirical variance structure of the repeated-measures data, what is the smallest true population effect size that the research design was adequately powered (at 80% or 90% power, two-tailed $\alpha = 0.05$) to detect?

In a repeated-measures Difference-in-Differences (DiD) framework where each respondent i provides an initial pre-treatment aesthetic rating (Y_{i1}) and a post-treatment evaluation (Y_{i2}) on a 7-point Likert scale, the standard error of the change score $\Delta Y_i = Y_{i2} - Y_{i1}$ depends directly on the test-retest correlation ρ between trials:

$$\sigma_{\Delta} = \sqrt{\sigma_1^2 + \sigma_2^2 - 2\rho\sigma_1\sigma_2} \quad (6)$$

When the test-retest correlation is high, within-subject differencing cancels out stable individual evaluative variance, yielding a change-score standard deviation substantially smaller than the cross-sectional baseline standard deviation ($\sigma_{\Delta} \ll \sigma_1$). In our survey sample, aesthetic judgments exhibit strong stability across repeated trials ($r = 0.8945$, with $\sigma_1 = 1.500$ and $\sigma_2 = 1.548$), yielding an empirical standard deviation of the change score, $\sigma_{\Delta} = 0.702$. This within-subject autocorrelation reduces the standard error of treatment comparisons by more than 50% relative to an independent post-test design without pre-treatment baselines.

For a two-sample Difference-in-Differences comparison between a treatment condition of size N_T and the unexposed Control Condition of size N_C , the standard error is:

$$SE_{\text{DiD}} = \sigma_{\Delta} \sqrt{\frac{1}{N_T} + \frac{1}{N_C}} \quad (7)$$

The Minimum Detectable Effect at significance level α (two-tailed) and statistical power $1 - \beta$ is calculated as:

$$\text{MDE}_{\text{raw}} = (z_{1-\alpha/2} + z_{1-\beta}) \times SE_{\text{DiD}} \quad (8)$$

where $z_{0.975} = 1.960$ for $\alpha = 0.05$, $z_{0.80} = 0.842$ for 80% power (multiplier = 2.802), and $z_{0.90} = 1.282$ for 90% power (multiplier = 3.242). To facilitate comparison across studies, effect sizes are expressed both in raw 7-point Likert scale units (MDE_{raw}) and as standardized effect sizes relative to baseline variance, Cohen’s $d = \text{MDE}_{\text{raw}}/\sigma_1$ (Cohen, 1988).

Table A.1 presents the complete MDE calculations across all analytical tiers of the experimental design, and Figure A.1 visualizes the theoretical precision curves alongside the empirical thresholds achieved in this study.

The empirical findings from the MDE analysis yield three primary methodological take-aways. First, regarding repeated-measures exposure drift (Panel A of Table A.1), the full

Table A.1: Ex-Post Minimum Detectable Effect (MDE) Analysis Across Experimental Specifications

Model Specification / Level	N_T	N_C	SE	80% Power		90% Power	
				MDE _{raw}	Cohen's d	MDE _{raw}	Cohen's d
<i>Panel A: Within-Subject Exposure Drift (Paired Comparison)</i>							
Full Sample Drift	2,253	—	0.015	0.041	0.028	0.048	0.032
Control Condition Drift	177	—	0.053	0.148	0.099	0.171	0.114
<i>Panel B: Pooled Difference-in-Differences Contrasts (vs. Control $N_C = 177$)</i>							
Taste-Only Conditions	172	177	0.075	0.210	0.140	0.244	0.162
Single Status Conditions	290	177	0.067	0.188	0.125	0.217	0.145
Pooled High-Status Conditions	575	177	0.060	0.169	0.113	0.196	0.130
<i>Panel C: Subgroup DiD Contrasts Within Status Consistency Quadrants</i>							
Working Class, No College	60–214	70	0.108	0.304	0.203	0.352	0.234
Middle Class, College	51–164	51	0.124	0.347	0.231	0.401	0.267
Middle Class, No College	32–146	33	0.150	0.421	0.281	0.487	0.325
Working Class, College	16–72	23	0.186	0.522	0.348	0.604	0.403

Note: Minimum Detectable Effects (MDE) are calculated assuming a two-tailed $\alpha = 0.05$ with standard statistical power thresholds of 80% ($z = 2.802$) and 90% ($z = 3.242$). Calculations incorporate the empirical test-retest correlation across repeated trials ($r = 0.8945$, yielding $\sigma_\Delta = 0.702$ relative to baseline $\sigma_1 = 1.500$). MDE_{raw} is expressed in units of the 7-point aesthetic rating scale; Cohen’s $d = \text{MDE}_{\text{raw}}/\sigma_1$. In Panel C, N_T indicates the range of treatment cell sizes across specific cue conditions, and reported SE and MDE values correspond to median cell sizes.

sample ($N = 2,253$) possessed extraordinary statistical power, capable of detecting shifts as minute as Cohen’s $d \geq 0.028$ (0.041 scale points). Crucially, even within the unexposed Control Condition ($N_C = 177$), the design was powered at 80% to detect a drift of $d \geq 0.099$ (0.148 scale points). This high baseline precision enabled us to reliably identify the empirical +0.107 exposure drift in the Control Condition ($t = 2.05, p = 0.041$), confirming that the upward drift was not an artifact of random noise.

Second, for the primary pooled Difference-in-Differences models compared against the unexposed Control Condition (Panel B of Table A.1), the design possessed 80% power to detect small treatment effects: $d \geq 0.113$ (0.169 points) for pooled high-status conditions ($N_T = 575$), $d \geq 0.125$ (0.188 points) for single-status conditions ($N_T \approx 290$), and $d \geq 0.140$ (0.210 points) for taste-only conditions ($N_T = 172$). All three pooled thresholds fall comfortably below Cohen’s standard threshold for a “small” effect size ($d = 0.20$). This confirms that the pooled models were well-equipped to detect subtle social influence cues.

Third, within the disaggregated status-consistency quadrants (Panel C of Table A.1), cell sizes are necessarily smaller, resulting in wider detection envelopes. In the two larger, status-consistent quadrants—Working Class without College ($N_C = 70$) and Middle Class with College ($N_C = 51$)—the median MDE thresholds were $d = 0.203$ (0.304 scale points) and $d = 0.231$ (0.347 scale points), providing adequate power for small-to-moderate effects. In contrast, in the smaller, status-inconsistent quadrants—Middle Class without College ($N_C = 33$) and Working Class with College ($N_C = 23$)—the median MDE thresholds increased to




Figure A.1: Minimum Detectable Effect (MDE) Power Analysis. Panel A displays theoretical MDE curves (Cohen's d) at 80% power ($\alpha = 0.05$) across cell sample sizes, illustrating the precision advantage of repeated-measures DiD across varying levels of test-retest autocorrelation (r). Panel B displays the empirical MDE thresholds achieved in this study across within-subject drift, pooled DiD models, and status-differentiated subgroup comparisons.

$d = 0.281$ (0.421 points) and $d = 0.348$ (0.522 points). Consequently, statistically significant shifts in these cells emerged only when behavioral adjustments were pronounced, such as the -0.394 negative divergence ($p < 0.001$) observed among college-educated working-class respondents exposed to High-Status Dislike cues. Methodologically, this indicates that non-significant coefficients in the smallest status cells should be interpreted as reflecting greater estimation uncertainty rather than definitive evidence of null effects.

B Survey Vignettes and Treatment Prompts

Below we transcribe the verbatim wording, instructions, question stems, and response options of the survey instrument administered in the 2012 study, based on the original Qualtrics protocol.

1. Baseline Measurement (Trial 1, Survey Page 4)

Respondents were presented with the high-resolution image of James McNeill Whistler's *Nocturne: Blue and Gold – Old Battersea Bridge* (1872) accompanied by the following verbatim introductory instruction and question prompt:

Screen Text / Instruction:

"Please look at this painting for a few seconds."

Question Prompt:

“Give us your first impression of this painting – how much do you like it?”

Responses were recorded using the following 8 single-choice options:

- 1 = I like it very much
- 2 = I like it
- 3 = I like it just a little
- 4 = I neither like it nor dislike it
- 5 = I dislike it just a little
- 6 = I dislike it
- 7 = I dislike it very much
- 8 = I can’t see the image

Items were reverse-coded during data processing (1 = “dislike very much”, 4 = “neither like nor dislike”, 7 = “like very much”) so that higher numerical values indicate greater aesthetic affinity. Respondents who selected “I can’t see the image” ($N = 6$) were excluded from the analytic sample.

2. Experimental Treatment Vignettes (Survey Page 5)

Following their initial evaluation, respondents in the experimental conditions were exposed to feedback regarding the evaluations of others. Notably, the prompt did *not* describe the evaluating group in relative terms to the participant; instead, it presented absolute descriptive characteristics of the alter group, allowing social proximity or distance to emerge relationally from the participant’s own social location.

A. Status-Differentiated Conditions (Class & Taste): Respondents assigned to conditions crossing status and valence read:

“Thank you for your response. Our survey so far indicates that the group of people most likely to say that they [LIKE / DISLIKE] the painting you saw have completed...”

where the alter characteristics varied by experimental condition:

- **Low-Status Others (LCC):** *“...high school but did not go to college, and work in a food service job.”*
- **High-Status Economic Specialists (ES):** *“...a Bachelor’s degree in Business but no additional degrees, and work at a financial investment firm.”*

- **High-Status Cultural Specialists (CS):** “... several advanced degrees including a PhD, and work at a college or university.”

Immediately following the vignette, respondents answered a subjective similarity probe:

“Given only what you know about this group’s education, occupation, and opinion about the painting, how similar do you think you are to the people in this group?”

with response options: *Very similar, Somewhat similar, Just a little similar, Not very similar, or Not at all similar.*

B. Undifferentiated Taste-Only Conditions (Control 2): Respondents in the Taste-Only conditions were told:

“Thank you for your response. Other people who have taken this survey recently have said that they [LIKE / DISLIKE] the painting you saw.”

“Given only what you know about these people’s opinion about the painting, how similar do you think you are to them?”

with the identical 5-point similarity response options.

C. Unexposed Control Condition (No Alter): Respondents assigned to the pure Control Condition received no feedback screen, evaluation valence, or information about others whatsoever, proceeding directly from Trial 1 to the intervening survey block.

3. Alter Attribution Probe (Survey Page 6, Treatment Respondents Only)

Immediately after the similarity question, respondents in the status-differentiated conditions completed an attribution check:

“In the last question, when you thought about how similar you are to the people in this group, which factor was most influential?”

with response options:

- Their education and/or occupation
- Their opinion about the painting
- A combination of both

4. Intervening Survey Battery (Cognitive Buffer, Survey Pages 7–15)

Between Trial 1 and Trial 2, all respondents completed substantive survey modules lasting 15 to 20 minutes, comprising full educational history (highest degree attained and college major field), intergenerational educational capital (parental bachelor’s degree attainment), employment status, detailed 22-category Standard Occupational Classification (SOC), annual personal pre-tax income (23 income brackets), subjective social class self-identification (Lower class, Working class, Middle class, Upper class), and an extensive musical taste battery assessing liking, listening frequency, favorite genre, and perceived audience demographic stereotypes across 20 distinct musical genres.

5. Post-Treatment Measurement (Trial 2, Survey Page 16)

Following the intervening modules, respondents were re-shown the identical artwork and asked:

Screen Text / Instruction:

“Please look at this painting for a few seconds again.”

Question Prompt:

“Give us your final impression of this painting – how much do you like it?”

recorded using the identical 8 single-choice options (from 1 = “I like it very much” to 7 = “I dislike it very much”, plus “I can’t see the image”), reverse-coded on the same 1–7 scale.

C Comparison of Unweighted and IPW-Weighted DiD Models

Table A.2 reports the side-by-side comparison between the primary unweighted DiD model and the IPW-weighted model balancing baseline age, gender, race/ethnicity, and parental education across the pooled sample. Both specifications produce consistent point estimates, standard errors, and substantive conclusions regarding valence asymmetry and exposure drift suppression.

D Complete Status-Differentiated DiD Estimates

Figure A.3 displays the complete 24-cell grid of Difference-in-Differences treatment effect estimates across all six treatment conditions and all four status consistency quadrants, estimated using Inverse Probability Weighting (IPW). Bars indicate point estimates and error bars denote 95% confidence intervals relative to the unexposed Control Condition within each status group.

Table A.2: Comparison of Unweighted and IPW-Weighted Difference-in-Differences Models (Pooled Sample)

Independent Variable	Unweighted DiD		IPW-Weighted DiD	
	Estimate (SE)	<i>t</i> value	Estimate (SE)	<i>t</i> value
Intercept (β_0)	4.689 (0.114)	40.99***	4.672 (0.115)	40.72***
Trial 2 Exposure Drift (β_1)	0.107 (0.052)	2.05*	0.156 (0.054)	2.90**
<i>Difference-in-Differences Treatment Effects (Trial $\times$ Condition):</i>				
Class & Taste / Like / +Status	+0.041 (0.060)	0.68	+0.006 (0.062)	0.09
Class & Taste / Like / -Status	-0.128 (0.067)	-1.93 [†]	-0.112 (0.069)	-1.61
Class & Taste / Dislike / +Status	-0.121 (0.060)	-2.02*	-0.201 (0.062)	-3.27**
Class & Taste / Dislike / -Status	+0.016 (0.066)	0.23	-0.034 (0.068)	-0.49
Taste Only / Like	+0.096 (0.075)	1.29	+0.068 (0.078)	0.87
Taste Only / Dislike	-0.125 (0.075)	-1.67 [†]	-0.172 (0.076)	-2.26*
Condition Main Effects Included	Yes		Yes	
Random Intercept Variance (σ_u^2)	2.074		2.026	
Residual Variance (σ_e^2)	0.243		1.022	
<i>N</i> (Observations / Individuals)	4,506 / 2,253		4,504 / 2,252	

Note: [†] $p < 0.10$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Reference condition: unexposed Control Condition.

E Sensitivity Analysis: Behavioral Choices Excluding Taste-Only Conditions

To verify that the discrete behavioral patterns documented in the main text (Figure 4) are specifically generated by status-differentiated social influence rather than by undifferentiated taste-consensus cues, we conducted a sensitivity analysis by excluding respondents assigned to the two Taste-Only conditions ($N = 344$). The resulting analytic sample ($N = 1,870$) restricts behavioral classification to respondents exposed to feedback that combines evaluation valence and others' occupational/educational status, alongside unexposed control stayers.

Figure A.4 displays the Average Marginal Effects from this sensitivity specification, compared against Working Class respondents without a college degree. As shown in the figure, all substantive patterns persist with exceptional stability: status-consistent college-educated respondents exhibit elevated evaluative stability (AME = +0.048, $p < 0.001$), status-inconsistent respondents exhibit heightened conformity (AME = +0.060, $p < 0.001$ for Working Class with College; AME = +0.043, $p < 0.001$ for Middle Class without College), and college education sharply reduces oppositional reactance (AME = -0.030, $p = 0.003$ and AME = -0.021, $p = 0.042$). Notably, under this status-restricted specification, the lower evaluative stability of status-inconsistent workers is statistically significant (AME = -0.030, $p = 0.038$), reinforcing our core finding that status crystallization anchors aesthetic judgments against external social influence.




Figure A.2: Difference-in-Differences Treatment Effects in the Pooled Sample (Unweighted Linear Mixed Model). Horizontal bars denote point estimates with 95% confidence intervals relative to the unexposed Control Condition (dashed line at zero). Colors indicate statistical significance: red ($p < 0.05$), orange ($p < 0.10$), and grey ($p \geq 0.10$).

F Robustness Check: Alternative Definition of Social Class

Restricting the subjective class definition strictly to respondents self-identifying as “working class” or “middle class” (excluding lower and upper class extremes) produces identical conclusions: status-consistent groups maintain the highest stay rates (Middle/College: 78.2%; Working/No College: 74.7%) and status-inconsistent groups maintain the highest conformity rates (Middle/No College: 17.9%; Working/College: 16.9%).




Figure A.3: Complete Grid of Difference-in-Differences Treatment Effects by Status Consistency Group (IPW-Weighted). Estimates benchmark condition shifts against the unexposed Control Condition. Colors indicate statistical significance: red ($p < 0.05$), orange ($p < 0.10$), and grey ($p \geq 0.10$).




Figure A.4: Sensitivity Analysis: Average Marginal Effects on Behavioral Choices (Stay, Conform, React) Across Status-Consistency Groups, Excluding Taste-Only Conditions ($N = 1,870$). Horizontal bars indicate differences in choice probability relative to Working Class respondents without a college degree, with 95% confidence intervals. Colors indicate statistical significance: red ($p < 0.05$) and grey ($p \geq 0.10$).