

Social Influence Effect on Aesthetic Judgments

Omar Lizardo & Sara Skiles

June 2026

Abstract

In this research note, we address the question, foundational to sociological studies of taste, of the extent to which others' aesthetic judgments influence the transient evaluation of cultural works, and the extent to which this influence depends on the social distance between self and other. Drawing on classic and recent cultural theory, we hypothesize that high (low) status respondents (as indexed by educational attainment) will be more likely to change their judgment of a cultural object in a negative direction when exposed to information about the negative aggregate judgment of other persons of high (low) status in comparison to when those judgments come from others of low (high) status. We use a quasi-experimental design to manipulate both exposure to others' evaluations and the status characteristics of those others.

1 Introduction

1.1 Theoretical Background

Taste expressions and consumption have long been shown to be a function of individuals’ desire to differentiate themselves from others (Simmel, 1957; DiMaggio, 1987). Bourdieu (1984) suggests that taste differences evoke a visceral reaction, and it is through the rejection of others’ aesthetic tastes that individuals engage in the symbolic boundary work that separates them. Thus, the expression of distastes serves to communicate social distance from certain social groups (Bryson, 1996; Lamont and Molnár, 2002). In contexts where individuals discover a similarity in taste with members of a meaningful outgroup, they frequently engage in “taste abandonment” to maintain social boundaries (Berger and Heath, 2008). Despite the robust literature on the relationship between cultural familiarity and network connections (Erickson, 1996; Lizardo, 2006), relatively little is known about the dynamic mechanisms underlying symbolic boundary work in real time.

How does access to the evaluations of others influence our own aesthetic evaluations of cultural objects? The simplest answer to this question sees the production of aesthetic judgments as a special case of judgment under uncertainty. This means that information about other people’s judgments becomes a key cog in a *generalized influence mechanism*. The basic idea behind generalized influence is simple: in any context where there is uncertainty about how to evaluate a cultural object, access to information about how others have evaluated it provides a heuristic for resolving that uncertainty (Strang and Macy, 2001). As such, according to this view, individuals will tend to adjust their own evaluations in a direction consistent with the valence of other people’s evaluations. For instance, discovering that a cultural object is broadly endorsed by others (regardless of those others’ characteristics) often increases its appeal (Positive Generalized Alignment), while discovering it is disliked decreases its appeal (Negative Generalized Alignment), as has been shown in previous work (Salganik et al., 2006).

However, cultural influence is rarely so monolithic; the status of the “others” matters. Exposure to the preferences of a high-status group might induce *conformity* (alignment), whereas exposure to a low-status or out-group might induce *reactance* (distancing). This approach is consistent with traditional models in the sociology of culture and consumption, which use a balance-like logic (Heider, 1946) to predict how information about other people’s choices (either explicit or implied) affects judgments of taste. These models suggest that access to positive evaluations from others will increase subjective evaluations of worth when those others are positively regarded, but not when they are negatively regarded. In the same way, access to other people’s negative evaluations will lead to a decrease in subjective evaluations of worth when those others are positively regarded (Lizardo and Skiles, 2016). Here influence operates as either conformity or “goodwill” (Bourdieu, 1984) when the presumed others are high-status, or as rejection or “symbolic boundary drawing” when others are low-status (Bryson, 1996).

A “Bourdieuian” (Bourdieu, 1984) version of this model adds the complicating factor of the focal person’s own status position, suggesting that goodwill and conformity to the opinions of presumed high-status others should be more prevalent among lower status respondents, while resistance and distancing from the judgments of presumed low-status others

should be more prevalent among higher status persons (Lizardo and Skiles, 2016).

While the Bourdieusian symbolic power model predicts goodwill (conformity on the part of lower-status persons to the judgments of high-status others) and equally possible pattern of results is that of generalized reactance, or the symbolic “war of all against all” where low-status others distance their judgments against those perceived to be higher status, while higher status people do the same against those of perceived lower status Tampubolon (2008).

In what follows, we put these considerations to the test by specifically examining how the respondent’s objective and subjective class status interact with the fictional group’s class status to produce theoretically distinct behaviors: *staying* (resisting influence), conforming (influence as a form of alignment), or reacting (influence as a form of distancing).

1.2 Hypotheses

Baseline Condition (Mere Exposure Effect):

Overall, we expect that mere exposure to the same cultural object across repeated trials results in an average improvement in evaluation, even in the absence of information about how other people evaluate the object (Reber et al., 1998). Access to information about other people’s evaluations can block or facilitate the mere exposure effect.

Generalized Influence in the Evaluation Only Condition:

- **Positive Evaluations:** Positive evaluations (liking) from generalized others should have an overall positive influence on the respondent’s second evaluation, facilitating the baseline drift.
- **Negative Evaluations:** Negative evaluations (disliking) from generalized others should have an overall negative influence on the respondent’s second evaluation, blocking the mere exposure effect.

When examining status plus evaluation conditions involving cross-status reactance and symbolic power, we advance three primary hypotheses. First, Hypothesis 1 (same-status alignment) posits that individuals will adjust their judgments to align with the aggregate judgment of their status peers. That is, both high- and low-status respondents will change their judgments in a positive direction when exposed to the positive judgments of others sharing their status, and in a negative direction when exposed to their negative judgments. Second, Hypothesis 2 (symbolic power and upward alignment) predicts that low-status respondents will align their judgments with the aggregate judgments of high-status others, adopting their positive evaluations and conforming to their negative evaluations. Finally, Hypothesis 3 (cross-status reactance) suggests that individuals will distance themselves from the aggregate judgments of out-group members, particularly low-status others. Thus, respondents will shift their judgments negatively when exposed to the positive judgment of an out-group, and positively when exposed to their negative judgment.

2 Data and Methods

2.1 Experimental Design and Data

We utilize data from the SSI-2012 survey. The core of the study is a within-subject experimental design where respondents rate an artwork on a 7-point scale (“taste1”), are exposed to a treatment condition, and then rate the artwork again (“taste2”).

All respondents viewed an image of the same painting: *Nocturne: Battersea Bridge*, by James McNeill Whistler (1872). Figure 1 displays the artwork used as the experimental stimulus. This specific painting was selected based on feedback gathered from an online pilot study conducted in June 2011, which tested images of eight paintings on a convenience sample of respondents. The Whistler painting was chosen for the main experiment because it elicited the highest variation in baseline opinions across all measured demographic characteristics, preventing floor or ceiling effects.



Figure 1: Experimental Stimulus: *Nocturne: Battersea Bridge*, by James McNeill Whistler (1872)

The experimental protocol proceeded as follows: Respondents first gave their initial impression of the painting on a 1–7 scale ranging from “like it very much” to “dislike it very much.” This expression serves as a baseline measure of taste, uninfluenced by external information. Next, respondents in the treatment groups were told that the group of people most likely to either like or dislike the painting possessed specific education and occupation characteristics.

The treatment conditions vary along three dimensions regarding this fictional alter group. The first dimension is taste evaluation, denoting whether the alter group liked or disliked the artwork. The second dimension is the status of the alter, contrasting high-status individuals (characterized by a college education and middle-class standing) with low-status individuals (working class and lacking a college degree). The final dimension is the control or no-status condition, which includes scenarios where only the alter’s taste is known without providing any class information, alongside a pure baseline control group.

2.2 Variables

The primary dependent variable in this study is the shift in aesthetic evaluation across the two trials. We analyze this outcome in several ways: as a continuous change score, through an ordinal structure using Cumulative Link Mixed Models (as detailed in the appendix), and as a nominal behavioral outcome categorized as staying, conforming, or reacting. The independent variables consist of the various experimental condition factors indicating the alter group’s status and evaluation. Furthermore, we include respondent status consistency as a moderating variable, captured by crossing working- or middle-class identification with the presence or absence of a college degree.

2.3 Analytic Strategy

Our data consists of repeated measures (before and after providing access to others’ evaluations) of the respondent’s aesthetic evaluation of a cultural object. Because repeated measures are correlated within subjects (but independent across subjects), we rely on random-effects models as a natural framework to simultaneously model within and between subjects variance. We use hierarchically nested versions of the mixed-effects model via maximum likelihood and use likelihood ratio tests to select the best fit.

To overcome the limitations of standard linear regression and fully test our hypotheses regarding alignment versus distancing, we additionally employ two advanced modeling techniques. First, we utilize multinomial logistic regression. While linear mixed models estimate the average magnitude and direction of shifts, they obscure the underlying distribution of behaviors. For example, an average shift of zero could mean complete rigidity, or it could mean polarized halves of conformity and reactance canceling each other out. To address this, we use a multinomial logistic regression to explicitly model the discrete probabilities of three theoretically distinct behaviors: maintaining one’s evaluation (staying), moving toward the group (conforming), or moving away from the group (reacting). Second, we employ inverse probability weighting (IPW) on all models to balance demographic covariates—including age, gender, race, and parental education—across the observational status groups. This ensures that the estimated moderating effects of class are not confounded by other demographic differences.

All results reported below utilize these IPW weights to ensure unbiased estimates of the status interactions.

3 Results

3.1 Covariate Balance

Before evaluating the hypotheses, we assessed the covariate balance across the observational status groups. Because respondents are not randomly assigned to their social class or educational attainment, these groups differ systematically on demographics (age, gender, race, parent’s education). We applied Propensity Score Weighting (IPW) to balance these covariates. As shown in Figure 2, weighting effectively eliminated demographic imbalances

across the status categories. All subsequent analyses utilize these weights to isolate the true moderating effect of social class and education.




Figure 2: Covariate Balance Before and After Weighting

3.2 Evaluating the Mere Exposure Effect

The baseline condition allows us to observe the counterfactual drift in aesthetic evaluations over repeated trials without external influence. Our analysis of the trial effects (Figure 3) provides mixed support for a universal mere exposure effect. While college-educated respondents (both Working and Middle class) exhibit a significant positive drift in evaluations across trials, respondents without a college degree exhibit no significant change. Thus, the mere exposure effect is not uniform but contingent on educational capital.

3.3 Evaluating Generalized Influence

In the “No Status” conditions, we hypothesized that positive evaluations from generalized others would facilitate the baseline drift, while negative evaluations would block it. The results partially support this: exposure to positive evaluations produces significant positive shifts, particularly for Middle Class individuals without a college degree. Conversely, exposure to generalized negative evaluations largely suppresses the positive baseline drift across all groups, resulting in null trial effects, exactly as predicted.




Figure 3: Estimated Trial Effect by Condition and Status. Error bars represent 95% confidence intervals.

3.4 Evaluating Status-Based Influence (Hypotheses 1–3)

When introducing the class status of the fictional alter group, we observe complex dynamics that challenge linear theories of social influence.

Hypothesis 1 (Same-status alignment) predicted that respondents would align their judgments with the positive and negative judgments of their own status peers. This finds *mixed support*. High-status respondents (Middle Class, College) show a significant positive shift when exposed to the “Like / High Status” condition, but no significant negative shift when exposed to the “Dislike / High Status” condition. In contrast, Low-status respondents (Working Class, No College) do not align in the “Like / Low Status” condition, and strikingly, they show a significant *positive* shift (distancing) when exposed to the “Dislike / Low Status” condition. Thus, low-status individuals actively react against the negative tastes of their own peers.

Hypothesis 2 (Symbolic power/upward alignment) predicted that Low-status respondents would align with High-status judgments. We find *partial support*. Low-status

respondents do exhibit a significant positive shift in the “Like / High Status” condition. However, they exhibit null shifts in the “Dislike / High Status” condition, rather than conformity.

Hypothesis 3 (Cross-status reactance) predicted distancing when exposed to out-group judgments. This is *not supported*. High-status respondents do not significantly distance (shift negatively) from the “Like / Low Status” condition, nor do they distance (shift positively) from the “Dislike / Low Status” condition. The strongest distancing behavior in the data actually occurs within-group (Low-status respondents reacting against the “Dislike / Low Status” condition), suggesting that reactance is driven more by avoiding low-status stigma than by strict cross-status boundaries.

3.5 Disaggregating Behavior: Conformity vs. Reactance

While the continuous change score models provide the average magnitude and direction of taste shifts, they obscure the underlying categorical behaviors. For instance, an average linear shift of zero could indicate that all respondents maintained their original evaluation, or it could indicate that half the respondents strongly conformed while the other half strongly reacted, canceling each other out in the aggregate. Furthermore, analyzing average shifts does not allow us to isolate the baseline propensity to simply maintain one’s taste (aesthetic rigidity) versus changing it.

To capture these distinct behaviors, we map the continuous shifts onto three discrete categories—Stay, Conform, and React—and fit a Multinomial Logistic Regression model.




Figure 4: Probability of Behavioral Response by Status Group

This behavioral disaggregation clarifies why the linear hypotheses (H1–H3) yielded mixed results: status inconsistency fundamentally alters the propensity to conform or react regard-

less of the specific alter. As Figure 4 reveals, Middle Class individuals without a college degree are highly prone to *conformity*, whereas Working Class individuals with a college degree display the highest baseline probability of *reactance*.

3.6 Sensitivity Analysis: Isolating Class Feedback

To ensure that the results are driven by the *class status* of the fictional alter groups, we run a sensitivity analysis dropping the “No Status” conditions.




Figure 5: Behavioral Response (Excluding ‘No Status’ Conditions)

The probability distributions of conformity versus reactance remain highly consistent with the main findings, confirming that the presence of class-based status information uniquely structures the evaluation shifts of respondents.

4 Discussion and Conclusion

Our analysis demonstrates that aesthetic judgments are deeply malleable, but not uniformly so, advancing the sociology of culture in several ways. Primarily, the findings highlight that directionality matters. By replacing a binary “stay/shift” conceptualization with a multinomial “stay/conform/react” framework, we reveal that cultural influence involves active processes of both emulation and boundary-drawing. Additionally, we identify status inconsistency as a critical driver of change. Respondents with inconsistent status markers, such as working-class individuals with a college degree or middle-class individuals without one, are the most aesthetically fluid. Crucially, they resolve this fluidity differently, leaning heavily into either reactance or conformity depending on their exact social location. Finally,

the methodological rigor introduced by incorporating propensity score weighting and validating with ordinal models establishes that these dynamics represent structural sociological realities, rather than statistical artifacts of ordinal scales or observational confounding.

Taste is indeed a social product, shaped continuously by class consciousness, educational capital, and the omnipresent evaluation of our peers.

References

- Berger, J. and Heath, C. (2008). Who drives divergence? identity signaling, outgroup dissimilarity, and the abandonment of cultural tastes. *Journal of Personality and Social Psychology*, 95(3):593–607.
- Bourdieu, P. (1984). *Distinction: A social critique of the judgement of taste*. Harvard university press.
- Bryson, B. (1996). “anything but heavy metal”: Symbolic exclusion and musical dislikes. *American Sociological Review*, 61(5):884–899.
- DiMaggio, P. (1987). Classification in art. *American Sociological Review*, pages 440–455.
- Erickson, B. H. (1996). Culture, class, and connections. *American Journal of Sociology*, 102(1):217–251.
- Heider, F. (1946). Attitudes and cognitive organization. *The Journal of psychology*, 21(1):107–112.
- Lamont, M. and Molnár, V. (2002). The study of boundaries in the social sciences. *Annual review of sociology*, 28(1):167–195.
- Lizardo, O. (2006). How cultural tastes shape personal networks. *American Sociological Review*, 71(5):778–807.
- Lizardo, O. and Skiles, S. (2016). Cultural objects as prisms: Perceived audience composition of musical genres as a resource for symbolic exclusion. *Socius*.
- Reber, R., Winkielman, P., and Schwarz, N. (1998). Effects of perceptual fluency on affective judgments. *Psychological science*, 9(1):45–48.
- Salganik, M. J., Dodds, P. S., and Watts, D. J. (2006). Experimental study of inequality and unpredictability in an artificial cultural market. *Science*, 311(5762):854–856.
- Simmel, G. (1957). Fashion. *American journal of sociology*, 62(6):541–558.
- Strang, D. and Macy, M. W. (2001). In search of excellence: Fads, success stories, and adaptive emulation. *American Journal of Sociology*, 107(1):147–182.
- Tampubolon, G. (2008). Revisiting omnivores in america circa 1990s: The exclusiveness of omnivores? *Poetics*, 36(2-3):243–264.

A Robustness Check: Cumulative Link Mixed Models (CLMM)

While we primarily report the results of linear mixed models in the main text to provide intuitive average effect sizes (analyzing taste shift as a continuous score), treating a 7-point Likert scale as continuous can result in biased standard errors. To ensure our findings are not statistical artifacts of this assumption, we employ a Cumulative Link Mixed Model (CLMM) as a rigorous robustness check.

By respecting ordinal scale boundaries, the CLMM ensures that respondents who already rated the painting a “7” or a “1” in Trial 1 are modeled correctly (i.e., they cannot shift higher or lower, respectively), preventing boundary-induced underestimation of treatment effects.

We re-estimated the trial interactions using the CLMM. The CLMM results mathematically confirm the findings of the linear models: the direction and significance of the condition-specific shifts remain consistent. This assures us that the observed behaviors—particularly the distancing effects among low-status respondents—are robust structural realities, not merely ceiling or floor artifacts caused by the bounded 1–7 survey scale.