

Stay, React, or Conform? Examining Social Influence Effects on Aesthetic Judgments

Omar Lizardo & Sara Skiles

Revised Manuscript — September 17, 2026

Abstract

Do others' aesthetic judgments causally alter our own, and how does this influence depend on the social distance between self and other? While cultural sociologists have long debated whether taste reflects emulation, distinction, or inertia, empirical research has relied almost exclusively on observational associations. We address this question using a national survey experiment ($N = 1,981$, adults aged ≥ 24) with a repeated-measures Difference-in-Differences design and inverse probability weighting. Manipulating exposure to aggregate peer evaluations (likes vs. dislikes) and alter status credentials across repeated trials, we find strong evidence of *valence asymmetry*: while re-evaluating artwork produces an intrinsic positive exposure drift in the absence of cues, negative social consensus completely suppresses this appreciation, functioning as an evaluative veto. In contrast, positive endorsements yield no increment beyond baseline exposure. Across educational strata, we find no evidence of upward cultural goodwill or same-status alignment; rather, status boundaries operate through reactive distancing, with postgraduate respondents abandoning appreciation when an artwork is endorsed by working-class peers. Finally, discrete behavioral modeling reveals that evaluative stability ("staying") is anchored by *structural status consistency*: consistent actors exhibit high inertia, whereas status-inconsistent actors display heightened conformity. Cultural influence thus functions via asymmetric vetoes and boundary policing conditioned by structural consistency.

1 Introduction

1.1 Personal Tastes versus Others' Tastes

In the sociology of taste, cultural tastes and practices have long been argued to be a function of individuals' socially induced motivation to differentiate themselves from others (Simmel, 1957; DiMaggio, 1987; Lizardo and Skiles, 2016). In the classic statement, Bourdieu (1984) suggests that differences in taste evoke a visceral reaction of affinity or disaffinity, and that in following their embodied *habitus* and choosing to engage with or avoid certain cultural objects, individuals simultaneously engage in symbolic boundary work and status distinction, because the objects they choose and reject are, in turn, chosen or rejected by others.

In this way, expressing *distastes* in particular serves to communicate social distance from others who choose the objects we reject (Bryson, 1996; Lamont and Molnár, 2002; Lizardo and Skiles, 2016). Accordingly, in contexts where individuals discover an unexpected similarity between their taste and that of members of an outgroup they wish to distance themselves from, they should abandon these tastes to maintain social boundaries (Berger and Heath, 2008; Simmel, 1957).

However, other people's tastes do not have to be solely a source of distinction and separation; they can also be a basis for building solidarity and discovering shared perspectives, especially among people considered similar in some socially relevant way (Pomès and Arsel, 2018). Via this process, expressing shared likes and dislikes can signal affiliation with others and serve as a mechanism for creating bounded groups (DiMaggio, 1987; Lizardo, 2006). Accordingly, contexts in which individuals find an unexpected *similarity* of tastes with others can foster community, solidarity, and productive exchange (Wohl, 2015; Vanzella-Yang and Abrutyn, 2022).

Note that in both ideal-typical stories, it is presumed that cultural choices and judgments of taste are shaped, directly or indirectly, by information about other people's tastes. In the first story, the tastes of others are a source of distinction and separation, while in the second story, they are a source of affiliation and solidarity. In both cases, knowing others' tastes is presumed to have a *causal* effect on our own tastes. Yet, despite the centrality of this assumption to sociological theories of taste, there is surprisingly little empirical evidence that others' tastes causally influence our own (but see Salganik et al., 2006). In this paper, we address this gap by experimentally testing for the causal effect (if any) of information about others' tastes on people's aesthetic evaluations.

Accordingly, in the research reported in this paper, we ask: *How does access to the evaluations influence our own aesthetic evaluations of cultural objects?*

1.2 How Do Others' Tastes Influence Our Own?

1.2.1 The Generalized Influence Mechanism

The simplest answer to the question we posed is to see the production of aesthetic judgments as a special case of producing a judgment under uncertainty (Kahneman, 2011). This means that information about other people's judgments becomes a key cog in a *generalized influence mechanism* in which some aggregation of others' tastes serves as a guidepost for our own (Salganik et al., 2006; Childress and Friedkin, 2012).

The basic idea behind generalized influence is simple: in any context where there is uncertainty about how to evaluate a cultural object, access to information about how others have evaluated it provides a heuristic for resolving that uncertainty, such that going along with the crowd is not necessarily a sub-optimal strategy (Strang and Macy, 2001; Mark, 2003). According to this view, individuals should tend to adjust their own evaluations in a direction consistent with the valence of others’ evaluations. For instance, discovering that a cultural object is broadly endorsed by others (regardless of those others’ characteristics) often increases its appeal (Positive Generalized Alignment), while discovering it is disliked decreases its appeal (Negative Generalized Alignment), as has been shown in previous work (Salganik et al., 2006).

However, cultural influence is rarely so monolithic; the status of the “others” matters. Exposure to the preferences of a high-status group might induce *conformity* (alignment), whereas exposure to the likes and dislikes of a low-status group might induce *reactance* (distancing). This approach is consistent with traditional models in the sociology of culture and consumption, which use a balance-like logic (Heider, 1946)—e.g., if a group you want to distance yourself from likes something, then you dislike it—to predict how information about other people’s choices (either explicit or implied) affects judgments of taste (Okada, 2017). These models suggest that access to positive evaluations from others will increase subjective evaluations of worth when those others are positively regarded, but not when they are negatively regarded. In the same way, access to other people’s negative evaluations will lead to a decrease in subjective evaluations of worth when those others are positively regarded (Lizardo and Skiles, 2016). Here influence operates as either conformity or “cultural goodwill” (Bourdieu, 1984) when the presumed others are high-status, or as rejection or “symbolic boundary drawing” when others are low-status (Bryson, 1996).

The “Bourdieuian” (Bourdieu, 1984) version of this model adds the complicating factor of the focal person’s own position in the social space, suggesting that goodwill and conformity to the opinions of presumed high-status others should be more prevalent among lower-status respondents, while resistance and distancing from the judgments of presumed low-status others should be more prevalent among higher-status persons (Lizardo and Skiles, 2016).

While the Bourdieusian symbolic power model predicts either goodwill (conformity on the part of lower-status persons to the judgments of high-status others), or asymmetric distancing (disalignment on the part of high-status people to the judgment of low-status others) an equally possible pattern of results is that of generalized reactance, or the symbolic “war of all against all” where low-status others distance their judgments against those perceived to be of higher status, while higher-status people do the same against those of perceived lower status (Tampubolon, 2008). The final logical possibility, intimated at by solidarity-based models (DiMaggio, 1987), is universal alignment with same-class others, so that our judgments will be swayed—in a negative or positive direction—by others perceived to be of the same status or social station, such that our judgments will become more negative when we learn that same-status others reject the object and more positive when we find that they embrace it.

1.3 Study Description

In what follows, we put these considerations to the test by specifically examining how the respondent’s objective and subjective class status interact with a fictional set of others’ class status to produce three theoretically distinct outcomes: *staying* (resisting influence), conforming (influence as a form of alignment or changing our judgment to be more similar to that of others), or reacting (influence as a form of distancing, or changing our judgments to be more dissimilar to that of others).

To that end, we designed a survey-experimental study in which participants were asked to report their feelings about a painting—Whistler’s *Nocturne*; see Figure 1—on a scale ranging from “dislike it very much” to “like it very much.” After we obtained this evaluation, in the experimental part of the survey, we told a subset of participants that our survey had so far shown that the group of participants that was most likely to either like (for one subset of participants) or dislike (for another subset) was either people with a PhD and working in higher education, people holding a bachelor’s degree in business and working in financial investment, or people with a high school education and working in food service. Yet a third group of participants was simply told that an unmarked set of “other” participants who had taken the survey had liked or disliked the painting. All participants then spent between 10 and 20 minutes answering various sociodemographic questions and completing survey modules. Participants were then shown the same painting one last time, prompted to take another look, and asked to provide their final take on it, using the same evaluative categories as at the beginning (a control group received no feedback and was asked to rate the painting twice).

1.4 Hypotheses

1.4.1 Mere Exposure and Positive Drift

For participants in the Control Condition who are simply asked to rate the painting twice without being exposed to the opinions of others, repeated exposure to a cultural object across trials is expected to produce an upward drift in evaluation even in the absence of external social information, a phenomenon in psychology known as the *mere exposure effect* (Zajonc, 1968; Reber et al., 1998).¹

Accordingly, our baseline hypothesis is that we should observe *positive drift* in the Control Condition, with participants asked to judge the same cultural object twice being more likely to produce a more positive judgment (on average) the second time.

- **Hypothesis 1 (Mere Exposure Effect):** *Participants who are asked to judge the painting twice, without being exposed to the opinions, will be more likely to give a more positive judgment the second time.*

Thus, if the mere exposure mechanism is at work, the Control Condition should show a positive drift in evaluation across trials, which we can use as a baseline to compare against the treatment conditions in which participants are exposed to the opinions.

¹For an application of the mere exposure effect to historical dynamics in the aesthetic evaluation of art see Cutting (2003).

1.4.2 Social Influence and Generalized Alignment

Generalized positive or negative influence processes should operate in tandem with mere exposure, either enhancing it (in the positive generalized influence condition) or suppressing it (in the negative generalized influence condition). In other words, we should expect that exposure to positive evaluations from others will enhance the positive drift in evaluation across trials, whereas exposure to negative evaluations from others will suppress it (Salganik et al., 2006).

- **Hypothesis 2a (Positive Evaluations):** *Endorsements (liking) from generalized others should exert an overall positive influence on the respondent's second evaluation, facilitating baseline drift.*
- **Hypothesis 2b (Negative Evaluations):** *Disapproval (disliking) from generalized others should exert an overall negative influence on the respondent's second evaluation, blocking the mere exposure effect.*

1.4.3 Social Influence and Status Specific Alignment

When it comes to hypotheses having to do with the interaction between a person's own status and the perceived status of others who either like or dislike the painting, we can consider three possibilities:

- **Hypothesis 3 (Same-Status Alignment):** *Individuals will tend to align their judgments with those of others who are perceived to be of similar status, such that they will tend to move towards more positive evaluations when provided information that similar-status others liked the cultural object, and towards negative evaluations when provided information that similar-status others disliked the painting.*
- **Hypothesis 4 (Cultural Goodwill):** *Lower-status individuals will tend to align their judgments with those of higher-status others, moving towards more positive evaluations when provided information that higher-status others liked the cultural object, and towards more negative evaluations when provided information that higher-status others disliked the cultural object.*
- **Hypothesis 5 (Status-Based Distancing Disalignment):** *Individuals will tend to distance themselves from the judgments of others who are perceived to be of different status, such that they will tend to move towards more negative evaluations when provided information that higher-status others liked the cultural object, and towards more positive evaluations when provided information that higher-status others disliked the painting.*

Of course it bears saying that some of these hypotheses do not imply mutual exclusivity. Individuals could certainly engage in both same-group taste alignment and outgroup taste distancing. In the same way, low-status individuals could engage in goodwill behaviors (aligning their tastes with those of their presumed status superiors) while simultaneously engaging in ingroup solidarity (aligning their tastes with those of same-group others). While

in the real world these causal effects are entangled in complex, sometimes even self-canceling or self-enhancing, causal interplay, the advantage of the experimental approach is that we can claim to separate them cleanly.

1.4.4 Status Inconsistency and Aesthetic Malleability

The previous hypotheses presumed that individuals occupy a single, consistent status position, either high or low. However, in reality, many individuals occupy a status position that is inconsistent across different dimensions of social stratification. For instance, a person may be working class but have a college degree, or middle class but without one. In the classic literature on the subject, this is usually referred to as status consistency and status crystallization (Hope, 1975).

The basic expectation stemming from this literature, along with more contemporary work on the social psychology of status (Correll and Ridgeway, 2003), is that individuals with consistent status profiles will be more likely to maintain their initial evaluation (“staying”) than those with inconsistent status profiles, who will be more likely to revise their initial evaluation in response to information about the aesthetic opinions of others. The reason for this is that status crystallization provides people with assurance as to their judgments (inuring them from influence) while individuals caught in contradictory slots in the social structure will be more likely to be uncertain about their own judgments and thus more susceptible to influence.

- **Hypothesis 6 (Status Consistency):** *Individuals with consistent status profiles (either high or low) will be more likely to maintain their initial evaluation (“staying”) than those with inconsistent status profiles, who will be more likely to revise their initial evaluation in response to information about the aesthetic opinions of others.*

2 Data and Methods

2.1 Sample Recruitment and Survey Administration

Data for this study were collected between February 28 and March 6, 2012, via an online experimental survey administered by Survey Sampling International (SSI), a leading research firm specializing in online panel recruitment and survey sampling. A total of 3,782 panel members accessed the survey instrument. Of these, 97 were excluded due to lack of informed consent ($N = 85$) or age under 18 ($N = 12$), 1,394 were screened out because demographic quota targets had already been met, and 16 dropped out during early screens, resulting in an analytic sample of $N = 2,275$ completed responses (with $N = 2,253$ providing complete initial and post-treatment evaluations). Participants received monetary compensation (\$3.92) via SSI, funded by an NSF Doctoral Dissertation Improvement Grant (Award SES-1203426) and the University of Notre Dame Institute for Scholarship in Liberal Arts.

To give respondents an opportunity to complete higher education and eliminate right-censoring in educational attainment among individuals still of standard undergraduate age (18–24), our primary analytic sample restricts the population to adults aged 24 or older (specifically, ≥ 25 given standard five-year Census survey brackets; $N = 1,981$, yielding 3,962

repeated observations across trials). In this mature adult population, educational attainment reflects completed institutional cultural capital, allowing for an unambiguous test of status-based influence without conflating non-completion with non-enrollment. As documented in Appendix G, re-estimating all models on the unrestricted full sample ($N = 2,253$) yields identical substantive findings and theoretical verdicts.

The sampling frame employed dynamic quota matching calibrated to U.S. Census benchmarks for gender, age, race/ethnicity, and geographic region to achieve a diverse and quasi-representative national sample of U.S. adults. To ensure adequate statistical power for examining status-differentiated cultural dynamics—particularly within high-cultural-capital occupational cells—the sample incorporated a deliberate oversample of college graduates (43.3% holding a bachelor’s degree or higher in the primary analytic sample versus 40.5% in the full sample and approximately 28% nationally in 2012). Table 1 presents the full demographic breakdown of the primary analytic sample, compared with U.S. population parameters, separated by experimental treatment condition.²

The randomization mechanism built into the *Qualtrics* survey engine ensures that internal validity is maximized: respondents were randomly assigned to treatment and control conditions, creating balanced experimental groups across all observed baseline characteristics (Table 1). At the same time, the demographic diversity across gender, age, race, and geographic region provides substantial ecological validity and heterogeneity compared to conventional undergraduate convenience samples.³

2.2 Experimental Design and Stimulus

All respondents viewed an image of the same artwork: *Nocturne: Blue and Gold – Old Battersea Bridge* by James McNeill Whistler (1872), displayed in Figure 1. The painting was selected based on an online pilot study ($N = 100$) that tested eight artworks; the Whistler painting was chosen because it produced the highest initial pre-treatment variance in aesthetic evaluations across all demographic subgroups.

The experiment used an incomplete factorial design with within-subject pre- and post-trials. In Trial 1 (pre-treatment baseline), respondents viewed the artwork accompanied by the prompt:

“Please look at this painting for a few seconds.”

“Give us your first impression of this painting – how much do you like it?”

Responses were recorded using a 7-point Likert scale (ranging from 1 = “I like it very much” to 7 = “I dislike it very much”, along with an “I can’t see the image” screening option; reverse-coded so that 7 = “like very much” and 1 = “dislike very much”). Next, during

²Because the survey used a quota-matched, opt-in online consumer panel rather than a strict probability-based sample with known inclusion probabilities (such as the GSS or ANES), the reported standard errors, p -values, and confidence intervals should be interpreted as reflecting *model-based and experimental design uncertainty within the sample*, rather than formal design-based population parameter estimates.

³Because cell sample sizes vary across experimental conditions and observational status consistency subgroups, we also report a comprehensive Minimum Detectable Effect (MDE) power analysis in Appendix A (Table A.1 and Figure A.1), explicitly delineating the statistical precision and empirical detection boundaries of the experimental design.

Table 1: Demographic Characteristics of Survey Sample (Age ≥ 24) Compared to U.S. Population Benchmarks

Demographic Variable	Sample N	Sample %	U.S. Pop %	Treatment	Control	Taste Only
Gender						
Women	1,036	52.3%	51.5%	51.9%	52.9%	54.0%
Men	945	47.7%	48.5%	48.1%	47.1%	46.0%
Race / Ethnicity						
White	1,284	64.8%	65.1%	65.0%	64.3%	64.3%
Hispanic / Latino	306	15.4%	15.8%	15.7%	14.6%	14.3%
Black / African American	250	12.6%	12.3%	12.5%	13.4%	13.0%
Asian / Pacific Islander	90	4.5%	4.5%	4.3%	5.7%	5.0%
Other / Multiracial	51	2.6%	2.3%	2.5%	1.9%	3.3%
Age Group						
25–34	407	20.5%	17.5%	20.7%	19.1%	20.7%
35–44	400	20.2%	17.5%	20.4%	15.3%	21.7%
45–54	437	22.1%	19.2%	22.0%	23.6%	21.3%
55–64	366	18.5%	15.6%	18.7%	17.2%	18.0%
65+	371	18.7%	17.2%	18.2%	24.8%	18.3%
Geographic Region						
Northeast	371	18.7%	17.9%	19.0%	17.8%	17.7%
South	689	34.8%	37.2%	35.6%	37.6%	29.3%
Midwest	445	22.5%	21.6%	21.5%	26.1%	25.7%
West	476	24.0%	23.3%	24.0%	18.5%	27.3%
Education (Oversampled)						
Less than High School	66	3.3%	13.3%	3.6%	1.9%	2.7%
High School Diploma	402	20.3%	30.4%	19.8%	22.3%	22.0%
Some College / AA	655	33.1%	28.5%	34.1%	30.6%	29.3%
Bachelor’s Degree	577	29.1%	18.1%	28.6%	29.9%	31.3%
Graduate / Professional	281	14.2%	9.6%	14.0%	15.3%	14.7%
Total N	1,981	100.0%		1,524	157	300

Note: Primary analytic sample is restricted to respondents aged 24 or older ($N = 1,981$), giving all participants an opportunity to complete college. U.S. population benchmarks are drawn from the 2010/2012 U.S. Census Bureau Current Population Survey (CPS). The analytic sample oversamples college-educated respondents to populate higher-status occupational cells. Full sample demographics ($N = 2,275$) including the 18–24 cohort are reported in Appendix Table A.5.

an intervening survey block, respondents completed demographic questions and a musical taste battery before treatment respondents read a brief vignette describing how a group of other people evaluated the artwork. Finally, in Trial 2 (post-treatment final impression), respondents were shown the painting again and asked:

“Please look at this painting for a few seconds again.”

“Give us your final impression of this painting – how much do you like it?”

on the identical rating scale. (The complete verbatim wording, layout, and sequencing across



Figure 1: Experimental Stimulus: *Nocturne: Blue and Gold – Old Battersea Bridge*, by James McNeill Whistler (1872)

all survey screens are transcribed in Appendix B.)

To prevent terminological confusion, we distinguish between the *pre-treatment baseline measurement* (Trial 1, which captures initial, uninfluenced aesthetic preference) and the *Control Condition* (the experimental group that receives no information about others between trials, allowing us to observe pure temporal re-evaluation in the absence of treatment). We refer to changes over repeated trials in this unexposed group as *exposure drift*.

2.3 Mitigation of Experimenter Demand Characteristics and Respondent Reactivity

A common methodological concern in within-subject experimental designs is whether respondents become aware that an evaluative shift is being measured, inducing experimenter demand compliance or artificial consistency. We implemented four procedural and design safeguards to mitigate demand awareness and assess respondent reactivity.

First, Trial 1 and Trial 2 were separated by an extensive cognitive and temporal buffer comprising 15 to 20 minutes of substantive survey batteries. Between aesthetic evaluations, respondents completed detailed modules on educational history, parental degree attainment, employment status, 22-category Standard Occupational Classification (SOC) coding, annual household income, and a comprehensive 20-genre musical taste battery evaluating preference ratings, listening frequency, and listener demographic stereotypes. This substantial

intervening distance prevented respondents from mechanically recalling or feeling compelled to repeat their Trial 1 ratings.

Second, the experimental feedback screens were framed naturalistically as observational descriptive summaries from prior survey waves rather than prescriptive social norms: respondents were told, “*Our survey so far indicates that the group of people most likely to say that they [LIKE / DISLIKE] the painting you saw have completed...*” Immediately following this prompt, respondents completed attribution and subjective similarity probes concerning the evaluating group, maintaining the perception that the survey was investigating social perception rather than testing their personal susceptibility to persuasion. When respondents subsequently reached the end of the survey and were re-shown the painting, they were simply asked for their “final impression” using the identical rating scale, without being reminded of their initial Trial 1 score and without any instructional prompt linking their second rating to the opinions of others or suggesting that conformity was expected.

Third, in the following analyses, all treatment shifts are compared against the unexposed Control Condition. Because control respondents completed the identical intervening survey battery without receiving any cue from others, comparing treatment conditions against this baseline cleanly separates genuine informational influence from any residual re-test reactivity or temporal fatigue.

Finally, the empirical distribution of discrete behavioral choices provides strong evidence against demand compliance. Over 74% of respondents across conditions exhibited complete evaluative stability (“staying”), and a substantial proportion displayed active oppositional reactance by shifting in the direction opposite to the cue. Had respondents been compliance-driven, conformity rates would have dominated across conditions, rather than emerging selectively in status-inconsistent configurations.

2.4 Variables and Measurement

Aesthetic evaluation was measured on a reversed 1–7 scale. In continuous models, the outcome is the repeated evaluation across trials. In behavioral models, change between trials is categorized into three mutually exclusive choices: *Stay* ($\Delta = 0$), *Conform* (moving in the direction of others’ evaluations), and *React* (moving in the opposite direction).

The treatment conditions crossed others’ evaluation (Like vs. Dislike) with others’ social status. High-status others were differentiated into Cultural Specialists (possessing multiple advanced degrees including a PhD and working in higher education) and Economic Specialists (holding a bachelor’s degree in business and working in financial investment). Low-status others were described as possessing a high school education and working in food service. In addition, two Taste-Only conditions provided other like-or-dislike evaluations without any status indicators, and a pure Control Condition provided no information about others between trials ($N = 157$ in the primary sample; $N = 182$ in the full sample).

We collapse the Economic Specialist and Cultural Specialist conditions into a single High-Status condition ($+Status$) for the primary analysis. This consolidation is justified on both statistical and substantive grounds. First, crossing respondents’ own status consistency quadrants by all nine experimental conditions would expand the design to 36 cells, reducing cell sizes in smaller subgroups to $N \approx 20$ –35 (and $N = 21$ in control), severely inflating the Minimum Detectable Effect threshold ($d > 0.45$). Second, formal linear hypothesis tests

confirm that responses to Economic Specialists and Cultural Specialists are statistically indistinguishable for both positive cues ($F = 0.29, p = 0.591$ in the pooled sample; $F = 0.92, p = 0.337$ among college graduates) and negative cues ($F = 0.06, p = 0.803$ in the pooled sample; $F = 0.04, p = 0.846$ among college graduates). Collapsing them into a single high-status category effectively doubles the statistical power for high-status cues ($N \approx 505$ per arm in the primary sample and $N \approx 575$ in the full sample) without obscuring any empirical heterogeneity.

To capture respondents’ structural social location, we examine objective educational attainment as the primary status dimension governing cultural capital and aesthetic distinction. Because our experimental treatment vignettes explicitly differentiated alter status through educational credentials and occupations (doctoral cultural specialists, business baccalaureates, and high school food service workers), stratifying respondents by their own educational attainment provides direct construct congruence between ego and alter. We operationalize educational attainment into four granular tiers: (1) High School or Less ($N = 468$), (2) Two-Year Associate Degree or Some College ($N = 655$), (3) Four-Year Bachelor’s Degree ($N = 577$), and (4) Postgraduate or Advanced Degree ($N = 281$). This four-tier specification serves as our primary analytical framework for testing Hypotheses 3–5 in Section 3.2 (Figure 4 and Appendix Table A.2).

Next, to test Hypothesis 6 concerning status consistency and crystallization (Section 3.3), we introduce respondents’ subjective social class self-identification (Middle Class vs. Working Class) and cross it with educational attainment to form a four-quadrant typology of status consistency: Working Class without College, Working Class with College, Middle Class without College, and Middle Class with College. In the base operationalization, “Working Class” includes a minority of respondents who self-identified as “Lower Class” ($N = 220$, 11.1%) and “Middle Class” includes a small proportion of respondents who self-identified as “Upper Class” ($N = 47$, 2.4%). To verify that our substantive conclusions are not sensitive to boundaries of subjective class, we fit models using an alternative categorization that restricts the sample only to self-identified working- and middle-class respondents (excluding lower- and upper-class extremes) as an additional sensitivity analysis in Appendix F.

2.5 Inverse Probability Weighting

A key methodological challenge in experimental design examining status-contingent influences stems from the hybrid nature of the design. While the experimental cues (others’ evaluation valence and others’ status) were randomly assigned, respondents’ own social locations—captured by educational attainment and subjective class identification—represent pre-existing, observational attributes (Cruces et al., 2013). Consequently, individuals do not have an equal probability of belonging to each status group: for example, respondents holding graduate degrees or middle-class identities are systematically older, more likely to have college-educated parents, and differ in racial/ethnic composition compared to non-college working-class individuals. If unaddressed, comparisons across status groups (and their interactions with the randomized treatments) risk confounding the effect of social status with these correlated background characteristics.

To solve this problem and strengthen causal inference without inducing overcontrol bias (avoiding conditioning on post-treatment socioeconomic outcomes of education, such as

personal income or adult migration), we implement Inverse Probability Weighting (IPW) using strictly exogenous pre-treatment social background characteristics (Rosenbaum and Rubin, 1983). The logic of IPW proceeds in two steps. First, for each respondent, we estimate a propensity model predicting status group membership conditional on exogenous baseline covariates $\mathbf{X}_i$ (age, gender, race/ethnicity, and parental educational capital): $e_s(\mathbf{X}_i) = P(S_i = s \mid \mathbf{X}_i)$. We estimate generalized multinomial propensity scores both across the four educational tiers (for the Difference-in-Differences models in Section 3.2) and across the four status consistency quadrants (for the discrete behavioral choice models in Section 3.3). Second, each observation is weighted by the inverse of its estimated propensity score to target the Average Treatment Effect (ATE) across the sample: $w_i = 1/e_s(\mathbf{X}_i) = 1/P(S_i = s \mid \mathbf{X}_i)$.

Intuitively, IPW creates a synthetic “pseudo-population” by up-weighting underrepresented demographic profiles within each status group (e.g., working-class respondents whose parents completed college) and down-weighting overrepresented profiles. This weighting scheme breaks the empirical association between the baseline covariates $\mathbf{X}_i$ and status group membership. Under the assumption of conditional exchangeability (no unmeasured confounders given $\mathbf{X}_i$) and positivity, IPW ensures that status groups are demographically comparable. As a result, subsequent differences in evaluative shifts, behavioral inertia, or reactance across status categories can be interpreted as the direct consequence of structural status location and status consistency, rather than demographic artifacts.

Figure 2 displays the standardized mean differences (SMD) for all baseline covariates across the status categories before and after Inverse Probability Weighting (IPW), generated using `cobalt` (Greifer, 2022). Prior to weighting, status groups exhibited substantial demographic disparities: for instance, college-educated middle-class respondents were older and had higher rates of parental college attainment (raw SMDs ranging from 0.35 to 0.50). Following IPW estimation, the weighted pseudo-population achieved balance, with all absolute standardized mean differences compressed well below the conventional conservative threshold of 0.10 standard deviations. This confirms that subsequent model estimates isolate the effect of social class and status consistency, net of observed demographic sorting.

2.6 Difference-in-Differences Analytic Strategy

To strictly isolate information-induced shifts from exposure drift, we formulate our analysis within a Difference-in-Differences (DiD) framework. Rather than examining simple pre-to-post changes within each condition in isolation—which conflates the effect of information from others with the natural exposure drift that occurs upon viewing an artwork a second time—the DiD strategy compares every treatment condition against the influence-unexposed Control Condition.

Accordingly, we specify a linear mixed-effects models with random intercepts of the general form:

$$\text{Taste}_{it} = \beta_0 + \beta_1 \text{Trial}_{it} + \sum_{k=2}^K \beta_k \text{Cond}_{ik} + \sum_{k=2}^K \delta_k (\text{Trial}_{it} \times \text{Cond}_{ik}) + u_i + \epsilon_{it} \quad (1)$$

where i indexes individual respondents ($i = 1, \dots, N$), $t \in \{1, 2\}$ indexes measurement trials (Trial 1 pre-treatment baseline vs. Trial 2 post-treatment evaluation), $k \in \{1, \dots, K\}$ indexes




Figure 2: Covariate Balance Across Observational Status Groups Before and After IPW Weighting (Balancing Age, Gender, Race/Ethnicity, and Parental Education). The dashed lines indicate the standard ± 0.10 threshold for acceptable balance.

the experimental conditions with the unexposed Control Condition ($k = 1$) serving as the omitted reference category, $u_i \sim \mathcal{N}(0, \sigma_u^2)$ is the individual-specific random intercept, and $\epsilon_{it} \sim \mathcal{N}(0, \sigma_\epsilon^2)$ is the trial-level residual error.

In the Equation 1 specification, the intercept (β_0) represents the expected baseline aesthetic evaluation at Trial 1 (pre-treatment) for respondents assigned to the unexposed Control Condition: $\beta_0 = E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]$. The Trial main effect (β_1) captures the counterfactual *exposure drift* from Trial 1 to Trial 2 in the absence of external information from others: $\beta_1 = E[\text{Taste} \mid \text{Trial} = 2, \text{Control}] - E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]$. When repeated viewing induces an intrinsic appreciation of the artwork (the mere exposure effect), β_1 should be positive and statistically significant.

The condition main effects (β_k) reflect pre-treatment differences between condition k and the Control Condition at Trial 1: $\beta_k = E[\text{Taste} \mid \text{Trial} = 1, \text{Cond} = k] - E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]$. Because respondents were randomly assigned to experimental conditions before viewing the painting, these coefficients assess baseline balance and are expected to be close to zero. The interaction coefficients (δ_k) represent the *causal Difference-in-Differences treatment effects*. Each δ_k measures whether the change in evaluations from Trial 1 to Trial 2 in treatment condition k differed significantly from the change observed in the unexposed Control Condition:

$$\begin{aligned} \delta_k = & (E[\text{Taste} \mid \text{Trial} = 2, k] - E[\text{Taste} \mid \text{Trial} = 1, k]) \\ & - (E[\text{Taste} \mid \text{Trial} = 2, \text{Control}] - E[\text{Taste} \mid \text{Trial} = 1, \text{Control}]) \end{aligned} \quad (2)$$

Substantively, δ_k answers the counterfactual question: *Did exposure to evaluation k cause respondents' aesthetic judgments to change significantly more or less than they would have changed under mere repeated exposure alone?* A positive estimate ($\delta_k > 0$) indicates that feedback from others facilitated aesthetic appreciation beyond natural exposure drift, whereas a negative estimate ($\delta_k < 0$) indicates that feedback from others suppressed or reversed the positive exposure drift that would have naturally occurred.

Estimating the simple pooled DiD model in Equation 1 without differentiating by social location yields mostly muted or null effects (reported in Appendix Table A.2 and Appendix Figure A.2). As our introductory considerations suggest, however, this pooled invariance is an artifact of aggregation, hiding underlying heterogeneity in the causal effect (Brand and Thomas, 2013). In essence, because aesthetic influence operates relationally based on the status congruence between ego and alter, pooling across socially dissimilar respondents washes out opposing directional adjustments. The substantive dynamics of social influence emerge when we examine models interacting experimental conditions with education, subjective social class, and status consistency.

Instead, to test hypotheses of how social influence is moderated by respondents' structural social location (educational attainment, subjective social class, and status consistency), we extend Equation 1 to a fully moderated Difference-in-Differences specification:

$$\begin{aligned}
\text{Taste}_{it} = & \beta_0 + \beta_1 \text{Trial}_{it} + \sum_{k=2}^K \beta_k \text{Cond}_{ik} + \sum_{m=2}^M \gamma_m \text{Status}_{im} \\
& + \sum_{k=2}^K \delta_k (\text{Trial}_{it} \times \text{Cond}_{ik}) + \sum_{m=2}^M \theta_m (\text{Trial}_{it} \times \text{Status}_{im}) \\
& + \sum_{k=2}^K \sum_{m=2}^M \lambda_{km} (\text{Cond}_{ik} \times \text{Status}_{im}) \\
& + \sum_{k=2}^K \sum_{m=2}^M \psi_{km} (\text{Trial}_{it} \times \text{Cond}_{ik} \times \text{Status}_{im}) + u_i + \epsilon_{it}
\end{aligned} \tag{3}$$

Where Status_{im} represents indicators for respondent i 's structural social location across M categories ($m = 1, \dots, M$, with $m = 1$ denoting the reference group). In this moderated framework, θ_m captures group-specific differences in counterfactual baseline exposure drift in the Control Condition ($\text{Trial}_{it} \times \text{Status}_{im}$), δ_k represents the DiD treatment effect for the reference status group, and ψ_{km} represents the three-way interaction coefficient ($\text{Trial}_{it} \times \text{Cond}_{ik} \times \text{Status}_{im}$) testing whether the causal DiD treatment effect of condition k diverges significantly across status groups. The status-specific causal DiD treatment effect for subgroup m is given by the linear combination $\delta_{km} = \delta_k + \psi_{km}$.⁴

⁴By the Frisch-Waugh-Lovell theorem, estimating Equation 1 (or Equation 3) via linear mixed models with individual random intercepts (u_i) is mathematically equivalent to regressing individual change scores ($\Delta \text{Taste}_i = \text{Taste}_{i2} - \text{Taste}_{i1}$) on the treatment and status interaction terms with the Control Condition as the omitted reference group:

$$\Delta \text{Taste}_i = \beta_1 + \sum_{m=2}^M \theta_m \text{Status}_{im} + \sum_{k=2}^K \delta_k \text{Cond}_{ik} + \sum_{k=2}^K \sum_{m=2}^M \psi_{km} (\text{Cond}_{ik} \times \text{Status}_{im}) + \epsilon_i \tag{4}$$

3 Results

3.1 Baseline Exposure Drift and Generalized Social Influence

We begin by establishing a status-free empirical baseline: testing simultaneously for any mere-exposure drift observed under repeated evaluation in the unexposed Control Condition and examining whether learning that unknown others liked or disliked the artwork shifts aggregate aesthetic evaluations via a generalized influence mechanism. To do so, we estimate a repeated-measures linear mixed model on the subset of respondents assigned to the Control Condition ($N = 157$), the Generalized Like condition ($N = 152$), and the Generalized Dislike condition ($N = 148$), totaling $N = 457$ respondents (914 observations across trials) with Inverse Probability Weighting (IPW):

$$\begin{aligned} \text{Taste}_{it} = & \beta_0 + \beta_1 \text{Trial}_{it} + \beta_2 \text{Like}_i + \beta_3 \text{Dislike}_i \\ & + \delta_{\text{Like}}(\text{Trial}_{it} \times \text{Like}_i) + \delta_{\text{Dislike}}(\text{Trial}_{it} \times \text{Dislike}_i) + u_i + \epsilon_{it} \end{aligned} \quad (5)$$

where the unexposed Control Condition serves as the omitted reference group. Figure 3 visualizes these baseline dynamics, reporting the net Difference-in-Differences treatment effects for generalized social influence relative to the counterfactual exposure drift observed in the unexposed Control Condition.




Figure 3: Difference-in-Differences Treatment Effects for Generalized Social Influence (Status-Free Primary Sample, $N = 457$). The figure displays net Difference-in-Differences treatment effects (δ_k) relative to the unexposed Control Condition exposure drift (dashed vertical line at zero), with 95% confidence intervals. Colors denote statistical significance: red ($p < 0.05$) and grey ($p \geq 0.10$).

This status-free baseline model establishes two main findings. First, consistent with *Hypothesis 1 (Mere Exposure Effect)*, in the complete absence of external social cues (the

where the intercept and status main effects estimate the control exposure drift within each status group, and the interaction terms estimate the status-specific DiD treatment effects.

Control Condition), repeated exposure generates a statistically significant upward appreciation of the artwork ($\beta_1 = +0.156, SE = 0.055, t = 2.86, p = 0.0044$), confirming that mere exposure alone produces positive drift across trials. Second, evaluating generalized social influence reveals a pronounced *valence asymmetry* between positive and negative consensus. When respondents are informed that unknown others liked the painting, raw evaluations increase from 4.65 to 4.85 ($+0.203, p = 0.009$); however, relative to the natural counterfactual drift observed in the control group, this additional increment is positive but not statistically distinguishable from mere exposure drift ($\delta_{\text{Like}} = +0.047, SE = 0.080, t = 0.59, p = 0.555$). Consequently, *Hypothesis 2a (Positive Evaluations)* is not supported, as generalized positive cues fail to significantly facilitate evaluative drift beyond the baseline exposure effect.

In contrast, learning that others disliked the painting completely overrides the mere exposure effect, reducing second-trial evaluations to a flat -0.030 ($p = 0.68$). Relative to the control baseline drift, this constitutes a statistically significant negative Difference-in-Differences suppression ($\delta_{\text{Dislike}} = -0.186, SE = 0.078, t = -2.40, p = 0.0169$), providing strong empirical support for *Hypothesis 2b (Negative Evaluations)*. Directly contrasting positive and negative cues reveals a significant directional swing of nearly a quarter point on the 7-point scale ($\Delta = +0.233, SE = 0.080, \chi^2 = 8.51, p = 0.0035$), demonstrating that while generalized endorsements fail to reliably boost appreciation over mere exposure, generalized disliking exerts a discernible suppressing effect.

While this status-free model demonstrates clear valence asymmetry favoring other people’s *negative* evaluations as the strongest signal exerting a causal effect on aesthetic judgments, evaluating social influence solely at the aggregate level conceals substantial causal heterogeneity based on social location (Brand and Thomas, 2013).⁵ We therefore turn to our status-differentiated models to examine how social influence is conditioned by educational cultural capital (Section 3.2) and structural status consistency (Section 3.3).

3.2 Status-Differentiated Social Influence Across Educational Tiers

Because our experimental treatment vignettes explicitly differentiated alter status through educational credentials and professional occupations (doctoral cultural specialists, business baccalaureates, and high school food service workers), respondents’ own educational attainment represents the primary structural axis governing cultural capital, symbolic boundaries, and aesthetic judgment, providing direct construct congruence between ego and alter.

Given our achieved analytic sample ($N = 1,981$ individuals; $3,962$ repeated observations) and the strong statistical precision of our repeated-measures Difference-in-Differences design ($r = 0.899, \sigma_\Delta = 0.688$), we disaggregate educational attainment into four granular tiers: (1) High School or Less ($N = 468$), (2) Two-Year Associate Degree or Some College ($N = 655$), (3) Four-Year Bachelor’s Degree ($N = 577$), and (4) Postgraduate or Advanced Degree ($N = 281$). We estimate Difference-in-Differences models via generalized multinomial Inverse Probability Weighting (IPW) balancing baseline age, gender, race/ethnicity, and parental education across all four educational tiers. Figure 4 visualizes the net DiD treatment effects relative to the counterfactual exposure drift in the Control Condition within each tier (full tabular regression estimates, standard errors, and sample sizes are reported in Appendix

⁵As shown in the full pooled models in Appendix Table A.4 and Appendix Figure A.2.

Table A.2).




Figure 4: Estimated Difference-in-Differences Treatment Effects Across Four Educational Tiers (Multinomial IPW-Weighted, Age ≥ 24 Primary Sample). Horizontal bars indicate point estimates with 95% confidence intervals relative to the unexposed Control Condition within each tier (dashed line at zero). Colors denote statistical significance: red ($p < 0.05$), orange ($p < 0.10$), and grey ($p \geq 0.10$). Complete regression coefficients are reported in Appendix Table A.2.

Evaluating our status-differentiated hypotheses using the DiD framework across these four institutionalized cultural capital tiers reveals a monotonic gradient of causal heterogeneity. The simplest interpretation is that *the overall sensitivity to negative evaluations from anonymous others is a trait of highly educated groups*; in essence, people with higher education tend to be particularly sensitive to information that others did not like the painting, even when these others are just an anonymous mass ($\text{DiD} = -0.442$, $SE = 0.151$, $p = 0.0034$).

This sensitivity to negative feedback among the highly educated, however, extends to status-differentiated information, supporting the negative version of the same-status alignment idea (*Hypothesis 3*) for high-status respondents. Accordingly, when bachelor's degree holders learn that high-status people dislike the painting, they downgrade their evaluation of it, and this effect is even stronger among respondents with a post-graduate degree

(DiD = -0.266 , $SE = 0.122$, $p = 0.0300$). Puzzlingly, the aversion of high-status respondents to negative feedback even extends to news that *low-status* others dislike it, with people with a post-graduate degree decreasing their liking for the painting after hearing this news (DiD = -0.345 , $SE = 0.123$, $p = 0.0049$), a result not predicted by any theory.

Beyond the one shining spot mentioned, the Same-Status Alignment (*Hypothesis 3*) had a rough time with these data and the study design, receiving little to no support, especially when providing people with information about others' *positive* aesthetic judgments (likes). Hypothesis 3 predicted that individuals would align their judgments with peers sharing their educational status. However, high-status respondents with a Bachelor's degree show no positive alignment when informed that high-status people liked the artwork (DiD = -0.030 , $SE = 0.123$, $p = 0.809$). Among postgraduates, exposure to high-status likes is associated with a marginally negative shift (DiD = -0.235 , $SE = 0.120$, $p = 0.0504$).

At the other end of the spectrum, respondents with a high school education or less also fail to exhibit same-status alignment, displaying null shifts when told that low-status peers liked the artwork (DiD = $+0.051$, $SE = 0.130$, $p = 0.697$) and positive oppositional shifts when told that low-status peers disliked it (DiD = $+0.240$, $SE = 0.134$, $p = 0.0737$). Because same-status positive cues fail to foster appreciation across all tiers, the data do not support Hypothesis 3.

A similar verdict applies to the Cultural Goodwill idea (*Hypothesis 4*), which is also decisively rejected. Hypothesis 4 predicted that lower-status individuals would defer to elite aesthetic authority, adopting the likes and conforming to the dislikes of high-status others. Among respondents with a high school education or less, neither exposure to High-Status Likes (DiD = $+0.032$, $SE = 0.113$, $p = 0.777$) nor exposure to High-Status Dislikes (DiD = -0.140 , $SE = 0.113$, $p = 0.215$) produces a statistically significant change. In this study, lower-educated respondents are completely unresponsive to feedback about elite tastes, displaying neither upward emulation (Simmel, 1957) nor aesthetic deference (Bourdieu, 1984). Instead, their evaluative responsiveness is selectively populist: they shift positively (and marginally) only when informed of broad popular majorities (Generalized Like: DiD = $+0.284$, $SE = 0.151$, $p = 0.0608$) or in opposition to low-status peer disliking.

Third, the Bourdieusian version of the Status-Based Distancing Disalignment idea (*Hypothesis 5*) is strongly supported and emerges as the defining behavioral marker of elite cultural capital. When respondents with a postgraduate degree learn that working-class others liked the Whistler painting, their evaluations plunge by nearly half a point on the 7-point scale relative to the control baseline (DiD = -0.491 , $SE = 0.135$, $t = -3.64$, $p = 0.0003$, Cohen's $d = 0.33$), representing the single largest negative treatment effect in the entire experiment. Because postgraduates naturally drift upward by $+0.375$ ($p = 0.006$) under mere exposure, observing that lower-status others appreciated the artwork more than they did reverses their judgment, providing experimental evidence of Bourdieusian symbolic distinction and Simmelian taste abandonment.

Finally, the intermediate tier (Two-Year Associate Degree or Some College, $N = 655$) displays complete evaluative inertia across all conditions ($p \geq 0.25$, with point estimates tightly bounded between -0.113 and $+0.145$). Ancillary tests confirm that terminal Associate degree holders ($N = 196$) and non-degreed college attenders ($N = 459$) are equally inert. Bachelor's degree holders (Tier 3) represent the structural threshold where downward sensitivity begins: they adjust downward under High-Status Dislike (DiD = -0.271 , $SE =$

0.125, $p = 0.0296$) and Generalized Dislike ($\text{DiD} = -0.308$, $SE = 0.154$, $p = 0.0455$), while ignoring all positive endorsements ($p \geq 0.75$).

In short, disaggregating education reveals that sub-baccalaureate credentials foster positive facilitation or inertia, whereas baccalaureate and post-baccalaureate cultural capital turn aesthetic judgment into an asymmetric filter sensitive almost exclusively to social disapproval and outgroup contamination.

3.3 Status Inconsistency and Discrete Behavioral Choices

While continuous Difference-in-Differences models identify average shifts across conditions, they mask non-linear behavioral choices where opposing individual movements cancel each other out in the aggregate. Furthermore, examining education alone presumes that individuals occupy an unambiguous status position. In reality, many individuals experience status inconsistency across structural dimensions—such as possessing a college degree while identifying as working-class (Friedman et al., 2021), or claiming a middle-class identity without a college credential. In the classic sociological literature, status consistency and status crystallization (Hope, 1975; Lenski, 1954) govern normative certainty and evaluative firmness.

To test *Hypothesis 6*, we introduce respondents’ subjective social class self-identification (Middle Class vs. Working Class) and cross it with educational attainment (College Degree vs. No College) to construct a four-quadrant typology of status consistency: (1) Working Class without College ($N = 680$, low-consistent), (2) Working Class with College ($N = 243$, inconsistent), (3) Middle Class without College ($N = 443$, inconsistent), and (4) Middle Class with College ($N = 615$, high-consistent).⁶ We categorize responses into three mutually exclusive behaviors: *Stay* ($\Delta = 0$, maintaining one’s initial evaluation), *Conform* (moving in the direction of others’ evaluation), and *React* (moving in the opposite direction of others’ evaluation) and estimate an IPW-weighted multinomial logistic regression across the four status consistency groups.

Figure 5 presents the Average Marginal Effects (AME) on the probability of each behavioral choice from an IPW-weighted multinomial logistic regression, compared against the reference group of Working Class respondents without a college degree. In the baseline reference group, the predicted probabilities are 73.8% for staying, 14.2% for conforming, and 12.0% for reacting.

The marginal effects provide empirical support for *Hypothesis 6* and offer new insights into the social dynamics of status inconsistency. First, in terms of evaluative stability (Panel A: Stay), status consistency serves as a source of aesthetic confidence and thus inertia in aesthetic judgment when faced with feedback on others’ tastes. Self-identified middle-class individuals with a college degree are significantly more likely to maintain their initial judgment than non-college-educated working-class folk ($\text{AME} = +0.073$, $SE = 0.013$, $z = 5.47$, $p < 0.001$), with an overall stay rate of 81.1%. In contrast, status-inconsistent individuals without a college degree exhibit lower stability ($\text{AME} = -0.026$, $SE = 0.014$, $z = -1.81$, $p = 0.070$), with their stay rate falling to 71.2%. College-educated self-identified working-class people, on the other hand, are not statistically distinguishable from non-college-educated working-class

⁶Complete 24-cell continuous DiD regression estimates across all experimental conditions and status consistency quadrants are reported in Appendix Table A.3 for reference.




Figure 5: Average Marginal Effects on Behavioral Choices (Stay, Conform, React) Across Status-Consistency Groups (Age ≥ 24 Primary Sample). Horizontal bars indicate differences in choice probability relative to the baseline reference category (Working Class, No College), with 95% confidence intervals. Colors indicate statistical significance: red ($p < 0.05$), orange ($p < 0.10$), and grey ($p \geq 0.10$).

people (AME = -0.012 , $p = 0.379$, stay rate 72.5%).

Second, in terms of conformity (Panel B), status inconsistency emerges as the primary driver of susceptibility to external social influence. Both status-inconsistent groups display large, highly significant increases in the probability of conforming to others' judgments. Working Class individuals with a college degree are +4.2% more likely to conform relative to the reference group (AME = $+0.042$, $SE = 0.012$, $z = 3.60$, $p = 0.0003$, overall conformity rate 18.4%), and Middle Class individuals without college are +4.3% more likely to conform (AME = $+0.043$, $SE = 0.012$, $z = 3.69$, $p = 0.0002$, conformity rate 18.5%). In stark contrast, high-status consistent individuals are significantly less likely to conform (AME = -0.033 , $SE = 0.011$, $z = -3.08$, $p = 0.0021$, conformity rate 10.9%).

Finally, regarding oppositional reactance (Panel C), the results show that non-college-educated working-class individuals are *more* likely to react *against* others' tastes than individuals in the other three status-consistency categories. Both college-educated groups are significantly less likely to move away from others' judgments compared to non-college working-class people: Working Class respondents with a college degree are -3.0% less likely to react (AME = -0.030 , $SE = 0.010$, $z = -3.06$, $p = 0.0022$), and Middle Class college graduates are -4.0% less likely to react (AME = -0.040 , $SE = 0.010$, $z = -4.20$, $p < 0.0001$). Meanwhile, middle-class individuals without a college degree exhibit a marginally lower reactance rate (AME = -0.018 , $p = 0.08$).

To confirm that these behavioral dynamics are driven specifically by status-laden social cues rather than by generic feedback about others' opinions, we re-estimated the multinomial model excluding the two Taste-Only conditions ($N = 1,643$). As displayed in Appendix Figure A.3, the results are virtually identical: status consistency continues to anchor evaluative inertia (Middle Class with College: $+7.6\%$, $p < 0.001$), status inconsistency heightens

conformity (Working Class with College: +4.9%, $p < 0.001$; Middle Class without College: +4.1%, $p < 0.001$), and oppositional reactance remains concentrated among non-college working-class people (−3.1% and −4.1% for college-educated groups, $p < 0.001$).

4 Discussion

4.1 Summary of Key Results

In this study, we investigated how access to others’ aesthetic evaluations causally influences individual taste judgments in real time. Using a national survey experiment with a Difference-in-Differences design and inverse probability weighting, we resolved several long-standing ambiguities present in observational literature. Table 2 provides a systematic overview of our focal hypotheses, the underlying theoretical expectations, and the resulting empirical verdicts.

Table 2: Summary of Hypotheses, Theoretical Expectations, and Empirical Verdicts

Hypothesis / Mechanism	Theoretical Expectation	Empirical Verdict
<i>Hypothesis 1: Mere Exposure Effect</i>	Upward drift in aesthetic evaluation across repeated trials in the absence of external social cues	Supported ($p < 0.01$)
<i>Hypothesis 2a: Generalized Influence (Positive Evaluations)</i>	Endorsements from generalized others facilitate and enhance baseline exposure drift	Not Supported ($p \geq 0.10$)
<i>Hypothesis 2b: Generalized Influence (Negative Evaluations)</i>	Disapproval from generalized others suppresses and blocks baseline exposure drift	Supported ($p < 0.05$)
<i>Hypothesis 3: Same-Status Alignment</i>	Individuals align judgments with aggregate evaluations of peers sharing their status (ingroup conformity)	Not Supported ($p \geq 0.10$)
<i>Hypothesis 4: Cultural Goodwill</i>	Lower-status individuals align judgments with higher-status others (upward cultural deference)	Not Supported ($p \geq 0.20$)
<i>Hypothesis 5: Status-Based Distancing Disalignment</i>	Individuals distance judgments from outgroup evaluations (outgroup taste abandonment)	Supported ($p < 0.001$)
<i>Hypothesis 6: Status Consistency</i>	Consistent status profiles exhibit higher stability (“staying”) than inconsistent profiles; status inconsistency heightens conformity	Supported ($p < 0.001$)

Note: Empirical verdicts are based on Difference-in-Differences and multinomial logistic regression models estimated on the primary analytic sample (Age ≥ 24 , $N = 1,981$). Status-subgroup models and behavioral specifications incorporate Inverse Probability Weighting (IPW) on strictly exogenous baseline covariates: age, gender, race/ethnicity, and parental education. Full regression tables and coefficients are provided in the text and Appendix Tables A.2, A.3, and A.4.

As summarized in Table 2, the Difference-in-Differences framework provides a clear verdict on the mechanisms governing aesthetic judgment and social influence, yielding four primary empirical insights that reshape our theoretical understanding.

First, we confirm that repeated aesthetic evaluation produces an intrinsic positive exposure drift in the complete absence of external social cues (*Hypothesis 1*). When individuals re-evaluate an artwork, familiarity alone generates an autonomous boost in appreciation. This finding is methodologically and theoretically consequential: because aesthetic engagement is inherently dynamic, observational and simple pre/post studies that fail to benchmark against an unexposed control baseline risk misattributing natural familiarization to positive social influence.

Second, testing generalized social cues reveals decisive evidence of a pronounced *valence asymmetry* in cultural influence. While positive endorsements fail to generate an evaluative gain significantly above baseline exposure drift (rejecting *Hypothesis 2a*), negative consensus completely halts and overrides this natural appreciation, exerting an authoritative downward veto (supporting *Hypothesis 2b*). Social consensus thus operates asymmetrically: positive feedback provides largely redundant validation for an already familiarizing object, whereas negative feedback acts as an informative alarm, alerting individuals to social risk and enforcing evaluative caution.

Third, our status-stratified models demonstrate that aesthetic boundaries are maintained not through positive emulation, but through defensive, reactive boundary policing. Neither peer imitation within one’s own stratum (*Hypothesis 3*) nor aspirational deference toward elite taste (*Hypothesis 4*) finds empirical support once baseline drift is accounted for; lower-educated individuals remain entirely indifferent to both same-status and elite endorsements. Instead, the primary status-driven dynamic is status-based distancing disalignment (*Hypothesis 5*): highly educated actors—specifically those holding postgraduate credentials—swiftly devalue and abandon an artwork when informed that working-class peers appreciated it. Symbolic distinction operates not through proactive cultivation, but through vigilant, real-time taste abandonment triggered whenever lower-status approval threatens cultural exclusivity.

Finally, discrete behavioral modeling provides robust confirmation that aesthetic conviction is anchored in structural status consistency (*Hypothesis 6*). Resistance to social influence (“staying”) is not the exclusive preserve of high-status actors; rather, both high-consistent and low-consistent individuals display remarkable evaluative inertia. In contrast, status-inconsistent individuals—those whose educational attainment and subjective class identification diverge—experience structural cross-pressures that translate into evaluative fluidity and heightened susceptibility to conformity.

4.2 Limitations and Suggestions for Future Work

Several limitations of the present study suggest fruitful avenues for future research at the intersection of cultural sociology, experimental methods, and stratification.

First, regarding stimulus selection and aesthetic ambiguity, our experiment relied on a single cultural object: James McNeill Whistler’s *Nocturne: Blue and Gold – Old Battersea Bridge*. We deliberately selected this painting because its liminal position between representational landscape and early modernist abstraction—a tension that famously provoked the 1877 Ruskin-Whistler libel trial—provides high baseline evaluative variance and interpretive plasticity. However, cultural goods vary widely in their formal codification, institutional legitimacy, and interpretive openness. In highly institutionalized cultural genres characterized by established canons (such as classical music or canonical literature), prior

socialization and cultural authority may insulate judgments from peer influence. Conversely, in emergent, commercial, or avant-garde fields lacking codified standards, social cues and status signaling may exert even more decisive influence (Childress and Friedkin, 2012; Wohl, 2015). Future research should deploy multi-stimulus factorial designs across diverse cultural domains—including figurative versus abstract visual art, popular versus art-house film, and commercial versus independent music—to systematically identify how aesthetic ambiguity moderates susceptibility to social influence.

Second, regarding the nature of the social influence cue, our design examined unidirectional exposure to aggregate descriptive feedback regarding prior respondents’ opinions. This feature was methodologically essential for cleanly identifying causal directionality: by providing fixed, exogenous peer consensus, the design eliminates the classical reflection problem that besets observational network studies (Centola, 2011; Aral et al., 2009; Shalizi and Thomas, 2011). Nevertheless, in everyday life, cultural influence operates through dynamic, bidirectional interaction within evolving social networks, where individuals simultaneously influence and are influenced by peers, cultural arbiters, and algorithmic recommendation engines (Salganik et al., 2006; Vanzella-Yang and Abrutyn, 2022). In networked environments, structural features such as network density, clustering, and tie strength may amplify informational cascades or foster localized pockets of subcultural resistance (Mark, 2003; Lewis and Kaufman, 2018). Integrating survey-experimental control into interactive online laboratory platforms represents a promising frontier for capturing how collective aesthetic consensus and symbolic boundaries co-evolve in real time.

Third, regarding the identification of status dynamics, while information about alters’ tastes and credentials was strictly randomized, respondents’ own structural positions (educational attainment and subjective social class) were observational. We addressed potential confounding by implementing Inverse Probability Weighting calibrated to strictly exogenous pre-treatment demographic characteristics (age, gender, race/ethnicity, and parental education), ensuring conditional exchangeability without conditioning on post-treatment socioeconomic outcomes (Rosenbaum and Rubin, 1983; Brand and Thomas, 2013). Nonetheless, observational status categories could correlate with unmeasured psychological dispositions, such as cognitive reflection styles or generalized tendencies toward non-conformity. Future laboratory experiments could complement these findings by coupling informational treatments with situational status induction—for example, experimentally manipulating temporary status hierarchies or priming situational status security to determine whether the evaluative fluidity observed among status-inconsistent actors can be causally activated in situ (Correll et al., 2007).

Fourth, regarding temporal granularity and durability, our experiment captured evaluative adjustments across repeated trials separated by an extensive 15-to-20 minute cognitive buffer. While this cognitive buffer proved highly effective in preventing verbatim recall and disrupting demand characteristics, the post-treatment evaluation was necessarily measured within the same experimental session. A critical open question concerns the durability of these socially induced evaluative shifts: does status-based distancing represent a fleeting, situational response that decays once external social cues recede, or does it become consolidated into durable aesthetic dispositions and incorporated into the respondent’s embodied habitus (Lizardo and Skiles, 2016; Reber et al., 1998)? Longitudinal panel designs that re-contact respondents days, weeks, or months following experimental exposure are needed to

map the decay curves and consolidation trajectories of socially induced taste modifications.

Finally, regarding institutional scope conditions, our study was conducted within a quota-matched sample of adults in the contemporary United States. American cultural stratification is widely recognized as featuring relatively permeable symbolic boundaries, pervasive cultural omnivorousness, and a pervasive ethos of egalitarian populism (Erickson, 1996; Bryson, 1996; Lizardo, 2006). In societies characterized by more centralized cultural accreditation, more rigid educational tracking, or stronger aristocratic traditions (such as France or certain East Asian societies), cultural goodwill and deference to highbrow arbiters may remain far more potent (Bourdieu, 1984; Olivos and Wang, 2023). Cross-national comparative experiments will be essential to delineate how macro-institutional regimes condition individual responsiveness to social influence.

4.3 Implications: Status Consistency and Evaluative Fluidity

Our findings carry broad implications for sociological theories of taste, cultural cognition, and social stratification, challenging several foundational assumptions in the sociology of culture while pointing toward a more dynamic, interactional framework.

Reconfiguring Bourdieusian Distinction and the Absence of Cultural Goodwill.

For four decades, Pierre Bourdieu’s *Distinction* (Bourdieu, 1984) has served as the baseline paradigm for the sociology of taste, positing a dual vertical dynamic: dominant classes cultivate exclusive aesthetic dispositions to mark status distinction, while dominated classes exhibit “cultural goodwill”—an aspirational deference toward legitimate culture that reinforces their own subordination. Our experimental findings fundamentally reconfigure this model for contemporary democratic societies. First, we find no empirical evidence of upward cultural goodwill: lower-educated and working-class respondents remain completely indifferent to high-status consensus, exhibiting neither upward deference nor aspirational emulation. In contemporary cultural fields, subordinate groups display aesthetic autonomy and self-containment rather than deference to elite arbiters.

Second, our results show that status distinction operates not as proactive cultural hoarding, but as a *reactive defense mechanism*. Highly educated individuals—specifically those with postgraduate credentials—do not merely patronize highbrow art in the abstract; they actively abandon and devalue cultural goods once informed that working-class peers appreciate them. This status-based distancing disalignment demonstrates that symbolic boundaries are vigilant and reactive: the value of a cultural object is not intrinsic to its formal properties, but is relational and contingent upon who else consumes it (Simmel, 1957; Berger and Heath, 2008). Symbolic distinction in contemporary society is less about displaying deference to traditional cultural authorities and more about maintaining distance from outgroups through real-time taste abandonment.

Status Consistency as the Structural Anchor of Aesthetic Conviction. The most significant theoretical contribution of our behavioral findings is resurrecting and modernizing status crystallization theory (Lenski, 1954; Hope, 1975; Stryker and Macke, 1978) to explain aesthetic malleability. Classic theories of stratification frequently presume that high socioe-

conomic status confers cultural confidence and autonomy, whereas lower status produces insecurity and deference. Our results demonstrate that evaluative stability (the propensity to “stay” and maintain initial aesthetic judgments) is governed by *structural status consistency* rather than vertical status rank. Consistent working-class individuals without a college degree exhibit virtually the same high evaluative inertia as consistent middle-class college graduates. Occupying concordant objective and subjective positions provides an unambiguous structural anchor, shielding actors from social persuasion.

Conversely, status-inconsistent individuals—college graduates who identify as working class, or middle-class identifiers without college credentials—experience acute normative cross-pressures and status ambiguity (Cruces et al., 2013; Correll and Ridgeway, 2003). Lacking an unambiguous reference group, their aesthetic convictions are inherently fragile, leaving them significantly more susceptible to external influence and conformity. This finding provides a structural sociological explanation for why certain individuals are malleable: evaluative fluidity is not an idiosyncratic personality trait or psychological defect, but an emergent structural consequence of contradictory locations in the stratification system.

Valence Asymmetry and Dual-Process Cultural Cognition. Our findings on valence asymmetry speak directly to emerging dual-process theories of cultural cognition and the sociology of distastes (Lizardo and Skiles, 2016; Kahneman, 2011; Wohl, 2015). When individuals encounter a cultural object repeatedly, automatic, Type 1 processing produces perceptual fluency, generating an unreflective, intuitive appreciation (mere exposure effect) (Zajonc, 1968; Reber et al., 1998). Positive endorsements from others simply reinforce this default trajectory, providing redundant information that yields no significant extra evaluative boost.

In contrast, discovering that others dislike an artwork introduces social risk and cognitive dissonance, interrupting automatic fluency and triggering deliberative, Type 2 reassessment. Negative consensus functions as an informative alarm, alerting individuals that admiring the object may signal poor taste or incur social disapproval. Consequently, negative feedback exerts an authoritative veto over aesthetic appreciation. This cognitive asymmetry explains why, in everyday life and lifestyle subcultures, symbolic boundaries are far more reliably organized around shared *distastes* and cultural taboos than around shared positive preferences (Bryson, 1996; Lamont and Molnár, 2002).

Toward a Dynamic, Interactional Sociology of Taste. Finally, this study bridges the long-standing divide between macro-structural theories of stratification and micro-level models of cultural choice. For decades, observational surveys have mapped static correlations between tastes, education, and social networks, but have struggled to disentangle whether these patterns stem from selection (homophily), structural socialization, or real-time social influence (DiMaggio, 1987; Salganik et al., 2006; Vanzella-Yang and Abrutyn, 2022). By experimentally isolating the causal mechanisms of influence across repeated trials, our study shows that aesthetic judgments are inherently relational, asymmetric, and structured by status consistency. Macro-level cultural hierarchies do not replicate themselves mechanically; they are continuously negotiated through real-time behavioral choices—staying, reacting, or conforming—as individuals navigate the social coordinates of cultural objects.

References

- Aral, S., Muchnik, L., and Sundararajan, A. (2009). Distinguishing influence-based contagion from homophily-driven diffusion in dynamic networks. *Proceedings of the National Academy of Sciences of the United States of America*, 106(51):21544–21549.
- Berger, J. and Heath, C. (2008). Who drives divergence? identity signaling, outgroup dissimilarity, and the abandonment of cultural tastes. *Journal of Personality and Social Psychology*, 95(3):593–607.
- Bloom, H. S. (1995). Minimum detectable effects: A simple way to measure the statistical precision of field experiments. *Evaluation Review*, 19(5):547–556.
- Bourdieu, P. (1984). *Distinction: A social critique of the judgement of taste*. Harvard university press.
- Brand, J. E. and Thomas, J. S. (2013). Causal effect heterogeneity. In *Handbook of causal analysis for social research*, pages 189–213. Springer.
- Bryson, B. (1996). “anything but heavy metal”: Symbolic exclusion and musical dislikes. *American Sociological Review*, 61(5):884–899.
- Centola, D. (2011). An experimental study of homophily in the adoption of health behavior. *Science*, 334(6060):1269–1272.
- Childress, C. and Friedkin, N. E. (2012). Cultural reception and production: The social construction of meaning in book clubs. *American sociological review*, 77(1):45–68. doi: 10.1177/0003122411428153.
- Cohen, J. (1988). *Statistical Power Analysis for the Behavioral Sciences*. Lawrence Erlbaum Associates, Hillsdale, NJ, 2nd edition.
- Correll, S. J., Benard, S., and Paik, I. (2007). Getting a job: Is there a motherhood penalty? *American Journal of Sociology*, 112(5):1297–1338.
- Correll, S. J. and Ridgeway, C. L. (2003). Expectation states theory. In *Handbook of social psychology*, pages 29–51. Springer.
- Cruces, G., Perez-Truglia, R., and Tetaz, M. (2013). Biases in perceptions of sorting: How home-ownership and inequality shape beliefs about position. *Journal of Public Economics*, 98:100–112.
- Cutting, J. E. (2003). Gustave caillebotte, french impressionism, and mere exposure. *Psychonomic Bulletin & Review*, 10(2):319–343.
- DiMaggio, P. (1987). Classification in art. *American Sociological Review*, pages 440–455.
- Erickson, B. H. (1996). Culture, class, and connections. *American journal of sociology*, 102(1):217–251.

- Friedman, S., O'Brien, D., and McDonald, I. (2021). Deflecting privilege: Class identity and the intergenerational self. *Sociology*, 55(4):716–733.
- Gelman, A. and Carlin, J. (2014). Beyond power calculations: Assessing type s (sign) and type m (magnitude) errors. *Perspectives on Psychological Science*, 9(6):641–651.
- Greifer, N. (2022). *cobalt: Covariate Balance Tables and Plots*. R package version 4.4.1.
- Heider, F. (1946). Attitudes and cognitive organization. *The Journal of psychology*, 21(1):107–112.
- Hoenig, J. M. and Heisey, D. M. (2001). The abuse of power: The pervasive fallacy of analyzing power for data analytic results. *The American Statistician*, 55(1):19–24.
- Hope, K. (1975). Models of status inconsistency and social mobility effects. *American Sociological Review*, pages 322–343.
- Kahneman, D. (2011). *Thinking, Fast and Slow*. Farrar, Straus and Giroux.
- Lamont, M. and Molnár, V. (2002). The study of boundaries in the social sciences. *Annual review of sociology*, 28(1):167–195.
- Lenski, G. E. (1954). Status crystallization: a non-vertical dimension of social status. *American Sociological Review*, 19(4):405–413.
- Lewis, K. and Kaufman, J. (2018). The conversion of cultural tastes into social network ties. *The American journal of sociology*, 123(6):1684–1742. doi: 10.1086/697525.
- Lizardo, O. (2006). How cultural tastes shape personal networks. *American Sociological Review*, 71(5):778–807.
- Lizardo, O. and Skiles, S. (2016). Cultural objects as prisms: Perceived audience composition of musical genres as a resource for symbolic exclusion. *Socius*.
- Mark, N. P. (2003). Culture and competition: Homophily and distancing explanations for cultural niches. *American sociological review*, 68(3):319–345.
- Okada, S. (2017). Structure of cultural rejection. *Sociological Perspectives*, 60(2):355–377.
- Olivos, F. and Wang, P. (2023). Cultural omnivorousness and status inconsistency in chile: The role of objective and subjective social status. *Sociological Research Online*, 28(3):736–758.
- Pomiès, A. and Arsel, Z. (2018). Retracing the history of the concept of taste. In Arsel”, Z. and Bean”, J., editors, *Taste, Consumption, and Markets*. Routledge, New York.
- Reber, R., Winkielman, P., and Schwarz, N. (1998). Effects of perceptual fluency on affective judgments. *Psychological science*, 9(1):45–48.
- Rosenbaum, P. R. and Rubin, D. B. (1983). The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41–55.

- Salganik, M. J., Dodds, P. S., and Watts, D. J. (2006). Experimental study of inequality and unpredictability in an artificial cultural market. *Science*, 311(5762):854–856.
- Shalizi, C. R. and Thomas, A. C. (2011). Homophily and contagion are generically confounded in observational social network studies. *Sociological methods & research*, 40(2):211–239.
- Simmel, G. (1957). Fashion. *American journal of sociology*, 62(6):541–558.
- Strang, D. and Macy, M. W. (2001). In search of excellence: Fads, success stories, and adaptive emulation. *American Journal of Sociology*, 107(1):147–182.
- Stryker, S. and Macke, A. S. (1978). Status inconsistency and role conflict. *Annual Review of Sociology*, 4(1):57–90.
- Tampubolon, G. (2008). Revisiting omnivores in america circa 1990s: The exclusiveness of omnivores? *Poetics*, 36(2-3):243–264.
- Vanzella-Yang, A. and Abrutyn, S. (2022). Tastes, emotions, and social cohesion: Toward a cultural theory of social exchange. *Journal for the Theory of Social Behaviour*, 52(2):315–335.
- Wohl, H. (2015). Community sense: The cohesive power of aesthetic judgment. *Sociological Theory*, 33(4):299–326.
- Zajonc, R. B. (1968). Attitudinal effects of mere exposure. *Journal of Personality and Social Psychology*, 9(2p2):1–27.

A Statistical Power and Minimum Detectable Effect (MDE) Analysis

To rigorously evaluate the statistical precision and empirical detection boundaries of the study, we conducted a formal ex-post Minimum Detectable Effect (MDE) analysis (Bloom, 1995). Methodological scholarship has established that retrospective or “post-hoc” power calculations based on observed effect sizes and sample estimates are fundamentally flawed and uninformative, functioning merely as a tautological transformation of the observed p -value (Gelman and Carlin, 2014; Hoenig and Heisey, 2001). By contrast, an ex-post MDE analysis answers a principled design question: given the achieved cell sample sizes and the empirical variance structure of the repeated-measures data, what is the smallest true population effect size that the research design was adequately powered (at 80% or 90% power, two-tailed $\alpha = 0.05$) to detect?

In a repeated-measures Difference-in-Differences (DiD) framework where each respondent i provides an initial pre-treatment aesthetic rating (Y_{i1}) and a post-treatment evaluation (Y_{i2}) on a 7-point Likert scale, the standard error of the change score $\Delta Y_i = Y_{i2} - Y_{i1}$ depends directly on the test-retest correlation ρ between trials:

$$\sigma_{\Delta} = \sqrt{\sigma_1^2 + \sigma_2^2 - 2\rho\sigma_1\sigma_2} \quad (6)$$

When the test-retest correlation is high, within-subject differencing cancels out stable individual evaluative variance, yielding a change-score standard deviation substantially smaller than the cross-sectional baseline standard deviation ($\sigma_{\Delta} \ll \sigma_1$). In our primary analytic sample, aesthetic judgments exhibit strong stability across repeated trials ($r = 0.8993$, with $\sigma_1 = 1.504$ and $\sigma_2 = 1.554$), yielding an empirical standard deviation of the change score, $\sigma_{\Delta} = 0.688$. This within-subject autocorrelation reduces the standard error of treatment comparisons by more than 50% relative to an independent post-test design without pre-treatment baselines.

For a two-sample Difference-in-Differences comparison between a treatment condition of size N_T and the unexposed Control Condition of size N_C , the standard error is:

$$SE_{\text{DiD}} = \sigma_{\Delta} \sqrt{\frac{1}{N_T} + \frac{1}{N_C}} \quad (7)$$

The Minimum Detectable Effect at significance level α (two-tailed) and statistical power $1 - \beta$ is calculated as:

$$\text{MDE}_{\text{raw}} = (z_{1-\alpha/2} + z_{1-\beta}) \times SE_{\text{DiD}} \quad (8)$$

where $z_{0.975} = 1.960$ for $\alpha = 0.05$, $z_{0.80} = 0.842$ for 80% power (multiplier = 2.802), and $z_{0.90} = 1.282$ for 90% power (multiplier = 3.242). To facilitate comparison across studies, effect sizes are expressed both in raw 7-point Likert scale units (MDE_{raw}) and as standardized effect sizes relative to baseline variance, Cohen’s $d = \text{MDE}_{\text{raw}}/\sigma_1$ (Cohen, 1988).

Table A.1 presents the complete MDE calculations across all analytical tiers of the experimental design, and Figure A.1 visualizes the theoretical precision curves alongside the empirical thresholds achieved in this study.

Table A.1: Ex-Post Minimum Detectable Effect (MDE) Analysis Across Experimental Specifications (Age ≥ 24 Sample)

Model Specification / Level	N_T	N_C	SE	80% Power		90% Power	
				MDE _{raw}	Cohen's d	MDE _{raw}	Cohen's d
<i>Panel A: Within-Subject Exposure Drift (Paired Comparison)</i>							
Primary Sample Drift	1,981	—	0.015	0.043	0.029	0.050	0.033
Control Condition Drift	157	—	0.055	0.154	0.102	0.178	0.118
<i>Panel B: Pooled Difference-in-Differences Contrasts (vs. Control $N_C = 157$)</i>							
Taste-Only Conditions	150	157	0.079	0.220	0.146	0.255	0.169
Single Status Conditions	258	157	0.070	0.195	0.130	0.226	0.150
Pooled High-Status Conditions	505	157	0.063	0.176	0.117	0.204	0.135
<i>Panel C: Subgroup DiD Contrasts Within Status Consistency Quadrants</i>							
Working Class, No College	51–181	55	0.120	0.335	0.223	0.388	0.258
Middle Class, College	48–160	50	0.123	0.346	0.230	0.400	0.266
Middle Class, No College	26–125	31	0.153	0.429	0.285	0.496	0.330
Working Class, College	15–61	21	0.190	0.532	0.354	0.616	0.409

Note: Minimum Detectable Effects (MDE) are calculated assuming a two-tailed $\alpha = 0.05$ with standard statistical power thresholds of 80% ($z = 2.802$) and 90% ($z = 3.242$). Calculations incorporate the empirical test-retest correlation across repeated trials ($r = 0.8993$, yielding $\sigma_\Delta = 0.688$ relative to baseline $\sigma_1 = 1.504$). MDE_{raw} is expressed in units of the 7-point aesthetic rating scale; Cohen's $d = \text{MDE}_{\text{raw}}/\sigma_1$. In Panel C, N_T indicates the range of treatment cell sizes across specific cue conditions, and reported SE and MDE values correspond to median cell sizes.

The empirical findings from the MDE analysis yield three primary methodological takeaways. First, regarding repeated-measures exposure drift (Panel A of Table A.1), the full sample ($N = 2,253$) possessed extraordinary statistical power, capable of detecting shifts as minute as Cohen's $d \geq 0.028$ (0.041 scale points). Crucially, even within the unexposed Control Condition ($N_C = 177$), the design was powered at 80% to detect a drift of $d \geq 0.099$ (0.148 scale points). This high baseline precision enabled us to reliably identify the empirical +0.107 exposure drift in the Control Condition ($t = 2.05, p = 0.041$), confirming that the upward drift was not an artifact of random noise.

Second, for the primary pooled Difference-in-Differences models compared against the unexposed Control Condition (Panel B of Table A.1), the design possessed 80% power to detect small treatment effects: $d \geq 0.113$ (0.169 points) for pooled high-status conditions ($N_T = 575$), $d \geq 0.125$ (0.188 points) for single-status conditions ($N_T \approx 290$), and $d \geq 0.140$ (0.210 points) for taste-only conditions ($N_T = 172$). All three pooled thresholds fall comfortably below Cohen's standard threshold for a "small" effect size ($d = 0.20$). This confirms that the pooled models were well-equipped to detect subtle social influence cues.

Third, within the disaggregated status-consistency quadrants (Panel C of Table A.1), cell sizes are necessarily smaller, resulting in wider detection envelopes. In the two larger, status-consistent quadrants—Working Class without College ($N_C = 70$) and Middle Class with College ($N_C = 51$)—the median MDE thresholds were $d = 0.203$ (0.304 scale points) and $d = 0.231$ (0.347 scale points), providing adequate power for small-to-moderate effects. In




Figure A.1: Minimum Detectable Effect (MDE) Power Analysis. Panel A displays theoretical MDE curves (Cohen's d) at 80% power ($\alpha = 0.05$) across cell sample sizes, illustrating the precision advantage of repeated-measures DiD across varying levels of test-retest autocorrelation (r). Panel B displays the empirical MDE thresholds achieved in this study across within-subject drift, pooled DiD models, and status-differentiated subgroup comparisons.

contrast, in the smaller, status-inconsistent quadrants—Middle Class without College ($N_C = 33$) and Working Class with College ($N_C = 23$)—the median MDE thresholds increased to $d = 0.281$ (0.421 points) and $d = 0.348$ (0.522 points). Consequently, statistically significant shifts in these cells emerged only when behavioral adjustments were pronounced, such as the -0.394 negative divergence ($p < 0.001$) observed among college-educated working-class respondents exposed to High-Status Dislike cues. Methodologically, this indicates that non-significant coefficients in the smallest status cells should be interpreted as reflecting greater estimation uncertainty rather than definitive evidence of null effects.

B Survey Vignettes and Treatment Prompts

Below we transcribe the verbatim wording, instructions, question stems, and response options of the survey instrument administered in the 2012 study, based on the original Qualtrics protocol.

1. Baseline Measurement (Trial 1, Survey Page 4)

Respondents were presented with the high-resolution image of James McNeill Whistler's *Noc-turne: Blue and Gold – Old Battersea Bridge* (1872) accompanied by the following verbatim introductory instruction and question prompt:

Screen Text / Instruction:

“Please look at this painting for a few seconds.”

Question Prompt:

“Give us your first impression of this painting – how much do you like it?”

Responses were recorded using the following 8 single-choice options:

- 1 = I like it very much
- 2 = I like it
- 3 = I like it just a little
- 4 = I neither like it nor dislike it
- 5 = I dislike it just a little
- 6 = I dislike it
- 7 = I dislike it very much
- 8 = I can’t see the image

Items were reverse-coded during data processing (1 = “dislike very much”, 4 = “neither like nor dislike”, 7 = “like very much”) so that higher numerical values indicate greater aesthetic affinity. Respondents who selected “I can’t see the image” ($N = 6$) were excluded from the analytic sample.

2. Experimental Treatment Vignettes (Survey Page 5)

Following their initial evaluation, respondents in the experimental conditions were exposed to feedback regarding the evaluations of others. Notably, the prompt did *not* describe the evaluating group in relative terms to the participant; instead, it presented absolute descriptive characteristics of the alter group, allowing social proximity or distance to emerge relationally from the participant’s own social location.

A. Status-Differentiated Conditions (Class & Taste): Respondents assigned to conditions crossing status and valence read:

“Thank you for your response. Our survey so far indicates that the group of people most likely to say that they [LIKE / DISLIKE] the painting you saw have completed. . . .”

where the alter characteristics varied by experimental condition:

- **Low-Status Others (LCC):** *“...high school but did not go to college, and work in a food service job.”*
- **High-Status Economic Specialists (ES):** *“...a Bachelor’s degree in Business but no additional degrees, and work at a financial investment firm.”*

- **High-Status Cultural Specialists (CS):** “... several advanced degrees including a PhD, and work at a college or university.”

Immediately following the vignette, respondents answered a subjective similarity probe:

“Given only what you know about this group’s education, occupation, and opinion about the painting, how similar do you think you are to the people in this group?”

with response options: *Very similar, Somewhat similar, Just a little similar, Not very similar, or Not at all similar.*

B. Undifferentiated Taste-Only Conditions (Control 2): Respondents in the Taste-Only conditions were told:

“Thank you for your response. Other people who have taken this survey recently have said that they [LIKE / DISLIKE] the painting you saw.”

“Given only what you know about these people’s opinion about the painting, how similar do you think you are to them?”

with the identical 5-point similarity response options.

C. Unexposed Control Condition (No Alter): Respondents assigned to the pure Control Condition received no feedback screen, evaluation valence, or information about others whatsoever, proceeding directly from Trial 1 to the intervening survey block.

3. Alter Attribution Probe (Survey Page 6, Treatment Respondents Only)

Immediately after the similarity question, respondents in the status-differentiated conditions completed an attribution check:

“In the last question, when you thought about how similar you are to the people in this group, which factor was most influential?”

with response options:

- Their education and/or occupation
- Their opinion about the painting
- A combination of both

4. Intervening Survey Battery (Cognitive Buffer, Survey Pages 7–15)

Between Trial 1 and Trial 2, all respondents completed substantive survey modules lasting 15 to 20 minutes, comprising full educational history (highest degree attained and college major field), intergenerational educational capital (parental bachelor’s degree attainment), employment status, detailed 22-category Standard Occupational Classification (SOC), annual personal pre-tax income (23 income brackets), subjective social class self-identification (Lower class, Working class, Middle class, Upper class), and an extensive musical taste battery assessing liking, listening frequency, favorite genre, and perceived audience demographic stereotypes across 20 distinct musical genres.

5. Post-Treatment Measurement (Trial 2, Survey Page 16)

Following the intervening modules, respondents were re-shown the identical artwork and asked:

Screen Text / Instruction:

“Please look at this painting for a few seconds again.”

Question Prompt:

“Give us your final impression of this painting – how much do you like it?”

recorded using the identical 8 single-choice options (from 1 = “I like it very much” to 7 = “I dislike it very much”, plus “I can’t see the image”), reverse-coded on the same 1–7 scale.

C Full Regression Tables for Moderated DiD Models

Tables A.2 and A.3 report the full set of Difference-in-Differences regression coefficients, standard errors, and sample sizes for the moderated models in the primary analytic sample (Age ≥ 24 , $N = 1,981$), corresponding to the four-tier education analysis (Figure 4 in Section 3.2) and the status consistency matrix.

D Comparison of Unweighted and IPW-Weighted DiD Models

Table A.4 reports the side-by-side comparison between the primary unweighted DiD model and the IPW-weighted model balancing baseline age, gender, race/ethnicity, and parental education across the primary analytic sample (Age ≥ 24). Both specifications produce consistent point estimates, standard errors, and substantive conclusions regarding valence asymmetry and exposure drift suppression.

Table A.2: Difference-in-Differences Treatment Effects Across Four Educational Tiers (Age ≥ 24 Sample, IPW-Weighted)

Treatment Condition	HS or Less	2-Yr / Some Col	Bachelor's	Postgraduate
Generalized Like	+0.284 [†] (0.151)	+0.086 (0.173)	-0.040 (0.145)	-0.232 (0.149)
High-Status Like	+0.032 (0.113)	+0.145 (0.127)	-0.030 (0.123)	-0.235 [†] (0.120)
Low-Status Like	+0.051 (0.130)	-0.113 (0.141)	+0.044 (0.139)	-0.491 ^{***} (0.135)
Generalized Dislike	-0.108 (0.147)	+0.069 (0.154)	-0.308 [*] (0.154)	-0.442 ^{**} (0.151)
High-Status Dislike	-0.140 (0.113)	-0.080 (0.126)	-0.271 [*] (0.125)	-0.266 [*] (0.122)
Low-Status Dislike	+0.240 [†] (0.134)	+0.036 (0.139)	-0.106 (0.136)	-0.345 ^{**} (0.123)
Control Exposure Drift	+0.053 (0.118)	+0.021 (0.103)	+0.149 (0.101)	+0.375 ^{**} (0.137)
Analytic Sample (N)	468	655	577	281

[†] $p < 0.10$, ^{*} $p < 0.05$, ^{**} $p < 0.01$, ^{***} $p < 0.001$. Standard errors in parentheses.

Note: Estimates represent Difference-in-Differences treatment contrasts relative to the unexposed Control Condition within each educational tier, estimated via multinomial Inverse Probability Weighting (IPW) balancing baseline age, gender, race/ethnicity, and parental education.

Table A.3: Difference-in-Differences Treatment Effects Across Status Consistency Quadrants (Age ≥ 24 Sample, IPW-Weighted)

Treatment Condition	Working Class Identification		Middle Class Identification	
	No College	College Degree	No College	College Degree
Generalized Like	+0.179 (0.164)	-0.145 (0.150)	+0.335 [†] (0.183)	-0.094 (0.156)
High-Status Like	+0.097 (0.120)	+0.035 (0.125)	+0.131 (0.135)	-0.142 (0.127)
Low-Status Like	-0.056 (0.139)	+0.122 (0.138)	-0.058 (0.150)	-0.308 [*] (0.147)
Generalized Dislike	-0.137 (0.148)	-0.349 [*] (0.156)	+0.155 (0.174)	-0.344 [*] (0.155)
High-Status Dislike	-0.045 (0.121)	-0.320 [*] (0.125)	-0.201 (0.134)	-0.243 [†] (0.131)
Low-Status Dislike	+0.149 (0.140)	-0.134 (0.134)	-0.011 (0.150)	-0.239 [†] (0.139)
Control Exposure Drift	+0.019 (0.104)	+0.278 ^{**} (0.107)	+0.082 (0.121)	+0.246 [*] (0.112)
Analytic Sample (N)	680	243	443	615

[†] $p < 0.10$, ^{*} $p < 0.05$, ^{**} $p < 0.01$, ^{***} $p < 0.001$. Standard errors in parentheses.

Note: Estimates represent Difference-in-Differences treatment contrasts relative to the unexposed Control Condition within each status consistency quadrant for the primary sample (Age ≥ 24 , $N = 1,981$), estimated using generalized multinomial Inverse Probability Weighting (IPW) balancing baseline age, gender, race/ethnicity, and parental education.

E Sensitivity Analysis: Behavioral Choices Excluding Taste-Only Conditions

To verify that the discrete behavioral patterns documented in the main text (Figure 5) are specifically generated by status-differentiated social influence rather than by undifferentiated taste-consensus cues, we conducted a sensitivity analysis by excluding respondents assigned to the two Taste-Only conditions ($N = 300$ in the primary sample). The resulting analytic sample ($N = 1,643$) restricts behavioral classification to respondents exposed to feedback

Table A.4: Comparison of Unweighted and IPW-Weighted Difference-in-Differences Models (Age ≥ 24 Sample)

Independent Variable	Unweighted DiD		IPW-Weighted DiD	
	Estimate (SE)	<i>t</i> value	Estimate (SE)	<i>t</i> value
Intercept (β_0)	4.624 (0.122)	37.93***	4.625 (0.122)	37.86***
Trial 2 Exposure Drift (β_1)	0.140 (0.054)	2.57*	0.156 (0.055)	2.82**
<i>Difference-in-Differences Treatment Effects (Trial $\times$ Condition):</i>				
Taste Only / Like	+0.064 (0.078)	0.82	+0.047 (0.081)	0.58
Class & Taste / Like / +Status	+0.018 (0.062)	0.29	+0.025 (0.063)	0.40
Class & Taste / Like / -Status	-0.124 [†] (0.069)	-1.79	-0.063 (0.072)	-0.88
Taste Only / Dislike	-0.160* (0.078)	-2.05	-0.186* (0.079)	-2.36
Class & Taste / Dislike / +Status	-0.158* (0.063)	-2.53	-0.207** (0.064)	-3.26
Class & Taste / Dislike / -Status	-0.037 (0.069)	-0.54	-0.062 (0.070)	-0.88
Condition Main Effects Included	Yes		Yes	
Random Intercept Variance (σ_u^2)	2.100		2.060	
Residual Variance (σ_e^2)	0.233		0.982	
<i>N</i> (Observations / Individuals)	3,962 / 1,981		3,962 / 1,981	

[†] $p < 0.10$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Reference condition: unexposed Control Condition.

that combines evaluation valence and others' occupational/educational status, alongside unexposed control stayers.

Figure A.3 displays the Average Marginal Effects from this sensitivity specification, compared against Working Class respondents without a college degree. As shown in the figure, all substantive patterns persist with exceptional stability: status-consistent college-educated respondents exhibit elevated evaluative stability (AME = +0.076, $p < 0.001$), status-inconsistent respondents exhibit heightened conformity (AME = +0.049, $p < 0.001$ for Working Class with College; AME = +0.041, $p < 0.001$ for Middle Class without College), and college education sharply reduces oppositional reactance (AME = -0.031, $p < 0.001$ and AME = -0.041, $p < 0.001$). Notably, under this status-restricted specification, the lower evaluative stability of status-inconsistent non-college respondents remains apparent, reinforcing our core finding that status crystallization anchors aesthetic judgments against external social influence.

F Robustness Check: Alternative Definition of Social Class

Restricting the subjective class definition strictly to respondents self-identifying as “working class” or “middle class” (excluding lower and upper class extremes) produces identical conclusions: status-consistent groups maintain the highest stay rates (Middle/College: 80.5%; Working/No College: 74.1%) and status-inconsistent groups maintain the highest conformity rates (Middle/No College: 18.2%; Working/College: 17.6%).




Figure A.2: Difference-in-Differences Treatment Effects in the Primary Sample (Age ≥ 24 , Unweighted Linear Mixed Model). Horizontal bars denote point estimates with 95% confidence intervals relative to the unexposed Control Condition (dashed line at zero). Colors indicate statistical significance: red ($p < 0.05$), orange ($p < 0.10$), and grey ($p \geq 0.10$).

G Robustness Check: Unrestricted Full Sample ($N = 2,253$) Estimates

In this appendix, we report the complete estimation results obtained when including the entire survey population without age restrictions ($N = 2,253$ complete cases, including the 18–24 cohort; 4,506 repeated observations across trials). As demonstrated below, all substantive conclusions, directional signs, and empirical hypothesis verdicts are completely invariant to the inclusion of younger respondents who may not yet have completed college.

Table A.5 reproduces the demographic distribution of the full unrestricted sample ($N = 2,275$). Tables A.6 and A.7 report the full sample estimates for the stepwise models (moderated separately by education and subjective class) and the four-quadrant status consistency matrix, respectively.




Figure A.3: Sensitivity Analysis: Average Marginal Effects on Behavioral Choices (Stay, Conform, React) Across Status-Consistency Groups, Excluding Taste-Only Conditions (Age ≥ 24 Sample, N = 1,643). Horizontal bars indicate differences in choice probability relative to Working Class respondents without a college degree, with 95% confidence intervals. Colors indicate statistical significance: red ($p < 0.05$) and grey ($p \geq 0.10$).

Table A.5: Demographic Characteristics of Full Unrestricted Sample ($N = 2,275$)

Demographic Variable	Sample N	Sample %	U.S. Pop %	Treatment	Control	Taste Only
Gender						
Women	1,184	52.1%	51.5%	51.7%	52.2%	53.9%
Men	1,091	47.9%	48.5%	48.3%	47.8%	46.1%
Race / Ethnicity						
White	1,463	64.3%	65.1%	64.7%	63.2%	63.1%
Hispanic / Latino	365	16.1%	15.8%	16.2%	14.8%	15.9%
Black / African American	279	12.3%	12.3%	12.1%	13.7%	12.1%
Asian / Pacific Islander	103	4.5%	4.5%	4.4%	5.5%	4.6%
Other / Multiracial	65	2.8%	2.3%	2.5%	2.7%	4.3%
Age Group						
18–24	272	12.0%	13.1%	11.9%	11.0%	12.7%
25–34	410	18.0%	17.5%	18.2%	16.5%	17.9%
35–44	402	17.7%	17.5%	17.9%	13.7%	18.7%
45–54	441	19.4%	19.2%	19.4%	20.9%	18.4%
55–64	372	16.3%	15.6%	16.5%	15.4%	16.1%
65+	377	16.6%	17.2%	16.1%	22.5%	16.2%
Geographic Region						
Northeast	417	18.3%	17.9%	18.8%	16.5%	17.0%
South	816	35.9%	37.2%	36.6%	39.6%	30.5%
Midwest	505	22.2%	21.6%	21.2%	26.4%	24.8%
West	536	23.6%	23.3%	23.4%	17.6%	27.7%
Education (Oversampled)						
Less than High School	89	3.9%	13.3%	3.8%	3.3%	4.6%
High School Diploma	474	20.8%	30.4%	20.3%	23.6%	21.9%
Some College / AA	791	34.8%	28.5%	36.0%	31.3%	30.6%
Bachelor's Degree	625	27.5%	18.1%	27.3%	27.5%	28.5%
Graduate / Professional	296	13.0%	9.6%	12.6%	14.3%	14.5%
Total N	2,275	100.0%		1,746	182	347

Note: U.S. population benchmarks are drawn from the 2010/2012 U.S. Census Bureau Current Population Survey (CPS).

Table A.6: Difference-in-Differences Estimates Moderated Separately by Education and Subjective Class (Full Sample, $N = 2, 253$)

Treatment Condition	Panel A: Education Alone		Panel B: Subjective Class Alone	
	No College	College Degree	Working Class	Middle Class
Generalized Like	+0.185 (0.114)	-0.061 (0.103)	+0.046 (0.106)	+0.147 (0.109)
High-Status Like	+0.117 (0.086)	-0.083 (0.086)	+0.076 (0.082)	-0.017 (0.089)
Low-Status Like	-0.086 (0.096)	-0.138 (0.097)	-0.082 (0.092)	-0.191 [†] (0.099)
Generalized Dislike	+0.004 (0.106)	-0.305** (0.107)	-0.212* (0.102)	-0.079 (0.109)
High-Status Dislike	-0.064 (0.085)	-0.282** (0.087)	-0.100 (0.082)	-0.172 [†] (0.089)
Low-Status Dislike	+0.159 (0.097)	-0.187* (0.094)	+0.051 (0.092)	-0.054 (0.099)
Control Condition Drift	+0.011 (0.075)	+0.247*** (0.075)	+0.097 (0.071)	+0.153 [†] (0.079)
Sample Size (N)	1,340	913	1,099	1,154

[†] $p < 0.10$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Note: Estimates represent Difference-in-Differences treatment contrasts relative to the unexposed Control Condition in the unrestricted full sample ($N = 2, 253$).

Table A.7: Difference-in-Differences Treatment Effects Across Status Consistency Quadrants (Full Sample, $N = 2, 252$, IPW-Weighted)

Treatment Condition	Working Class Identification		Middle Class Identification	
	No College	College Degree	No College	College Degree
Generalized Like	+0.138 (0.158)	-0.125 (0.145)	+0.338 [†] (0.173)	-0.055 (0.152)
High-Status Like	+0.130 (0.115)	-0.024 (0.121)	+0.038 (0.133)	-0.118 (0.126)
Low-Status Like	-0.110 (0.132)	+0.028 (0.133)	-0.111 (0.146)	-0.296* (0.144)
Generalized Dislike	-0.120 (0.140)	-0.396* (0.154)	+0.134 (0.168)	-0.264 [†] (0.151)
High-Status Dislike	+0.006 (0.115)	-0.394*** (0.119)	-0.209 (0.132)	-0.195 (0.128)
Low-Status Dislike	+0.185 (0.132)	-0.171 (0.131)	+0.089 (0.149)	-0.227 [†] (0.136)
Control Exposure Drift	-0.013 (0.057)	+0.311* (0.134)	+0.098 (0.102)	+0.241** (0.070)
Analytic Sample (N)	832	267	507	646

[†] $p < 0.10$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$. Standard errors in parentheses.

Note: Estimates represent Difference-in-Differences treatment contrasts relative to the unexposed Control Condition in the unrestricted full sample ($N = 2, 252$).