

Response to Reviewers and Editor

Manuscript Number: POETICS-D-26-00490

Title: *Stay, React, or Conform? Examining Social Influence Effects on Aesthetic Judgments*

Authors: Omar Lizardo & Sara Skiles

Journal: *Poetics*

Editor-in-Chief: Hannah Wohl

Date: September 18, 2026

Overview and Letter to the Editor

Dear Dr. Wohl,

Thank you very much for the opportunity to revise and resubmit our manuscript, “*Stay, React, or Conform? Examining Social Influence Effects on Aesthetic Judgments*” (POETICS-D-26-00490), for consideration in *Poetics*. We are deeply grateful to you and the two anonymous reviewers for the constructive, insightful, and encouraging feedback.

We have carefully considered all comments from the editorial decision and the review reports. In response, we have developed a thorough revision addressing the methodological, conceptual, and empirical points raised. Specifically, our revision:

1. **Clarifies the Causal Warrant & Empirical Hook:** We explicitly position the study as providing the first clean causal, experimental test of taste influence and distancing dynamics that prior sociological literature had only demonstrated cross-sectionally or observationally.
2. **Adopts a Difference-in-Differences (DiD) Analytic Framework:** We formalize our repeated-measures models into a DiD framework comparing treatments against the no-information control condition across pre- and post-treatment trials, disambiguating baseline measurement from the control group.
3. **Streamlines Theoretical Framing:** We restructured Section 1 into a coherent, linear progression and reframed the findings around **valence asymmetry** (negativity suppressing positive drift vs. positivity facilitating it conditionally).
4. **Articulates a Mechanism for Status Inconsistency & Evaluative Stability:** We provide a clear theoretical mechanism explaining why **status consistency** (in both high- and low-status groups) fosters evaluative certainty/stability (“staying”), whereas status inconsistency generates evaluative fluidity.
5. **Expands Methodological Transparency & Appendix Materials:** We have incorporated ex-post Minimum Detectable Effect (MDE) power diagnostics, survey sampling details, descriptive statistics, exact stimulus/treatment wording, full regression model tables, and expanded Inverse Probability Weighting (IPW) covariate balance diagnostics.

Below, we provide a point-by-point response to all comments and suggestions from the Editor and both Reviewers.

Response to the Editor (Dr. Hannah Wohl)

1. Format Choice: Research Note vs. Full Article

While you originally submitted the manuscript as a “research note,” I offer two possible paths forward. You can revise the manuscript as a research note, honing in on the empirical contribution, or you can flesh out the manuscript and submit it as a full article, which would involve expanding on the theoretical contribution.

Status: [Addressed]

Location in Revised Manuscript: Throughout revised manuscript (`manuscript_R1.tex`)

Response / Actions Taken:

- We are grateful to the Editor for this strategic choice. In accordance with the guidance of both reviewers, we have pursued a path that captures the best of both options:
- We have retained a concise, tightly focused manuscript structure centered squarely on the core empirical facts, while substantially deepening the theoretical framework surrounding status crystallization, valence asymmetry, and causal identification.
- To maintain narrative clarity and momentum in the main text, all extensive technical models, survey instrument vignettes, ex-post MDE power analyses, full regression tables, and auxiliary sensitivity checks have been organized into a comprehensive Supplementary Appendix.

2. Empirical Novelty and Warrant for Research Question

Reviewer 2 suggested that the empirical novelty and warrant for the research question could be made clearer. I think you could be more explicit about what specifically in the question, “How does access to the evaluations of others influence our own aesthetic evaluations of cultural objects?” has not been answered and/or how the literature offers competing answers. As I read it, you are offering the first empirical test of what can be inferred from a synthesis of the literature, but has not yet been empirically proven as such. If this is the case, I would make that clearer and update the literature (per Reviewer 1’s suggestion).

Status: [Addressed]

Location in Revised Manuscript: Section 1 (Introduction), Section 1.1–1.2 (Theoretical Background)

Response / Actions Taken:

- We thank the Editor for this incisive framing. We have revised the Introduction and Section 1 to center directly on this empirical warrant.
- While decades of sociological scholarship (from Bourdieu and Veblen to contemporary cultural sociology) establish that cultural tastes correlate strongly with social status and peer networks, this literature has almost exclusively relied on static, cross-sectional, and observational data.

Consequently, observational homophily conflates selection, common exposure, and actual interpersonal influence.

- We now explicitly frame the contribution as providing the **first experimental causal identification** of how evaluations of others—varying in both valence (like vs. dislike) and status congruence—directly alter, suppress, or reinforce individual aesthetic judgments in real time.

3. Streamlining the Theoretical Front-End

Reviewer 2 also notes that the takeaway and utility for researchers could be clearer. I think the issue might be that, because you brought together theories from multiple literatures to construct multiple hypotheses, your findings then affirm/complicate/challenge these multiple theories from multiple literatures, leaving the contribution a bit muddled. Streamlining the front end might help with this.

Status: [Addressed]

Location in Revised Manuscript: Section 1.1–1.2 (Theoretical Framework), Section 1.4 (Hypotheses), Section 4 (Discussion)

Response / Actions Taken:

- We agree with the Editor and Reviewer 2 that weaving together disparate theoretical threads (informational cascade, status distinction, Bourdieu, and status inconsistency) created a meandering opening.
- We have streamlined Section 1 into three tight, logically cumulative pillars:
 1. **Informational vs. Status-Based Influence:** How exposure to others' evaluations operates generally versus when attached to status cues.
 2. **Valence Asymmetry:** Why negative evaluations exert a potent suppressive veto while positive evaluations facilitate appreciation conditionally.
 3. **Structural Status Consistency:** Why congruent structural positions (high-high or low-low) produce evaluative certainty and behavioral stability ("staying"), whereas inconsistency induces fluidity.

4. Methodological Rigor, DiD, and Demand Characteristics

Finally, a successful revision will involve responding to both reviewers' methodology concerns, and the results will need to remain valid or be sufficiently updated. Reviewer 1 posed various methodological questions and offered suggestions for analyses, such as using the difference-in-differences approach, and asked for more context regarding the methodology. Reviewer 2's point about whether the experimental design might have been obvious to participants was something I wondered too, and providing more methodological detail might help alleviate these

concerns.

Status: [Addressed]

Location in Revised Manuscript: Section 2.1–2.7 (Data and Methods), Section 3.1–3.4 (Results), Section 4 (Discussion)

Response / Actions Taken:

- We have addressed all methodological suggestions raised by both reviewers:
 - Re-formalized our repeated-measures linear mixed models into a **Difference-in-Differences (DiD)** framework contrasting treatment groups against the no-information control condition across pre- and post-treatment trials (addressing R1 #1).
 - Clarified the survey flow, temporal spacing, filler demographic/taste modules, and naturalistic vignette framing that prevented demand characteristics in Section 2.4 (addressing R2 #3).
 - Demonstrated that empirical behavioral patterns (e.g., high inertia / staying and active reactance) directly contradict passive demand compliance.

Response to Reviewer #1

Point 1: Analytical Strategy & Difference-in-Differences (DiD) Approach

1. *The analytical strategy may not be the most appropriate. To better isolate the treatment effects, the authors may consider applying a difference-in-differences approach, using the no-information condition as the control group and exploiting the pre-treatment measurement. Without such a comparison, the observed changes could potentially be attributed to repeated measurement rather than the information treatment itself. The authors refer to two “baselines”—the no-information condition and the measurement conducted before treatment—which suggests that a difference-in-differences design may be feasible.*

Status: [Addressed]

Location in Revised Manuscript: Section 2.7 (Difference-in-Differences Analytic Strategy), Section 3.1 (Baseline Exposure Drift and Generalized Social Influence), Section 3.2 (Status-Differentiated Social Influence Across Educational Groups), Figure 3, Figure 4, Table 2, Appendix Table A.2, and Appendix Table A.3

Response / Actions Taken: We want to warmly and wholeheartedly thank the reviewer for their exceptional perceptiveness. Pushing us to adopt the Difference-in-Differences (DiD) framework was by far the most consequential and transformative recommendation we received. It not only rescued the paper from a fundamental inferential trap that we had failed to recognize in our original submission, but it also uncovered striking causal patterns that were previously masked.

1. What the Original Analysis Was Doing vs. What DiD Accomplishes

In our original submission (`manuscript.tex`), we evaluated simple pre-to-post change scores within each experimental cell:

$$H_0 : \Delta_{\text{within}} = (\bar{Y}_{k,\text{Trial 2}} - \bar{Y}_{k,\text{Trial 1}}) = 0$$

This one-sample test simply asked: *Did respondents in condition k change their aesthetic rating between Trial 1 and Trial 2?*

As the reviewer brilliantly anticipated, this specification fell directly into the trap of conflating the **treatment effect** of external social information with the **mere exposure effect**—the natural upward appreciation that occurs when people re-evaluate an artwork across repeated trials.

Following the reviewer’s guidance, we formulated a formal **Difference-in-Differences (DiD) linear mixed model**:

$$\text{Taste}_{it} = \beta_0 + \beta_1 \text{Trial}_{it} + \sum_{k=2}^K \beta_k \text{Cond}_{ik} + \sum_{k=2}^K \delta_k (\text{Trial}_{it} \times \text{Cond}_{ik}) + u_i + \epsilon_{it} \quad (1)$$

where the unexposed **Control Condition** (which received no information between trials) serves as the omitted reference category. Here, the interaction coefficients δ_k directly isolate the causal treatment effect relative to the counterfactual temporal drift:

$$\delta_k = (\bar{Y}_{k,\text{Trial 2}} - \bar{Y}_{k,\text{Trial 1}}) - (\bar{Y}_{\text{Control},\text{Trial 2}} - \bar{Y}_{\text{Control},\text{Trial 1}})$$

2. Diagnosing What Changed: The “False Positive” Trap of Exposure Drift

When we examined the unexposed **Control Condition** in our primary analytic sample (adults aged ≥ 24 , $N = 1,981$), the reviewer’s intuition was confirmed: respondents who were given **zero information** nevertheless drifted upward substantially upon second viewing.

- In the status-free baseline sample ($N = 457$), Control Condition respondents drifted upward by $+0.156^{**}$ ($SE = 0.055, t = 2.86, p = 0.0044$).
- Across educational strata, exposure drift is heavily stratified: Postgraduates drift upward by $+0.375^{**}$ ($SE = 0.137, p = 0.0064$) and Bachelor’s degree holders by $+0.149$ ($SE = 0.101, p = 0.138$), whereas non-degreed college attenders/AA holders drift only $+0.021$ ($p = 0.841$) and high school respondents drift $+0.053$ ($p = 0.651$).
- Across status consistency quadrants, Middle Class / College control respondents drift up by $+0.250^*$ ($p = 0.025$) and Working Class / College by $+0.200$ ($p = 0.259$), whereas non-college working-class respondents show virtually zero drift ($\beta_1 = +0.016, p = 0.871$).

Why did several positive results from the original draft lose statistical significance?

In the original draft, we reported that respondents exposed to positive high-status cues showed raw upward shifts, which we interpreted as evidence of upward conformity. However, once we benchmark these raw shifts against their counterfactual control group—who drifted upward on their own without seeing any information—the net DiD estimates are statistically indistinguishable from zero (e.g., Bachelor’s degree holders exposed to High-Status Likes: $DiD = -0.030, p = 0.809$). What appeared to be “conformity” in our original submission was almost entirely an artifact of natural exposure drift.

Furthermore, because DiD benchmarks against the Control Condition ($N = 157$ in the primary sample), standard errors appropriately incorporate the variance of the control group ($Var(DiD) = Var(\Delta_k) + Var(\Delta_{Control})$), increasing from ≈ 0.06 to ≈ 0.11 – 0.15 . Consequently, weak positive drifts drop below conventional significance thresholds.

3. What Was Gained: New, Statistically Significant Causal Discoveries

While spurious upward drift vanished, netting out the counterfactual exposure drift **uncovered major causal mechanisms that were completely invisible in the original submission**:

1. Status-Based Distancing & Bourdieusian Taste Abandonment (Hypothesis 5):

In the original draft, we lacked the resolution to observe status distancing. In our revised 4-tier educational analysis, respondents with a **Postgraduate Degree** exposed to *Low-Status Likes* exhibit an immediate, severe evaluative collapse of nearly half a point: **DiD** $= -0.491^{***}$ ($SE = 0.135, t = -3.64, p = 0.0003$, **Cohen’s** $d = 0.33$)—the single largest negative treatment effect in the entire experiment! Because postgraduates naturally drift upward by $+0.375$ under mere exposure, discovering that lower-status peers liked the artwork reverses their appreciation, providing vivid causal evidence of taste abandonment. Similarly, in the status consistency matrix, Middle Class / College respondents exposed to Low-Status Likes drop by **DiD** $= -0.296^*$ ($SE = 0.144, p = 0.0391$).

2. Valence Asymmetry as a Suppressive Veto (Hypothesis 2b):

Across strata, negative evaluations act as an authoritative veto, completely overriding positive

exposure drift:

- *Generalized Dislike (Status-Free Baseline)*: DiD = -0.186^* ($SE = 0.078, p = 0.0169$).
- *High-Status Dislike (Bachelor's Tier)*: DiD = -0.271^* ($SE = 0.125, p = 0.0296$).
- *Generalized Dislike (Bachelor's Tier)*: DiD = -0.308^* ($SE = 0.154, p = 0.0455$).
- *High-Status Dislike (Postgraduate Tier)*: DiD = -0.266^* ($SE = 0.122, p = 0.0298$).
- *Low-Status Dislike (Postgraduate Tier)*: DiD = -0.345^{**} ($SE = 0.123, p = 0.0051$).
- *Generalized Dislike (Postgraduate Tier)*: DiD = -0.442^{**} ($SE = 0.151, p = 0.0035$).

In contrast, positive endorsements produce small, non-significant increments above control drift ($p \geq 0.55$). Negative feedback exerts a dependable veto on appreciation, while positive feedback is largely redundant.

3. Decisive Rejection of Cultural Goodwill (Hypothesis 4) & Same-Status Alignment (Hypothesis 3):

Lower-educated respondents (High School or Less) are completely unresponsive to elite cues (High-Status Like: DiD = $+0.032, p = 0.777$; High-Status Dislike: DiD = $-0.140, p = 0.215$), displaying aesthetic autonomy rather than Bourdieusian deference. Similarly, respondents across all tiers fail to positively align with same-status peers (Bachelor's with High-Status Like: DiD = $-0.030, p = 0.809$; High School with Low-Status Like: DiD = $+0.051, p = 0.697$).

4. Primary Difference-in-Differences Treatment Effects Table

The table below presents the Difference-in-Differences treatment effects across the four educational tiers in the primary analytic sample (Age ≥ 24 , $N = 1,981$), estimated via multinomial Inverse Probability Weighting (IPW) balancing baseline age, gender, race/ethnicity, and parental education (Appendix Table A.2):

Difference-in-Differences Treatment Effects Across Four Educational Tiers (Age ≥ 24 Sample, IPW-Weighted)

Treatment Condition	HS or Less ($N = 468$)	2-Yr / Some Col ($N = 655$)	Bachelor's ($N = 577$)	Postgraduate ($N = 281$)
Generalized Like	$+0.284^\dagger$ (0.151)	$+0.086$ (0.173)	-0.040 (0.145)	-0.232 (0.149)
High-Status Like	$+0.032$ (0.113)	$+0.145$ (0.127)	-0.030 (0.123)	$-0.235^\dagger$ (0.120)
Low-Status Like	$+0.051$ (0.130)	-0.113 (0.141)	$+0.044$ (0.139)	-0.491^{***} (0.135)
Generalized Dislike	-0.108 (0.147)	$+0.069$ (0.154)	-0.308^* (0.154)	-0.442^{**} (0.151)
High-Status Dislike	-0.140 (0.113)	-0.080 (0.126)	-0.271^* (0.125)	-0.266^* (0.122)
Low-Status Dislike	$+0.240^\dagger$ (0.134)	$+0.036$ (0.139)	-0.106 (0.136)	-0.345^{**} (0.123)
Control Exposure Drift	$+0.053$ (0.118)	$+0.021$ (0.103)	$+0.149$ (0.101)	$+0.375^{**}$ (0.137)

$^\dagger p < 0.10$, $^* p < 0.05$, $^{**} p < 0.01$, $^{***} p < 0.001$. Standard errors in parentheses. Estimates represent DiD treatment contrasts relative to the unexposed Control Condition within each tier. The complete 24-cell status consistency quadrant matrix is reported in Appendix Table A.3.

5. Streamlined and Publication-Ready Visualizations

Rather than presenting a cluttered 24-cell grid in the main text, our revised empirical architecture presents three focused, publication-grade figures:

- **Figure 3:** Baseline exposure drift and generalized influence ($N = 457$), illustrating the mere exposure effect and the suppressive veto of Generalized Dislike.
- **Figure 4:** 4-panel horizontal bar plot displaying DiD treatment effects across the four educational tiers ($N = 1,981$), directly evaluating Hypotheses 3, 4, and 5.
- **Figure 5:** 3-panel horizontal bar plot of Average Marginal Effects for discrete behavioral choices (Stay, Conform, React) across status consistency quadrants, directly evaluating Hypothesis 6.

We also added **Table 2** in the main text, systematically summarizing every hypothesis (H1–H6), its theoretical mechanism, empirical estimate, and statistical verdict. Full regression tables are preserved in Appendix C (Tables A.2 and A.3).

Point 2: Disambiguating “Baseline” vs. “Control Condition” Terminology

2. I also suggest avoiding the term “baseline” when referring to the no-information condition, because the study includes a pre-treatment measurement that may also be understood as a baseline. Using the same term for both is confusing. The authors could instead refer to the no-information condition as the control condition.

Status: [Addressed]

Location in Revised Manuscript: Throughout manuscript (manuscript_R1.tex, Section 2.3, Section 3.1, and Section 3.2)

Response / Actions Taken:

- We agree completely with the reviewer that using “baseline” for both the temporal pre-test and the unexposed experimental cell created unnecessary confusion.
- We have systematically overhauled our terminology throughout the manuscript:
 - The no-information experimental group is now strictly referred to as the **“Control Condition”** (or “no-information control group”).
 - The initial, unmanipulated Trial 1 rating (t_1) is referred to as the **“Pre-Treatment Baseline Measurement”**.
 - Changes over repeated trials in the absence of external social cues are explicitly termed **“Exposure Drift”** (or the “Mere Exposure Effect”).
- We have inserted an explicit clarifying note in Section 2.3 of the revised manuscript:

“To prevent terminological confusion, we strictly distinguish between the pre-treatment baseline measurement (Trial 1, which captures initial, uninfluenced aesthetic preference) and the Control Condition (the experimental group that receives no peer information between trials, allowing us to observe pure temporal re-evaluation in the absence

of treatment). We refer to changes over repeated trials in this unexposed group as exposure drift.”

Point 3: Survey Sampling Procedure, Inference Scope, and Descriptive Statistics Table

3. Please provide more details about the survey sampling procedure. Because this is an experimental study, external validity is less central than internal validity. However, if the survey used probability sampling, this would strengthen the conclusions. This information is also important for assessing what the reported confidence intervals represent—for example, whether they reflect only experimental uncertainty within the sample or can support population-level inference. Relatedly, the authors should include a table presenting descriptive statistics for the sample.

Status: [Addressed]

Location in Revised Manuscript: Section 2.1 (Sample Recruitment and Survey Administration), Section 2.2 (Inferential Scope and Interpretation of Confidence Intervals), Table 1, and Appendix Table G.1

Response / Actions Taken: We have substantially expanded Section 2.1 and Section 2.2 to provide comprehensive details regarding the survey sampling procedure, participant recruitment, quota filtering, the primary analytic sample restriction, and the inferential scope of reported statistics:

1. **Data Collection & Panel Administration:** Data were collected between **February 28 and March 6, 2012**, using an online survey administered by **Survey Sampling International (SSI)**. SSI constructed the sample from its verified, opt-in online consumer research panel. Participants gave informed consent and were compensated **\$3.92** for completing the survey (funded by an NSF Doctoral Dissertation Improvement Grant SES-1203426 and the Notre Dame Institute for Scholarship in Liberal Arts).
2. **Quota Settings & Participant Screening:** A total of **3,782 panel members** accessed the survey. Of these, 85 were excluded for non-consent, 12 were blocked for being under 18 years old, and 1,394 were screened out because demographic quota targets (balanced to 2010/2012 U.S. Census benchmarks for age, gender, race/ethnicity, and geographic region) had already been met. An additional 16 participants dropped out during early screens, resulting in a completed sample of $N = 2,275$ **respondents** ($N = 2,253$ providing complete pre- and post-treatment aesthetic judgments).
3. **Primary Analytic Sample Restriction (Age ≥ 24):** To ensure all participants have completed the traditional college-going window and eliminate right-censoring in educational attainment among the 18–24 cohort (where only 20.3% held a bachelor’s degree), our primary analytic sample is restricted to respondents aged 24 and older ($N = 1,981$ complete cases; 3,962 repeated observations across trials). College degree holders represent 43.3% ($N = 858$) of this primary sample. The full unrestricted sample ($N = 2,275$) is preserved as a complete sensitivity check in Appendix G (Tables G.1, G.2, and G.3), confirming that all empirical findings are robust across the full age distribution.

4. **Deliberate Oversampling of College Graduates:** To ensure adequate statistical power within high-cultural-capital occupational cells (Cultural Specialists and Economic Specialists), the study implemented a deliberate oversample of college graduates (43.3% bachelor's degree or higher in the primary sample vs. $\sim 28\%$ nationally).
5. **Scope of Confidence Intervals & Statistical Inference:** Because the survey utilized **quota-matched non-probability online panel sampling** rather than probability sampling with known selection probabilities, the reported standard errors, confidence intervals, and p -values reflect **model-based and experimental design uncertainty** within the sample, rather than design-based population parameter inference. However, the demographic diversity across gender, age, race, and geographic region provides far greater heterogeneity and ecological validity than standard convenience samples.
6. **Descriptive Statistics Table:** We have added **Table 1** to Section 2.1, displaying the demographic characteristics of the primary analytic sample alongside national U.S. benchmarks and across experimental treatment/control conditions:

Point 4: Sample Sizes & Ex-Post Minimum Detectable Effect (MDE) Analysis

4. The sample sizes within conditions are relatively small. If an a priori power analysis was not conducted during the study design, I recommend reporting an ex post minimum detectable effect (MDE) analysis in the appendix. This would help readers assess which effect sizes the study was capable of detecting. The authors may find the following discussion useful: “Why ex post power using estimated effect sizes is bad, but an ex post MDE is not”.

Status: [Addressed]

Location in Revised Manuscript: Section 2.2 (Inferential Scope and Interpretation of Confidence Intervals), Appendix A (Statistical Power and Minimum Detectable Effect Analysis), Table A.1, and Figure A.1

Response / Actions Taken:

- We are deeply grateful to the reviewer for raising this thoughtful methodological recommendation and directing us to the literature distinguishing flawed retrospective *post-hoc* power from principled *ex-post* Minimum Detectable Effect (MDE) analysis (e.g., Bloom 1995; Gelman & Carlin 2014; Hoenig & Heisey 2001; McKenzie 2012).
- Because our sample size was fixed by grant funding during the 2012 data collection, an ex-post MDE analysis provides the appropriate framework to evaluate what effect sizes the study was powered to detect, without conditioning on observed effect estimates or p -values.
- We have added a comprehensive ex-post MDE power analysis in Appendix A of `manuscript_R1.tex`, supported by a dedicated simulation and calculation script (`R/mde_power_analysis.R`), a formal summary table (Table A.1), and a dual-panel figure (Figure A.1) illustrating theoretical power curves and achieved empirical thresholds.
- Crucially, our analysis accounts for the repeated-measures Difference-in-Differences structure of the data:

Table 1: Demographic Characteristics of Survey Sample (Age ≥ 24) Compared to U.S. Population Benchmarks

Demographic Variable	Sample N	Sample %	U.S. Pop %	Treatment ($N = 1,524$)	Control ($N = 157$)	Taste Only ($N = 300$)
Gender						
Women	1,036	52.3%	51.5%	51.9%	52.9%	54.0%
Men	945	47.7%	48.5%	48.1%	47.1%	46.0%
Race / Ethnicity						
White	1,284	64.8%	65.1%	65.0%	64.3%	64.3%
Hispanic / Latino	306	15.4%	15.8%	15.7%	14.6%	14.3%
Black / African American	250	12.6%	12.3%	12.5%	13.4%	13.0%
Asian / Pacific Islander	90	4.5%	4.5%	4.3%	5.7%	5.0%
Other / Multiracial	51	2.6%	2.3%	2.5%	1.9%	3.3%
Age Group						
25–34	407	20.5%	17.5%	20.7%	19.1%	20.7%
35–44	400	20.2%	17.5%	20.4%	15.3%	21.7%
45–54	437	22.1%	19.2%	22.0%	23.6%	21.3%
55–64	366	18.5%	15.6%	18.7%	17.2%	18.0%
65+	371	18.7%	17.2%	18.2%	24.8%	18.3%
Geographic Region						
Northeast	371	18.7%	17.9%	19.0%	17.8%	17.7%
South	689	34.8%	37.2%	35.6%	37.6%	29.3%
Midwest	445	22.5%	21.6%	21.5%	26.1%	25.7%
West	476	24.0%	23.3%	24.0%	18.5%	27.3%
Education (Oversampled)						
Less than High School	66	3.3%	13.3%	3.6%	1.9%	2.7%
High School Diploma	402	20.3%	30.4%	19.8%	22.3%	22.0%
Some College / AA	655	33.1%	28.5%	34.1%	30.6%	29.3%
Bachelor’s Degree	577	29.1%	18.1%	28.6%	29.9%	31.3%
Graduate / Professional	281	14.2%	9.6%	14.0%	15.3%	14.7%
Total N	1,981	100.0%		1,524	157	300

Note: Full unrestricted sample demographics for $N = 2,275$, including the 18–24 cohort, are reported in Appendix Table G.1.

- Pre-treatment and post-treatment aesthetic judgments exhibit exceptionally high test-retest autocorrelation ($r = 0.8993$, with $\sigma_1 = 1.504$ and $\sigma_2 = 1.547$).
- This strong individual-level stability shrinks the standard deviation of change scores to $\sigma_\Delta = 0.688$, cutting the standard error of treatment comparisons by more than 54% relative to an independent post-test design ($\sigma_\Delta/\sigma_1 = 0.457$).
- Across the analytical tiers of the study in the primary analytic sample (assuming two-tailed $\alpha = 0.05$ and standard 80% statistical power), the design was powered to detect:
 1. **Within-Subject Exposure Drift:** Cohen’s $d \geq 0.029$ (0.043 scale points) in the primary analytic sample ($N = 1,981$) and $d \geq 0.102$ (0.154 scale points) in the unexposed Control Condition ($N_C = 157$). This confirmed that the study had high statistical power to detect the +0.156 control drift ($p = 0.0044$).
 2. **Pooled Difference-in-Differences Contrasts (vs. Control $N_C = 157$):**

- Pooled High-Status Conditions ($N_T = 505$): Cohen’s $d \geq 0.117$ (0.176 scale points).
- Single-Status Conditions ($N_T = 258$): Cohen’s $d \geq 0.130$ (0.195 scale points).
- Taste-Only Conditions ($N_T = 150$): Cohen’s $d \geq 0.146$ (0.220 scale points).

All pooled DiD models fall comfortably below Cohen’s benchmark for a “small” effect ($d = 0.20$), confirming adequate power to detect modest cultural cues.

3. Subgroup DiD Contrasts Within Status Consistency Quadrants:

- Working Class, No College ($N_C = 55, N_T = 51$ –181): Median MDE of $d \geq 0.223$ (0.335 scale points).
 - Middle Class, College ($N_C = 50, N_T = 48$ –160): Median MDE of $d \geq 0.230$ (0.346 scale points).
 - Middle Class, No College ($N_C = 31, N_T = 26$ –125): Median MDE of $d \geq 0.285$ (0.429 scale points).
 - Working Class, College ($N_C = 21, N_T = 15$ –61): Median MDE of $d \geq 0.354$ (0.532 scale points).
- This analysis provides clear methodological clarity: the pooled models and status-consistent subgroups were adequately powered to detect small-to-moderate effects, whereas the smaller status-inconsistent cells required moderate-to-large behavioral shifts to achieve statistical significance. As we note in the revised manuscript, null findings in the smallest cells reflect wider estimation uncertainty rather than affirmative evidence of zero influence.

Appendix Table A.1: Ex-Post Minimum Detectable Effect (MDE) Analysis (Age ≥ 24 Sample)

Model Specification / Level	N_T	N_C	SE	80% Power		90% Power	
				MDE _{raw}	Cohen's d	MDE _{raw}	Cohen's d
<i>Panel A: Within-Subject Exposure Drift (Paired Comparison)</i>							
Primary Sample Drift	1,981	—	0.015	0.043	0.029	0.050	0.033
Control Condition Drift	157	—	0.055	0.154	0.102	0.178	0.118
<i>Panel B: Pooled Difference-in-Differences Contrasts (vs. Control $N_C = 157$)</i>							
Taste-Only Conditions	150	157	0.079	0.220	0.146	0.255	0.169
Single Status Conditions	258	157	0.070	0.195	0.130	0.226	0.150
Pooled High-Status Conditions	505	157	0.063	0.176	0.117	0.204	0.135
<i>Panel C: Subgroup DiD Contrasts Within Status Consistency Quadrants</i>							
Working Class, No College	51–181	55	0.120	0.335	0.223	0.388	0.258
Middle Class, College	48–160	50	0.123	0.346	0.230	0.400	0.266
Middle Class, No College	26–125	31	0.153	0.429	0.285	0.496	0.330
Working Class, College	15–61	21	0.190	0.532	0.354	0.616	0.409

Point 5: Exact Treatment Wording & Vignette Prompts in Appendix

5. Please report the exact wording of the information treatments in the appendix. For example, readers need to know whether the wording indicated that the reference group had the same or a different social status from the participants.

Status: **[Addressed]**

Location in Revised Manuscript: Appendix B (Survey Vignettes and Treatment Prompts)

Response / Actions Taken:

- We have added the complete verbatim survey instrument flow, prompt wording, and treatment vignettes to Appendix B of `manuscript_R1.tex`, transcribing directly from the original Qualtrics protocol.
- Specifically, to address the reviewer’s query about whether the reference group was described as having the *same* or *different* status relative to the participant:
 - **The prompt did NOT describe the reference group in relative terms to the participant** (i.e., it did not say “people of the same status as you” or “people of higher status than you”).
 - Instead, the survey provided **objective, absolute descriptions** of the fictional peer group’s education and occupational characteristics:
 - * *Low Cultural Capital (LCC)*: “...completed high school but did not go to college, and work in a food service job.”
 - * *Economic Specialists (ES / +Status)*: “...completed a Bachelor’s degree in Business but no additional degrees, and work at a financial investment firm.”
 - * *Cultural Specialists (CS / +Status)*: “...completed several advanced degrees including a PhD, and work at a college or university.”
 - * *Taste Only (No Status)*: “Other people who have taken this survey recently have said that they [LIKE / DISLIKE] the painting you saw.” (No educational or occupational information provided).
- Immediately following this vignette, respondents answered an attribution item: “*Given only what you know about this group’s education, occupation, and opinion about the painting, how similar do you think you are to the people in this group?*” and were asked which dimension was most influential in their assessment of similarity (education/occupation, opinion of the painting, or both).
- The complete text and options for Trial 1 (Pre-treatment baseline), the intervening demographic and musical modules, and Trial 2 (Post-treatment evaluation) are now fully documented in Appendix B.

Point 6: Disaggregating Education and Subjective Class: Reorganizing Results and Expanding to Four Educational Tiers

6. Before moving to the analysis of combined status, please report the results for subjective social class and education separately. These results could be presented as additional subfigures or included in the appendix. The robustness check using an alternative categorization is useful, as this possibility came to mind immediately when reading the analysis. It would also be relevant to mention this robustness check when the combined-status variable is first described.

Status: [Addressed – Major Reorganization]

Location in Revised Manuscript: Section 2.5 (Variables and Measurement), Section 3.2 (Status-Differentiated Social Influence Across Educational Groups: Testing Hypotheses 3–5), Figure 4, Section 3.3 (Status Inconsistency and Discrete Behavioral Choices: Testing Hypothesis 6), Figure 5, Appendix Table A.2, Appendix Table A.3, Appendix Table G.2, and Appendix F

Response / Actions Taken:

- We are exceptionally grateful to the reviewer for this suggestion. Following the recommendation to examine educational attainment and subjective social class separately prompted us to re-evaluate how structural status location is best operationalized across our empirical pipeline, leading to a major, substantive reorganization that has vastly improved the clarity, theoretical coherence, and empirical power of the manuscript.
- When we followed the reviewer’s suggestion to separate the dimensions, two crucial insights became immediately apparent:
 1. **Subjective Social Class Alone is Diffuse:** Examining subjective class identification without regard to education yielded an attenuated, diffuse picture with limited explanatory leverage on its own.
 2. **Educational Attainment is the Decisive Structural Fault Line:** By contrast, educational attainment emerged as the primary, unambiguous structural axis governing cultural capital, symbolic boundary work, and aesthetic judgment. This aligns perfectly with our experimental design, where the treatment vignettes explicitly differentiated alter status through educational credentials and occupations (doctoral cultural specialists, business baccalaureates, and high school food service workers). Stratifying by education provides direct construct congruence between the status of the respondent (ego) and the status of the evaluator (alter).
- **The Realization of Statistical Power Across Four Educational Tiers:**
 - Guided by the reviewer’s prompt, we examined whether our design was restricted to a blunt binary classification (College vs. No College).
 - Given our primary adult sample ($N = 1,981$ individuals; 3,962 repeated observations) and the exceptional test-retest precision of our repeated-measures Difference-in-Differences design ($r = 0.899, \sigma_{\Delta} = 0.688$), we realized that **we possess sufficient statistical power to move beyond a binary split and disaggregate educational attainment into four granular tiers:**

1. *High School or Less* ($N = 468$)
 2. *Two-Year Associate Degree or Some College* ($N = 655$)
 3. *Four-Year Bachelor's Degree* ($N = 577$)
 4. *Postgraduate or Advanced Degree* ($N = 281$)
- Estimating the DiD models across these four tiers via generalized multinomial Inverse Probability Weighting (IPW) revealed a **striking, monotonic cultural capital gradient** that elevates the empirical narrative of the paper.

• **How the Empirical Results are Now Reorganized in the Manuscript:**

- **Section 3.2: Testing Hypotheses 3–5 via the Four Education Groups (Figure 4 & Table A.2):**
 - * *Hypothesis 3 (Same-Status Alignment) is cleanly rejected:* High-status respondents (Bachelor's and Postgraduates) do not align with high-status likes ($\text{DiD} = -0.030, p = 0.81$ and $\text{DiD} = -0.235, p = 0.05$). Lower-status respondents (High School or Less) do not align with low-status likes ($\text{DiD} = +0.051, p = 0.70$) or dislikes ($\text{DiD} = +0.240, p = 0.074$).
 - * *Hypothesis 4 (Cultural Goodwill) is cleanly rejected:* Respondents with a high school education or less exhibit zero deference or upward alignment with elite consensus, displaying flat shifts under High-Status Likes ($\text{DiD} = +0.032, p = 0.78$) and High-Status Dislikes ($\text{DiD} = -0.140, p = 0.22$). Instead, their responsiveness is selectively populist, shifting positively under Generalized Likes ($\text{DiD} = +0.284, p = 0.061$).
 - * *Hypothesis 5 (Status-Based Distancing Disalignment) is decisively confirmed:* Among Postgraduates (Tier 4), learning that low-status peers liked the artwork produces an immediate evaluative collapse of nearly half a point ($\text{DiD} = -0.491, SE = 0.135, t = -3.64, p = 0.0003$, Cohen's $d = 0.33$). This is the single largest negative effect in the entire experiment, providing vivid confirmation of Bourdieusian symbolic distinction and taste abandonment. Postgraduates also display profound negative sensitivity across all dislike cues (-0.266 to $-0.442, p < 0.03$).
 - * *Evaluative Buffering in Tier 2:* The intermediate two-year Associate / Some College cohort ($N = 655$) displays complete evaluative inertia across all conditions ($p \geq 0.25$), acting as a structural buffer between populist facilitation and elite boundary-drawing.
- **Section 3.3: Reserving Subjective Class for Status Consistency and Discrete Choices (Testing Hypothesis 6):**
 - * Rather than presenting an uninformative singleton subjective class panel, we introduce subjective class self-identification for its true theoretical purpose: crossing it with education to construct the four-quadrant **status consistency typology** (Working/No College, Working/College, Middle/No College, Middle/College).
 - * In accordance with status crystallization theory (Lenski 1954; Hope 1975), status consistency governs normative certainty versus evaluative fluidity. We evaluate this mechanism directly through discrete behavioral choices via multinomial logistic regression (Figure 5):
 - *Inertia (Stay):* Status consistency anchors stability—Middle Class college graduates are +7.3% more likely to stay ($\text{AME} = +0.073, z = 5.47, p < 0.001$, stay rate 81.1%), whereas inconsistent actors display lower stability.
 - *Conformity:* Status-inconsistent actors are significantly more susceptible to exter-

nal influence (Working Class with College +4.2%, $p = 0.0003$; Middle Class without College +4.3%, $p = 0.0002$), while high-status consistent actors are significantly less likely to conform (−3.3%, $p = 0.002$).

- *Reactance*: Oppositional reactance is concentrated among non-college workers, while institutional cultural capital inhibits reactance (−3.0% and −4.0%, $p < 0.003$).

- **Elimination of the Old 24-Cell Continuous Grid (Old Figure 5)**: The old 24-cell continuous interaction matrix in Section 3.3 diluted statistical power, invited multiple-comparison skepticism, and created narrative redundancy with the discrete multinomial models. By eliminating the old Figure 5 from the main text and retaining its tabular estimates in Appendix Table A.3 for completeness, the empirical results flow seamlessly from baseline influence (Section 3.1) to educational status tiers (Section 3.2) to discrete status consistency choices (Section 3.3). Full separate models for Education alone and Subjective Class alone are reported in Appendix Table G.2.

Point 7: Observational Nature of Status Categories & IPW Logic

7. Although the information treatment was randomly assigned, membership in categories such as “low-high” and “low-low” is also determined by respondents’ pre-existing education or subjective social class. For example, individuals with lower levels of education may be more likely to occupy a lower position than the reference group—another reason why the exact treatment wording should be provided. Consequently, respondents do not necessarily have the same probability of entering each combined-status category, and comparisons across these categories may partly capture pre-existing differences between respondents rather than only the causal effect of the experimental information. This issue is common in studies that provide respondents with information about their position in the social structure (e.g., Cruces et al., 2013). The authors acknowledge this problem and address it using inverse probability weighting (IPW), but it would be valuable to devote a few additional sentences to explaining the logic of this approach and its implications for causal interpretation. Such a discussion would be a valuable contribution to the field and would help readers understand the rationale for the analytical strategy.

Status: [Addressed]

Location in Revised Manuscript: Section 2.6 (Inverse Probability Weighting), Figure 2

Response / Actions Taken:

- We are deeply grateful to the reviewer for raising this vital methodological distinction and for pointing us to the insightful framing in Cruces et al. (2013).
- In Section 2.6 of the revised manuscript (`manuscript_R1.tex`), we have added a dedicated, pedagogical explanation of the hybrid observational-experimental setup, the formal mechanics of Inverse Probability Weighting (IPW), and its exact implications for causal interpretation:
 1. **The Structural Confounding Problem**: While the experimental treatments (peer like vs. dislike and peer status vignettes) were randomized, respondents’ own structural status locations (objective education $\times$ subjective class) are observational attributes. Con-

sequently, respondents from different status groups differ systematically on baseline demographic characteristics—for instance, Middle Class college graduates are substantially older and disproportionately have college-educated parents compared to non-college Working Class respondents. Unadjusted comparisons across status groups risk attributing differences in evaluative shifts to social class when they may instead reflect demographic confounding.

2. **Pedagogical Mechanics of IPW:** Following Rosenbaum and Rubin (1983), we address this in two formal steps:

- *Step 1 (Generalized Propensity Scores):* We fit a multinomial logistic regression predicting each respondent’s probability of belonging to their observed status consistency quadrant $S_i \in \{1, 2, 3, 4\}$ (or educational group) conditional on exogenous background covariates $\mathbf{X}_i$ (age, gender, race/ethnicity, and parental educational capital):

$$e_s(\mathbf{X}_i) = P(S_i = s \mid \mathbf{X}_i)$$

- *Step 2 (Constructing Balancing Weights):* We assign each respondent a weight equal to the inverse of their estimated propensity score:

$$w_i = \frac{1}{e_s(\mathbf{X}_i)} = \frac{1}{P(S_i = s \mid \mathbf{X}_i)}$$

3. **Intuitive Logic & Pseudo-Population:** Intuitively, IPW creates a synthetic “**pseudo-population**” by up-weighting underrepresented demographic profiles within each status group (e.g., working-class individuals whose parents attended college) and down-weighting overrepresented profiles. This severs the statistical dependency between the baseline covariates $\mathbf{X}_i$ and status group membership.
4. **Causal Interpretation:** Under the standard assumptions of conditional exchangeability (no unmeasured confounders given $\mathbf{X}_i$) and positivity ($0 < e_s(\mathbf{X}_i) < 1$), the weighted pseudo-population satisfies covariate balance. Consequently, subsequent differences in evaluative drift, behavioral inertia (“staying”), and reactive distancing across status categories can be causally attributed to structural status location and consistency, rather than demographic sorting.

Point 8: Covariate Balance Diagnostics

8. The covariate-balance analysis should be explained more clearly. Does it compare the no-information control condition with all the other categories combined? Balance could also be assessed separately across all experimental groups, with the complete results reported in the appendix.

Status: [Addressed]

Location in Revised Manuscript: Section 2.1, Table 1, Section 2.6 (Inverse Probability Weighting), Figure 2

Response / Actions Taken:

- We have revised Section 2.1 and Section 2.6 to clarify the exact comparisons evaluated in the covariate balance analysis:

- **Balance Across Experimental Arms (Random Assignment):** As reported in Table 1, demographic characteristics (gender, race/ethnicity, age, region, and education) are balanced across experimental conditions (Treatment, Control, Taste Only). Formal tests confirm that random assignment was fully successful, with no baseline imbalances across experimental arms (all $p > 0.40$).
- **Balance Across Observational Status Categories (IPW):** By contrast, Figure 2 specifically evaluates balance across the **observational status categories** (Middle Class/College, Middle Class/No College, Working Class/College, Working Class/No College), comparing each group against the pooled population before and after weighting using `cobalt` (Greifer 2022).
- Prior to weighting, raw standardized mean differences (SMDs) for age and parental education ranged between 0.35 and 0.50 standard deviations. After weighting, all absolute SMDs are compressed well below the conservative 0.10 threshold across all covariates.

Point 9: Figure Polish: Dashed Line & Capitalization

9. In Figure 2, please explain what the dashed line represents and capitalize the variable labels consistently.

Status: [Addressed]

Location in Revised Manuscript: Figure 2, Figure 3, Figure 4, Figure 5, Appendix Figure A.1, Appendix Figure A.2, Appendix Figure A.3

Response / Actions Taken:

- We have systematically overhauled all manuscript figures:
 - Replaced all legacy forest plots with clean, publication-standard horizontal bar plots with color-matched 95% error bars, colored by statistical significance ($p < 0.05$ red, $p < 0.10$ orange, and grey for not significant).
 - The dashed vertical line is explicitly defined in every figure caption and plotted at zero (DiD = 0), representing the null reference benchmark of the unexposed Control Condition (or ± 0.10 threshold for covariate balance in Figure 2).
 - Standardized condition ordering and title capitalization are applied consistently across all figures.

Point 10: Clarification of Generalized Influence Findings

10. In discussing generalized influence, the authors state that “exposure to generalized negative evaluations largely suppresses the positive baseline drift across all groups.” If I understand correctly, readers are expected to interpret the “dislike (no status)” effect in Figure 3. However, only one effect appears to be statistically significant, and it is in the opposite

direction. A similar issue arises for “like (no status).” I therefore wonder whether this interpretation is based on an analysis of the full sample that ignores the subgroup partition but is not currently reported or explained. The authors should clarify which results support this conclusion.

Status: [Addressed]

Location in Revised Manuscript: Section 3.1 (Baseline Exposure Drift and Generalized Social Influence), Equation 3, Figure 3, Table 2

Response / Actions Taken:

- We appreciate the reviewer pointing out the ambiguity in the original draft. The reviewer’s intuition was entirely correct: the earlier narrative conflated pooled models with subgroup interactions.
- In response, we have created a dedicated subsection—**Section 3.1: Baseline Exposure Drift and Generalized Social Influence**—that evaluates generalized influence cleanly on the status-free subset of the primary analytic sample ($N = 457$, comprising the Control Condition $N = 157$, Generalized Like $N = 152$, and Generalized Dislike $N = 148$):
 1. **Mere Exposure Drift in Control (Hypothesis 1):** Repeated viewing of the artwork in the absence of external social cues produces a statistically significant upward appreciation ($\beta_1 = +0.156^{**}$, $SE = 0.055$, $t = 2.86$, $p = 0.0044$).
 2. **Valence Asymmetry:**
 - *Generalized Like (Hypothesis 2a):* Learning that anonymous others liked the artwork yields an increase in raw evaluations ($+0.203^{**}$), but relative to the counterfactual $+0.156$ control drift, this increment is positive and modest ($\delta_{\text{Like}} = +0.047$, $SE = 0.079$, $t = 0.59$, $p = 0.555$).
 - *Generalized Dislike (Hypothesis 2b):* In stark contrast, learning that anonymous others disliked the artwork completely eliminates exposure appreciation (raw change -0.030), resulting in a statistically significant negative DiD suppression relative to control drift ($\delta_{\text{Dislike}} = -0.186^*$, $SE = 0.078$, $t = -2.40$, $p = 0.0169$).
 - *Direct Contrast:* Directly contrasting positive and negative cues yields a significant swing of nearly a quarter point ($\Delta = +0.233^{**}$, $SE = 0.080$, $\chi^2 = 8.51$, $p = 0.0035$).
- This status-free baseline model (visualized in Figure 3) provides unambiguous empirical support for Hypothesis 1 and Hypothesis 2b, confirming that negative consensus exerts an asymmetric veto power against natural aesthetic familiarization.

Point 11: Role of Disaggregated Behavioral Analysis

11. The contribution of the disaggregated analysis is unclear to me. The categorical analysis also reports average effects, as do all the regression models in the paper. If the authors intend to move beyond average effects, they would need to apply methods that explicitly examine dis-

tributional differences or treatment-effect heterogeneity. I recommend removing this analysis, which would also free up space to develop other aspects of the research note.

Status: [Addressed]

Location in Revised Manuscript: Section 3.3 (Status Inconsistency and Discrete Behavioral Choices), Figure 5, Appendix E, Appendix Figure A.3

Response / Actions Taken:

- We thank the reviewer for this thoughtful reflection. We carefully considered removing the behavioral analysis, but retained it in streamlined form because it resolves a fundamental limitation of continuous mean models:
 - In continuous models, an average treatment effect near zero cannot distinguish between complete evaluative stability (everyone staying) versus polarized churn (equal numbers conforming and reacting in opposite directions).
 - By categorizing choices into mutually exclusive discrete modes—**Stay** ($\Delta = 0$), **Conform** (shifting with cue), and **React** (shifting against cue)—multinomial logit models directly evaluate distributional behavioral choices via Average Marginal Effects (AMEs).
- This reveals two critical empirical patterns that continuous means completely obscure, directly testing **Hypothesis 6 (Status Consistency)**:
 1. **Structural Inertia (Stay):** Status consistency anchors stability. High-status consistent individuals (Middle Class / College) exhibit the highest stay rate (81.1%) and are +7.3% more likely to stay (AME = +0.073***, $SE = 0.013$, $z = 5.47$, $p < 0.001$). Low-status consistent actors (Working Class / No College) also exhibit high inertia (73.8% stay rate). By contrast, status-inconsistent actors display lower stability (Middle Class / No College stay rate 71.2%, AME = $-0.026^\dagger$, $p = 0.070$).
 2. **Elevated Conformity Among Inconsistent Actors:** Status-inconsistent individuals are significantly more susceptible to external influence: Working Class with College is +4.2% more likely to conform (AME = +0.042***, $p = 0.0003$), and Middle Class without College is +4.3% more likely to conform (AME = +0.043***, $p = 0.0002$). High-status consistent respondents are significantly less likely to conform (AME = -0.033^{**} , $p = 0.0021$).
 3. **Inhibition of Oppositional Reactance:** College education sharply inhibits oppositional reactance (-3.0% and -4.0% , $p < 0.003$), whereas reactance is concentrated among non-college workers (12.0%).
- Furthermore, to address the reviewer’s concern about robustness, we report a dedicated sensitivity analysis in Appendix E (Appendix Figure A.3) excluding the Taste-Only conditions, confirming that these discrete behavioral distributions replicate with exceptional fidelity when only status-laden cues are evaluated.

Point 12: Full Regression Model Tables in Appendix

12. Report tables with full models in the appendix.

Status: [Addressed]

Location in Revised Manuscript: Table 1, Table 2, Appendix Table A.1, Appendix Table A.2, Appendix Table A.3, Appendix Table A.4, Appendix Table G.1, Appendix Table G.2, Appendix Table G.3

Response / Actions Taken:

- We have fully satisfied this request by incorporating complete regression tables and diagnostic metrics across the manuscript and Appendix:
 - **Table 1:** Primary analytic sample descriptive statistics ($N = 1,981$) vs. 2012 Census benchmarks across experimental conditions.
 - **Table 2:** Comprehensive hypothesis evaluation summary table (H1–H6), theoretical mechanisms, empirical estimates, and statistical verdicts.
 - **Appendix Table A.1:** Complete ex-post Minimum Detectable Effect (MDE) power analysis table.
 - **Appendix Table A.2:** Full Difference-in-Differences regression coefficients across Four Educational Tiers ($N = 1,981$).
 - **Appendix Table A.3:** Complete 24-cell DiD matrix across Status Consistency Quadrants ($N = 1,981$).
 - **Appendix Table A.4:** Side-by-side comparison of Unweighted and IPW-Weighted pooled DiD models with random effect variance components.
 - **Appendix Table G.1:** Full unrestricted sample descriptive statistics ($N = 2,275$).
 - **Appendix Table G.2:** Full unrestricted sample stepwise DiD models (Education alone and Subjective Class alone, $N = 2,253$).
 - **Appendix Table G.3:** Full unrestricted sample 24-cell status consistency matrix ($N = 2,252$).

Point 13: Engagement with Recent Taste & Influence Debates

13. My only comment on the theoretical discussion is that the authors should engage more directly with the latest debates on taste and social influence, including more recent papers and book chapters by Lizardo and other scholars working in this area.

Status: [Addressed]

Location in Revised Manuscript: Section 1 (Introduction), Section 1.1–1.2 (Theoretical Background), Section 4 (Discussion), References

Response / Actions Taken:

- We have updated our theoretical framing and bibliography to engage recent scholarship on cultural cognition, dual-process habitus, aesthetic signaling, and symbolic boundaries (e.g., Lizardo 2016; Lizardo et al. on cognitive foundations of cultural tastes; Wohl 2015; Vanzella 2022).
- This contemporary literature reinforces our core argument: aesthetic evaluations do not merely reflect internal hedonic reactions, but serve as cognitive anchors that interface with external social cues to signal group belonging, status distinction, and structural boundary policing.

Response to Reviewer #2

Point 1: Research Question Rationale & Causal Hook

In section 1.1, the paper reviews past research that shows that people tend to share tastes with similar others and tend to have distaste for the tastes of dissimilar others. (I would agree that this is the state of the literature.) The authors then state the research question: “How does access to the evaluations of others influence our own aesthetic evaluations of cultural objects?” It’s not clear why this research question needs to be asked, given the literature that just got reviewed and the authors’ characterization of that literature. However, I think there is room to ask whether these relationships have been rigorously demonstrated with data that are designed to assess causal relationships. I think it might be the case that past work is observational and cross-sectional. Perhaps that should be the hook that highlights the novelty and contribution of the findings?

Status: [Addressed]

Location in Revised Manuscript: Section 1 (Introduction), Section 1.1–1.2 (Theoretical Background)

Response / Actions Taken:

- We are deeply grateful to Reviewer 2 for this crucial insight. We have restructured the entire opening of the paper around this exact methodological and theoretical hook.
- We now explicitly emphasize that while observational studies establish homophilous taste patterns, they cannot determine whether shared tastes reflect prior sorting, shared structural environments, or active social influence and reactive distancing. Our study provides the first experimental identification of these micro-interactional dynamics.

Point 2: Streamlining Section 1

The third paragraph of section 1.1 reviews work showing that high status groups’ tastes tend to induce conformity in others’ tastes and low status groups’ tastes tend to induce distancing. The fourth paragraph brings the status of the self back in, which gets us back to the starting point, where taste conformity results from social similarity between groups and taste distancing results from social dissimilarity. It is not clear to me that the self was ever not in the analysis. The last paragraph of section 1.1 adds in a new consideration: the distinction between objective and subjective class status. For me, this distinction is itself distinct from the material that has come before, which only deals with objective social similarity and dissimilarity. I think section 1.1 has to be rewritten to be more focused and streamlined on the concepts central to the data analysis. Right now it is meandering.

Status: [Addressed]

Location in Revised Manuscript: Section 1.1–1.3 (Theoretical Background)

Response / Actions Taken:

- We agree with the reviewer that weaving together multiple theoretical threads produced conceptual loops. We have rewritten Section 1 to follow a linear, cumulative progression across three central concepts:
 1. **Concept 1: Informational vs. Status-Based Influence:** Cues from others act as heuristics under aesthetic ambiguity, but influence is fundamentally structured by valence asymmetry and alter status.
 2. **Concept 2: Relational Status Matching:** Influence is governed by congruence between ego’s social location and alter’s status (same-status alignment vs. cross-status reactance).
 3. **Concept 3: Status Consistency vs. Inconsistency:** Structural alignment between objective credentials and subjective class anchors cognitive certainty and evaluative inertia (“staying”), whereas status inconsistency produces evaluative fluidity.

Point 3: Demand Characteristics and Respondent Reactivity

My understanding of the data being analyzed is that they are responses to a question on liking of the painting, compared to a repeated question on liking that followed the provision of information about who else liked/disliked the painting, with the occupation and education of those others being varied. I’m wondering whether this design made respondents aware that a shift in their preference was being measured, and how that might influence their responses. Where any steps taken to minimize demand awareness or respondent reactivity?

Status: [Addressed]

Location in Revised Manuscript: Section 2.4 (Mitigation of Experimenter Demand Characteristics and Respondent Reactivity)

Response / Actions Taken:

- We thank Reviewer 2 for raising this important methodological consideration. We have added a dedicated subsection—Section 2.4—detailing four procedural and design safeguards that minimized demand awareness:
 1. **Intervening Survey Buffer:** Pre- and post-evaluations were separated by an extensive 15-to-20 minute survey battery (employment, demographics, and a comprehensive 20-genre musical taste battery), preventing mechanical recall.
 2. **Naturalistic Information Framing:** Social feedback was presented naturalistically as observational survey summaries, followed by subjective similarity and attribution probes, keeping the perceived purpose centered on social perception rather than persuasion.
 3. **Unexposed Control Benchmark:** Comparing all treatment conditions against the pure Control Condition isolates treatment influence from any potential repeated-testing reactivity.
 4. **Empirical Evidence Against Demand Compliance:** If respondents were complying with experimenter demand, conformity would dominate across conditions. Instead, over 74%

exhibited complete evaluative stability (“staying”), and significant oppositional reactance emerged among discordant status groups.

Point 4: Framing of Valence Findings: Valence Asymmetry

In section 4, the authors write, “While previous research suggests that information about others’ choices acts as a heuristic to resolve uncertainty (Strang and Macy, 2001; Salganik et al., 2006), we find that valence matters.” Isn’t it true that valence mattered in past research as well, insofar as others’ likes vs dislikes mattered? Instead, isn’t the novel finding here is that there is an asymmetry to the valence?

Status: [Addressed]

Location in Revised Manuscript: Abstract, Section 1.2, Section 3.1, Section 4.1

Response / Actions Taken:

- We thank Reviewer 2 for this sharp and constructive correction. We have revised our framing throughout the manuscript to focus specifically on **valence asymmetry** rather than suggesting that prior literature ignored valence.
- We now highlight that while positive peer feedback produces selective, status-contingent shifts, negative peer feedback acts as an asymmetric veto that broadly halts positive exposure drift across groups.

Point 5: Theoretical Mechanism for Status Inconsistency as a Driver of Evaluative Fluidity

Although the discussion very briefly describes status inconsistency as a driver of evaluative fluidity, that finding is not sufficiently integrated with the hypotheses and there is no clear mechanism attributed to it. The authors should formally test whether status inconsistency moderates the experimental conditions. Is the paper also meant to demonstrate that there is a general propensity to change that is distinct from conformity to particular status groups? If so, there needs to be more explanation in the discussion to describe that and to explain the significance of that.

Status: [Addressed]

Location in Revised Manuscript: Section 1.2, Section 1.4 (Hypothesis 6), Section 3.3, Section 4.3

Response / Actions Taken:

- We have fully articulated the theoretical mechanism grounded in Lenski’s (1954) status crystallization theory and contemporary cultural sociology:

- **Status Consistency Anchors Cognitive Certainty:** Actors occupying congruent structural locations (high-high or low-low) are embedded within coherent social milieus with unambiguous normative expectations, generating high subjective certainty and behavioral stability (“staying”).
- **Status Inconsistency Generates Cognitive Ambiguity:** When actors possess college degrees while identifying as working class, or claim middle-class identity without institutional credentials, they experience cross-pressured reference norms. This ambiguity lowers evaluative inertia and heightens susceptibility to external cues (elevated conformity and reactivity).
- We formally tested this mechanism in Section 3.3 via Average Marginal Effects from multinomial logit models, confirming that status-inconsistent groups are significantly more likely to conform (Working Class / College +4.2%, $p = 0.0003$; Middle Class / No College +4.3%, $p = 0.0002$) and display significantly lower evaluative stability.

Point 6: Clarification of “Stay” / Inertia Across Consistent Groups

This sentence needs clarification: “For instance, we see that status consistent individuals, and particularly high-status consistent individuals have the highest propensity to stay with their initial judgment of taste (Middle Class/College: 79.0%; Working Class/No College: 74.4%), a classic behavioral indicator of high status (Correll and Ridgeway, 2003).” The lowest status individuals have a much higher propensity to stay than the two inconsistent groups who both have higher status. What does this imply about the link between high status and staying with a taste judgment? And it amplifies the need for a mechanism to explain the link between status inconsistency and evaluative fluidity.

Status: [Addressed]

Location in Revised Manuscript: Section 3.3, Section 4.3 (Discussion)

Response / Actions Taken:

- Reviewer 2 makes an exceptional point. We have corrected this narrative to clarify that evaluative stability (“staying”) is not a monotonic function of high status alone, but rather an outcome of **structural status consistency and normative certainty**.
- Both high-consistent individuals (Middle Class / College: 81.1% stay rate) and low-consistent individuals (Working Class / No College: 73.8% stay rate) occupy coherent, crystallized positions within the social structure, giving them high evaluative confidence. In contrast, status-inconsistent individuals (Middle Class / No College: 71.2% stay rate, AME = $-0.026^\dagger$, $p = 0.070$; Working Class / College: elevated conformity +4.2%, $p = 0.0003$) experience structural cross-pressures, leading to lower inertia and higher evaluative fluidity.

Point 7: Clarity and Utility for Researchers

For me, one of the virtues of a research note is that it presents clear findings that other researchers can rely on for having established an empirical set of facts. If find that this research note is lacking that clarity, so I'm not sure how readily it will be useful to other researchers.

Status: [Addressed]

Location in Revised Manuscript: Abstract, Section 3.4, Table 2, Section 4.1 (Discussion)

Response / Actions Taken: To ensure maximum utility for other researchers, we have synthesized the study's core empirical findings into **Table 2 (Summary of Hypotheses, Expectations, and Empirical Verdicts)** in Section 3.4, presenting an unambiguous summary of:

1. **Counterfactual Exposure Drift:** In the absence of cues, liking drifts upward significantly in the status-free baseline sample ($\beta_1 = +0.156^{**}, p = 0.0044$), driven by postgraduates ($\beta_1 = +0.375^{**}, p = 0.0064$) and bachelor's degree holders ($+0.149$), while non-college respondents show zero drift ($+0.053, p = 0.651$).
2. **Valence Asymmetry as an Asymmetric Veto:** Negative evaluations from others actively suppress positive exposure drift ($\delta_{\text{Dislike}} = -0.186^*, p = 0.0169$; High-Status Dislike among BA holders $\text{DiD} = -0.271^*, p = 0.0296$; High-Status Dislike among postgraduates $\text{DiD} = -0.266^*, p = 0.0298$), whereas positive evaluations fail to reach significance beyond exposure drift ($\delta_{\text{Like}} = +0.047, p = 0.555$).
3. **Bourdiesian Cross-Status Reactance & Taste Abandonment:** High-status postgraduates devalue and abandon an artwork when informed that working-class peers liked it ($\text{DiD} = -0.491^{***}, p = 0.0003$, Cohen's $d = 0.33$). Similarly, Middle Class / College respondents drop by $\text{DiD} = -0.296^* (p = 0.0391)$.
4. **Status Consistency as an Inertial Anchor:** Evaluative stability is anchored by structural status consistency (stay rates of 81.1% and 73.8%), while status inconsistency elevates conformity ($+4.2\%$ and $+4.3\%, p < 0.001$).