Response to Reviewers and Editor

Manuscript Number: POETICS-D-26-00490
Date: October 11, 2026

Overview and Letter to the Editor

Dear Dr. Wohl,

Thank you very much for the opportunity to revise and resubmit our manuscript, *"Stay, React, or Conform? Examining Social Influence Effects on Aesthetic Judgments"* (POETICS-D-26-00490), for consideration in *Poetics*. We are deeply grateful to you and the two anonymous reviewers for the constructive, insightful, and encouraging feedback.

We have carefully considered all comments from the editorial decision and the review reports. In response, we have developed a thorough revision addressing the methodological, conceptual, and empirical points raised.

Most centrally, our revision: (1) *clarifies the causal warrant and empirical hook* by explicitly positioning the study as providing an experimental causal test of taste influence and distancing dynamics that prior sociological literature had shown primarily cross-sectionally or observationally; (2) *adopts a Difference-in-Differences (DiD) analytic framework* that benchmarks treatment effects against the unexposed control condition across trials, purging the estimates of counterfactual temporal drift and disambiguating baseline measurement from the control group; (3) *streamlines the theoretical framing* into a linear progression organized around *valence asymmetry*, wherein negative evaluations act as an authoritative veto while positive evaluations facilitate appreciation only conditionally; (4) *articulates a sociological mechanism for status consistency*, drawing on status crystallization theory to explain why congruent positions foster evaluative certainty and behavioral stability ("staying"), whereas status inconsistency breeds fluidity; and (5) *expands methodological transparency* by providing ex-post Minimum Detectable Effect (MDE) power diagnostics, survey sampling details, descriptive statistics, exact treatment wording, full regression model tables, and expanded Inverse Probability Weighting (IPW) covariate balance diagnostics in a comprehensive Supplementary Appendix.

Below, we provide a point-by-point response to all comments and suggestions from the Editor and both Reviewers.

Response to the Editor (Dr. Hannah Wohl)

1. Format Choice: Research Note vs. Full Article

> *While you originally submitted the manuscript as a "research note," I offer two possible paths forward. You can revise the manuscript as a research note, honing in on the empirical contribution, or you can flesh out the manuscript and submit it as a full article, which would involve expanding on the theoretical contribution.*

Location in Revised Manuscript: Throughout revised manuscript (`manuscript_R1.tex`)

Response: We are grateful to the Editor for giving us this choice. Following the guidance of both reviewers, we have pursued a path that captures the best of both options. First, we have kept a concise, focused manuscript structure centered squarely on the core empirical findings, while deepening the theoretical framework surrounding status crystallization, valence asymmetry, and causal identification.

To maintain narrative clarity and momentum in the main text, all extensive technical models, survey instrument vignettes, ex-post MDE power analyses, full regression tables, and auxiliary sensitivity checks have been organized into a Supplementary Appendix.

2. Empirical Novelty and Warrant for Research Question

> *Reviewer 2 suggested that the empirical novelty and warrant for the research question could be made clearer. I think you could be more explicit about what specifically in the question, "How does access to the evaluations of others influence our own aesthetic evaluations of cultural objects?" has not been answered and/or how the literature offers competing answers. As I read it, you are offering the first empirical test of what can be inferred from a synthesis of the literature, but has not yet been empirically proven as such. If this is the case, I would make that clearer and update the literature (per Reviewer 1's suggestion).*

Location in Revised Manuscript: Section 1 (Introduction), Section 1.1–1.2 (Theoretical Background)

Response: We thank the Editor for suggesting this framing. We have revised the Introduction and Section 1 to focus directly on this empirical warrant. There, we note that decades of sociological scholarship (from Bourdieu and Simmel to contemporary cultural sociology) suggest that people's tastes should be directly influenced by information about others' tastes. We now explicitly frame the contribution as providing an experimental causal test of how evaluations of others—varying in both valence (like vs. dislike) and status congruence—directly alter, suppress, or reinforce individual aesthetic judgments.

3. Streamlining the Theoretical Front-End

> *Reviewer 2 also notes that the takeaway and utility for researchers could be clearer. I think the issue might be that, because you brought together theories from multiple literatures to construct multiple hypotheses, your findings then affirm/complicate/challenge these multiple theories from multiple literatures, leaving the contribution a bit muddled. Streamlining the front end might help with this.*

Location in Revised Manuscript: Section 1.1–1.2 (Theoretical Framework), Section 1.4 (Hypotheses), Section 4 (Discussion)

Response: We agree with the Editor and Reviewer 2 that weaving together disparate theoretical threads (such as informational cascades, Bourdieusian distinction, and status inconsistency) created a meandering opening in the original submission. In response, we have thoroughly streamlined the theoretical front-end of the manuscript into three integrated, logically cumulative sections.

First, we distinguish between purely informational and status-based social influence, clarifying how exposure to others' evaluations operates in the abstract compared to when evaluations are tethered to explicit status cues. Second, we formalize the logic of *valence asymmetry*, articulating why negative consensus can act as a suppressor, whereas positive consensus facilitates aesthetic appreciation only under certain conditions. Third, we ground our treatment of subjective class and education in Lenski's theory of status crystallization, explaining why structurally congruent positions (*high-high* or *low-low*) can foster evaluative certainty and behavioral stability ("staying"), whereas status inconsistency can generate evaluative fluidity. This streamlined organization establishes a direct and coherent throughline from theoretical expectations to our empirical hypotheses and results.

4. Methodological Rigor, DiD, and Demand Characteristics

> *Finally, a successful revision will involve responding to both reviewers' methodology concerns, and the results will need to remain valid or be sufficiently updated. Reviewer 1 posed various methodological questions and offered suggestions for analyses, such as using the difference-in-differences approach, and asked for more context regarding the methodology. Reviewer 2's point about whether the experimental design might have been obvious to participants was something I wondered too, and providing more methodological detail might help alleviate these concerns.*

Location in Revised Manuscript: Section 2.1–2.6 (Data and Methods), Section 3.1–3.3 (Results), Section 4 (Discussion)

Response: We have comprehensively addressed the methodological concerns raised by both reviewers while ensuring our core findings remain empirically rigorous and properly contextualized. Most substantially, we followed Reviewer 1's recommendation to reformulate our repeated-measures linear mixed models into a formal *Difference-in-Differences (DiD)* framework. This specification explicitly compares each experimental condition against the unexposed no-information control condition across pre- and post-treatment trials, purging the estimates of counterfactual temporal drift and revealing causal patterns that were previously masked.

In addition, we have addressed concerns regarding potential demand characteristics and respondent reactivity (raised by both the Editor and Reviewer 2) in Section 2.3. There, we document the four procedural safeguards embedded in the study design: the substantial 15-to-20 minute intervening survey buffer, the naturalistic framing of aggregate feedback, the unexposed control benchmark, and empirical behavioral patterns—namely, high baseline stability and active oppositional reactance—that run counter to the hypothesis of unchecked compliance with experimenter expectations.

Response to Reviewer 1

Point 1: Analytical Strategy & Difference-in-Differences (DiD) Approach

> *1. The analytical strategy may not be the most appropriate. To better isolate the treatment effects, the authors may consider applying a difference-in-differences approach, using the no-information condition as the control group and exploiting the pre-treatment measurement. Without such a comparison, the observed changes could potentially be attributed to repeated measurement rather than the information treatment itself. The authors refer to two "baselines"—the no-information condition and the measurement conducted before treatment—which suggests that a difference-in-differences design may be feasible.*

Location in Revised Manuscript: Section 2.6 (Difference-in-Differences Analytic Strategy), Section 3.1, Section 3.2, Figures 3–4, Section 4.1 (Table 2), and Appendix Tables A.2–A.3

Response: We want to thank the reviewer for their really good methodological advice on this score. The recommendation to adopt the Difference-in-Differences (DiD) framework was particularly consequential, as it not only rescued the paper from a fundamental inferential trap we had failed to recognize in our original submission but also uncovered striking causal patterns that were previously masked.

In our original submission, we evaluated simple pre-to-post change scores within each experimental cell:

$$H_0 : \Delta_{\text{within}} = (\bar{Y}_{k,\text{Trial 2}} - \bar{Y}_{k,\text{Trial 1}}) = 0$$

This one-sample test assessed whether respondents in condition k changed their aesthetic rating from Trial 1 to Trial 2. As the reviewer incisively observed, this one-group pre-post design conflates the true causal treatment effect of social feedback with the *mere exposure effect*—the natural upward drift in appreciation that occurs when respondents re-evaluate an artwork upon repeated viewing.

Following the reviewer's guidance, we formulated a formal *Difference-in-Differences (DiD) linear mixed model*:

$$\text{Taste}_{it} = \beta_0 + \beta_1 \text{Trial}_{it} + \sum_{k=2}^{K} \beta_k \text{Cond}_{ik} + \sum_{k=2}^{K} \delta_k (\text{Trial}_{it} \times \text{Cond}_{ik}) + u_i + \epsilon_{it} \tag{1}$$

where the unexposed *Control Condition* (which received no information between trials) serves as the omitted reference category. Here, the interaction coefficients δ_k isolate the causal treatment effect relative to counterfactual temporal drift:

$$\delta_k = (\bar{Y}_{k,\text{Trial 2}} - \bar{Y}_{k,\text{Trial 1}}) - (\bar{Y}_{\text{Control,Trial 2}} - \bar{Y}_{\text{Control,Trial 1}})$$

When we examined the unexposed Control Condition in our revised primary analytic sample (adults aged ≥ 25, $N = 1,981$), the reviewer's intuition was confirmed: respondents who were given no information drifted upward on a second viewing. In the status-free baseline sample ($N = 457$), Control Condition respondents exhibited an upward exposure drift of $+0.156$ points ($SE =$

$0.055, t = 2.87, p = 0.0041$). This explains why several apparent positive effects from the original submission lost statistical significance in the DiD framework. In our initial draft, we reported that respondents exposed to high-status endorsements exhibited raw upward shifts, which we interpreted as evidence of upward conformity. However, once compared to the control group—who drifted upward on their own without seeing any social information—the net DiD estimates are statistically indistinguishable from zero. What originally appeared to be active social conformity was, in reality, almost entirely an artifact of unmanipulated exposure drift. Furthermore, because DiD benchmarks directly against the control group ($N = 157$ in the primary sample), standard errors appropriately incorporate the variance of the control condition ($\text{Var}(\text{DiD}) = \text{Var}(\Delta_k) + \text{Var}(\Delta_{\text{Control}})$).

While spurious upward drift was properly eliminated, netting out counterfactual exposure drift revealed causal mechanisms that were entirely obscured in the original submission. First, the DiD framework provides evidence consistent with *status-based distancing and taste abandonment* (current Hypothesis 5). In our revised four-tier educational analysis, respondents with a postgraduate degree exposed to *Low-Status Likes* exhibited an immediate evaluative collapse of nearly half a point—the single largest negative treatment effect in the experiment. Because postgraduates naturally drift upward under mere exposure, discovering that working-class others endorsed the artwork actively reverses their appreciation.

Second, the analysis confirms *valence asymmetry in the patterns of change in aesthetic judgment we observe* (Hypothesis 2b). Across educational strata, negative evaluations consistently outweigh positive ones in their effects on judgments. In the status-free baseline sample, Generalized Dislike exerts a significant downward drag relative to control drift. Among Bachelor's degree holders, both High-Status Dislikes and Generalized Dislikes result in significant aesthetic devaluation. Among postgraduates, this effect is pervasive across all negative cues, including High-Status Dislikes, Low-Status Dislikes, and Generalized Dislikes. By contrast, positive endorsements yield modest increments that do not reach statistical significance beyond control drift.

Third, the DiD estimates provide decisive empirical evidence rejecting both *cultural goodwill (Hypothesis 4)* and *same-status alignment (Hypothesis 3)*. Lower-educated respondents (high school or less) remain completely unresponsive to elite cues, displaying aesthetic autonomy rather than deference. Likewise, respondents across all educational groups fail to positively align with same-status others.

Point 2: Disambiguating "Baseline" vs. "Control Condition" Terminology

> *2. I also suggest avoiding the term "baseline" when referring to the no-information condition, because the study includes a pre-treatment measurement that may also be understood as a baseline. Using the same term for both is confusing. The authors could instead refer to the no-information condition as the control condition.*

Location in Revised Manuscript: Throughout the manuscript (Section 2.2, Section 3.1, and Section 3.2)

Response: We agree completely with the reviewer that using the term "baseline" to denote both the temporal pre-test and the unexposed experimental cell created unnecessary ambiguity. In response, we have systematically overhauled our terminology throughout the manuscript.

Specifically, the no-information experimental group is now strictly referred to as the *Control Con-*

dition (or the *no-information control group*). The initial, unmanipulated evaluation at Trial 1 (t_1) is designated as the *pre-treatment baseline measurement*, and temporal changes over repeated trials in the absence of peer feedback are explicitly termed *exposure drift* (or the *mere exposure effect*). To ensure total clarity for readers from the outset, we have added an explicit note in Section 2.2 of the revised manuscript:

> *"To prevent terminological confusion, we strictly distinguish between the pre-treatment baseline measurement (Trial 1, which captures initial, uninfluenced aesthetic preference) and the Control Condition (the experimental group that receives no peer information between trials, allowing us to observe pure temporal re-evaluation in the absence of treatment). We refer to changes over repeated trials in this unexposed group as exposure drift."*

Point 3: Survey Sampling Procedure, Inference Scope, and Descriptive Statistics Table

> *3. Please provide more details about the survey sampling procedure. Because this is an experimental study, external validity is less central than internal validity. However, if the survey used probability sampling, this would strengthen the conclusions. This information is also important for assessing what the reported confidence intervals represent—for example, whether they reflect only experimental uncertainty within the sample or can support population-level inference. Relatedly, the authors should include a table presenting descriptive statistics for the sample.*

Location in Revised Manuscript: Section 2.1 (Sample Recruitment and Survey Administration), Table 1, and Appendix Table A.6 (Appendix G)

Response: We have substantially expanded Sections 2.1 and 2.2 to provide comprehensive details on the survey sampling procedure, participant recruitment, quota filtering, the primary analytic sample restriction, and the inferential scope of our statistical estimates.

To ensure that all respondents had completed the traditional college-going window and eliminate right-censoring in educational attainment among the 18–24 cohort (where only 18.1% held a bachelor's degree, and 20.3% held a college degree of any kind), we restrict our primary analytic sample to respondents aged 25 and older ($N = 1,981$ complete cases; 3,962 repeated trial observations). Within this primary sample, college degree holders represent 43.3% ($N = 858$). The study implemented a deliberate oversampling of college graduates relative to the national benchmark ($\sim 28\%$) to ensure adequate statistical power across high-cultural-capital occupational cells (Cultural Specialists and Economic Specialists). For completeness, the full unrestricted sample ($N = 2,275$) is preserved as a sensitivity check in Appendix G (Tables A.6, A.7, and A.8), confirming that all empirical findings and hypothesis verdicts are identical across the entire age spectrum.

Regarding the inferential scope of our statistics, because the survey used quota-matched nonprobability panel sampling rather than probability sampling with known selection probabilities, the reported standard errors, confidence intervals, and p-values reflect *model-based and experimental design uncertainty* within the sample rather than design-based inference about population parameters. Nonetheless, the demographic diversity across gender, age, race, and geographic region provides

substantially greater heterogeneity and ecological validity than conventional university convenience samples.

To document the sample composition transparently, we have added Table 1 in Section 2.1, displaying the demographic characteristics of the primary analytic sample alongside U.S. Census benchmarks and across experimental treatment and control conditions (Table 1).

Point 4: Sample Sizes & Ex-Post Minimum Detectable Effect (MDE) Analysis

> *4. The sample sizes within conditions are relatively small. If an a priori power analysis was not conducted during the study design, I recommend reporting an ex post minimum detectable effect (MDE) analysis in the appendix. This would help readers assess which effect sizes the study was capable of detecting. The authors may find the following discussion useful: "Why ex post power using estimated effect sizes is bad, but an ex post MDE is not".*

Location in Revised Manuscript: Section 2.1 (Sample Recruitment and Survey Administration), Appendix A (Statistical Power and Minimum Detectable Effect Analysis), Table A.1, and Figure A.1

Response: This is a thoughtful methodological recommendation and a helpful direction to the literature distinguishing flawed retrospective *post-hoc* power from more principled *ex-post* Minimum Detectable Effect (MDE) analysis. Because our sample size was fixed by grant funding during the 2012 survey administration, an ex-post MDE analysis provides the appropriate framework for assessing which effect sizes the study was capable of detecting, without conditioning on observed effect estimates or sample p-values.

In response, we have added a comprehensive ex-post MDE power analysis in Appendix A of the revised manuscript, supported by a dedicated simulation and calculation script (to be shared in a GitHub repository in the event of publication), a summary table (Table A.1), and a dual-panel figure (Figure A.1) illustrating theoretical power curves alongside achieved empirical thresholds.

Point 5: Exact Treatment Wording & Vignette Prompts in Appendix

> *5. Please report the exact wording of the information treatments in the appendix. For example, readers need to know whether the wording indicated that the reference group had the same or a different social status from the participants.*

Location in Revised Manuscript: Appendix B (Survey Vignettes and Treatment Prompts)

Response: We have added the complete verbatim survey instrument flow, prompt wording, and treatment vignettes to Appendix B of the revised manuscript, transcribed directly from the original Qualtrics protocol.

In particular, addressing the reviewer's query regarding whether the reference group was framed as possessing the *same* or *different* status relative to the respondent, we clarify that the experimental vignettes *never described the reference group in relative terms* (i.e., participants were never told that others had "the same status as you" or "higher status than you"). Instead, the survey provided

descriptions of the fictional others' educational credentials and occupational positions. In the Taste-Only condition, no socioeconomic information was provided ("Other people who have taken this survey recently have said that they [LIKE / DISLIKE] the painting you saw"). The complete text and response scales for Trial 1, the intervening demographic and musical modules, and Trial 2 are now fully documented in Appendix B.

Point 6: Disaggregating Education and Subjective Class: Reorganizing Results and Expanding to Four Educational Tiers

> *6. Before moving to the analysis of combined status, please report the results for subjective social class and education separately. These results could be presented as additional subfigures or included in the appendix. The robustness check using an alternative categorization is useful, as this possibility came to mind immediately when reading the analysis. It would also be relevant to mention this robustness check when the combined-status variable is first described.*

Location in Revised Manuscript: Section 2.4 (Variables and Measurement), Section 3.2 (Status-Differentiated Social Influence: Testing Hypotheses 3–5), Figure 4, Section 3.3 (Status Inconsistency and Discrete Behavioral Choices: Testing Hypothesis 6), Figure 5, Appendix Table A.2, Appendix Table A.3, Appendix Table A.4, Appendix Figure A.2, Appendix Table A.7 (Appendix G), and Appendix F

Response: We are exceptionally grateful to the reviewer for this suggestion. Disaggregating educational attainment and subjective social class prompted us to fundamentally reconsider how we operationalize structural status location in the empirical analysis. This led to a major, substantive reorganization that has significantly improved the manuscript's clarity, theoretical coherence, and empirical power.

When we separated the two dimensions, two crucial insights became immediately apparent. First, examining subjective social class alone yielded an attenuated, diffuse picture with limited explanatory power. Second, educational attainment emerged as the primary, unambiguous structural axis governing cultural capital, symbolic boundary work, and aesthetic judgment. This directly aligns with our experimental design, where the vignettes explicitly differentiated alter status through educational credentials and occupations (doctoral cultural specialists, business baccalaureates, and high school food service workers). Stratifying by education provides direct construct congruence between the status of the respondent (ego) and the status of the evaluator (alter).

Guided by the reviewer's recommendation, we examined whether our design was constrained to a blunt binary classification (College vs. No College). Given our primary adult sample ($N = 1,981$ individuals; 3,962 repeated observations) and the high test-retest precision of our repeated-measures Difference-in-Differences design ($r = 0.899, \sigma_\Delta = 0.688$), we realized we had plenty of statistical power to move beyond a binary split and disaggregate educational attainment into four groups: (1) High School or Less ($N = 468$); (2) Two-Year Associate Degree or Some College ($N = 655$); (3) Four-Year Bachelor's Degree ($N = 577$); and (4) Postgraduate or Advanced Degree ($N = 281$).

Estimating DiD models across these four groups via generalized multinomial Inverse Probability Weighting (IPW) revealed a clear monotonic cultural capital gradient, presented in Section 3.2

(Figure 4 and Appendix Table A.2). First, *Hypothesis 3 (Same-Status Alignment)* is rejected: high-status respondents do not align with high-status likes, and lower-status respondents do not align with low-status likes. Second, *Hypothesis 4 (Cultural Goodwill)* is also rejected: respondents with a high school education or less exhibit no deference to elite consensus, showing flat shifts under High-Status Likes and High-Status Dislikes. Instead, their responsiveness is selectively populist, shifting positively under Generalized Likes. Third, *Hypothesis 5 (Status-Based Distancing)* is strongly supported: among Postgraduates, learning that working-class peers liked the artwork triggers an immediate evaluative collapse of nearly half a point—the single largest negative effect in the entire experiment, providing vivid causal evidence of taste abandonment. Postgraduates also exhibit severe negative sensitivity across all dislike cues.

Having established the educational gradient, we use subjective social class identification for the status inconsistency analysis in Section 3.3: cross-cutting education to construct the four-quadrant status consistency typology (Working Class + No College, Working Class + College, Middle Class + No College, Middle Class + College). Grounded in Lenski's (1954) status crystallization framework, status consistency governs the tension between normative certainty and evaluative fluidity, which we assess directly through discrete behavioral choices using multinomial logistic regression (Figure 5). Status consistency anchors evaluative stability: Middle-class college graduates are +7.3% more likely to maintain their initial judgment, whereas inconsistent actors exhibit lower stability. Conversely, status-inconsistent actors are significantly more susceptible to external influence (Working Class with College: +4.2%; Middle Class without College: +4.3%), while high-status consistent actors are significantly less likely to conform. Oppositional reactance is concentrated among non-college workers (12.0%), whereas institutional cultural capital inhibits reactance.

This reorganization also allowed us to remove the old, unwieldy 24-cell continuous interaction matrix (old Figure 5) from the main text, which had diluted statistical power and invited multiple-comparison concerns. By relocating the 24-cell matrix to Appendix Table A.3 and reporting separate models for Education alone and Subjective Class alone in Appendix Table A.7 (Appendix G), the empirical narrative flows more seamlessly from baseline influence (Section 3.1) to educational status tiers (Section 3.2) to discrete status consistency choices (Section 3.3).

We additionally report this same education-alone/subjective-class-alone specification estimated directly on the primary analytic sample itself (Age $\geq$ 25, N = 1,981) in Appendix Table A.4 and Appendix Figure A.2. This confirms, within the focal sample used throughout the main text, that educational attainment is the more decisive single-axis moderator of social influence (with significant negative suppression under Generalized Dislike and High-Status Dislike among college-educated respondents), while subjective class identification alone yields a comparatively diffuse pattern—consistent with the full-sample robustness check in Appendix G and further motivating our decision to foreground the four-tier educational model in the main text.

Point 7: Observational Nature of Status Categories & IPW Logic

7. Although the information treatment was randomly assigned, membership in categories such as "low–high" and "low–low" is also determined by respondents' pre-existing education or subjective social class. For example, individuals with lower levels of education may be more likely to occupy a lower position than the reference group—another reason why the

exact treatment wording should be provided. Consequently, respondents do not necessarily have the same probability of entering each combined-status category, and comparisons across these categories may partly capture pre-existing differences between respondents rather than only the causal effect of the experimental information. This issue is common in studies that provide respondents with information about their position in the social structure (e.g., Cruces et al., 2013). The authors acknowledge this problem and address it using inverse probability weighting (IPW), but it would be valuable to devote a few additional sentences to explaining the logic of this approach and its implications for causal interpretation. Such a discussion would be a valuable contribution to the field and would help readers understand the rationale for the analytical strategy.

Location in Revised Manuscript: Section 2.5 (Inverse Probability Weighting), Figure 2

Response: We thank the reviewer for raising this vital methodological distinction and for directing us to the framing in Cruces et al. (2013). In Section 2.5 of the revised manuscript, we have added a more detailed explanation of the hybrid observational-experimental setup, the formal mechanics of Inverse Probability Weighting (IPW), and its exact implications for causal inference.

Point 8: Covariate Balance Diagnostics

8. The covariate-balance analysis should be explained more clearly. Does it compare the no-information control condition with all the other categories combined? Balance could also be assessed separately across all experimental groups, with the complete results reported in the appendix.

Location in Revised Manuscript: Section 2.1, Table 1, Section 2.5 (Inverse Probability Weighting), Figure 2

Response: We have revised Sections 2.1 and 2.5 to clarify the exact comparisons evaluated in our balance diagnostics, distinguishing between the balance achieved across experimental conditions via randomization and the balance achieved across observational status groups via IPW weighting.

First, regarding the experimental conditions, Table 1 shows that random assignment successfully balanced demographic characteristics (gender, race/ethnicity, age, parental education, and educational attainment) across the treatment, control, and taste-only conditions, with formal tests confirming no baseline imbalances in these characteristics (all $p > 0.20$), along with the baseline aesthetic evaluation itself ($p = 0.46$). The geographic region distribution shows a marginal variation across arms ($p = 0.04$), a byproduct of quota-filling in the opt-in panel that is unrelated to the treatment contrasts and is absorbed by the experimental design.

Second, regarding the observational status categories, Figure 2 evaluates covariate balance across the four status-consistency quadrants (Middle Class/College, Middle Class/No College, Working Class/College, Working Class/No College). There, we compare each quadrant against the pooled population before and after weighting using the `cobalt` package (Greifer 2022). Prior to weighting, the maximum raw standardized mean differences (SMDs) across the status groups ranged from 0.29 (gender) to 0.59 (parental education) standard deviations. After applying IPW weights, all absolute

SMDs are compressed well below the conservative 0.10 threshold across all covariates, confirming high balance in the weighted sample.

Point 9: Figure Polish: Dashed Line & Capitalization

> *9. In Figure 2, please explain what the dashed line represents and capitalize the variable labels consistently.*

Location in Revised Manuscript: Figure 2, Figure 3, Figure 4, Figure 5, Appendix Figure A.1, Appendix Figure A.2, Appendix Figure A.3, Appendix Figure A.4

Response: We have systematically overhauled all figures across the manuscript and appendix. All legacy forest plots have been replaced with clean horizontal bar plots featuring color-matched 95% error bars, color-coded by statistical significance ($p < 0.05$ in red, $p < 0.10$ in orange, and non-significant in grey).

The dashed vertical line is now explicitly defined in every figure caption and plotted at zero (DiD = 0), representing the null benchmark of the unexposed Control Condition (or the standard ± 0.10 threshold for covariate balance in Figure 2). In addition, standardized condition ordering and consistent title capitalization have been applied uniformly across all figures.

Point 10: Clarification of Generalized Influence Findings

> *10. In discussing generalized influence, the authors state that "exposure to generalized negative evaluations largely suppresses the positive baseline drift across all groups." If I understand correctly, readers are expected to interpret the "dislike (no status)" effect in Figure 3. However, only one effect appears to be statistically significant, and it is in the opposite direction. A similar issue arises for "like (no status)." I therefore wonder whether this interpretation is based on an analysis of the full sample that ignores the subgroup partition but is not currently reported or explained. The authors should clarify which results support this conclusion.*

Location in Revised Manuscript: Section 3.1 (Baseline Exposure Drift and Generalized Social Influence), Equation 5, Figure 3, and Section 4.1 (Table 2)

Response: We appreciate the reviewer identifying this ambiguity in our initial draft. The earlier narrative conflated pooled models with subgroup interactions, creating confusion about which sample supported the claims of generalized influence.

To resolve this, we introduced a dedicated subsection—*Section 3.1: Baseline Exposure Drift and Generalized Social Influence*—that evaluates generalized influence cleanly on the status-free subset of the primary analytic sample ($N = 457$, comprising the Control Condition $N = 157$, Generalized Like $N = 152$, and Generalized Dislike $N = 148$). This analysis reveals two key empirical findings that were occluded in the previous version of the manuscript. First, in support of *Hypothesis 1 (Mere Exposure Drift)*, repeated viewing in the absence of external social cues yields a statistically

significant increase in appreciation. Second, the data reveal a powerful *valence asymmetry*: while learning that anonymous others liked the artwork increases raw evaluations, this change is modest relative to the counterfactual control drift, failing to support Hypothesis 2a. In contrast, learning that anonymous others disliked the artwork completely eliminates exposure appreciation, producing a statistically significant negative DiD suppression relative to control drift, supporting the current *Hypothesis 2b*.

Point 11: Role of Disaggregated Behavioral Analysis

> *11. The contribution of the disaggregated analysis is unclear to me. The categorical analysis also reports average effects, as do all the regression models in the paper. If the authors intend to move beyond average effects, they would need to apply methods that explicitly examine distributional differences or treatment-effect heterogeneity. I recommend removing this analysis, which would also free up space to develop other aspects of the research note.*

Location in Revised Manuscript: Section 3.3 (Status Inconsistency and Discrete Behavioral Choices), Figure 5, Appendix E, Appendix Figure A.4

Response: We thank the reviewer for this thoughtful reflection. We carefully weighed whether to remove the behavioral analysis, but chose to retain it in a streamlined form because it addresses a critical limitation of continuous average effect models. Specifically, in continuous models, an average treatment effect near zero cannot distinguish between complete evaluative stability (everyone staying) versus polarized churn (equal numbers conforming and reacting in opposite directions). By categorizing individual trajectories into mutually exclusive discrete modes—*Stay* ($\Delta = 0$), *Conform* (shifting in the direction of the cue), and *React* (shifting against the cue)—multinomial logit models directly evaluate distributional shifts via Average Marginal Effects (AMEs).

This approach reveals three empirical findings that continuous average effects completely obscure, directly testing *Hypothesis 6 (Status Consistency)*. First, status consistency anchors evaluative stability: high-status consistent individuals (Middle Class / College) exhibit the highest stay rate (81.1%) and are significantly more likely to stay (AME = +0.073), while low-status consistent actors (Working Class / No College) also show high inertia (73.8% stay rate). In contrast, status-inconsistent actors display lower stability (Middle Class / No College stay rate 71.2%). Second, status-inconsistent individuals are significantly more susceptible to external influence: Working-Class college graduates are +4.2% more likely to conform, and Middle-Class non-college respondents are +4.3% more likely to conform, whereas high-status-consistent respondents are significantly less likely to conform. Third, institutional cultural capital sharply inhibits oppositional reactance (-3.0% and $-4.0\%, p < 0.003$), whereas reactance is concentrated among non-college working-class individuals (12.0%).

To address concerns about robustness, we also report a dedicated sensitivity analysis in Appendix E (Appendix Figure A.4) that excludes the Taste-Only conditions. This confirms that the discrete behavioral distributions replicate with exceptional fidelity when only status-laden cues are evaluated.

Point 12: Full Regression Model Tables in Appendix

> *12. Report tables with full models in the appendix.*

Location in Revised Manuscript: Table 1, Table 2, Appendix Table A.1, Appendix Table A.2, Appendix Table A.3, Appendix Table A.4, Appendix Table A.5, Appendix Table A.6, Appendix Table A.7, Appendix Table A.8 (Appendix G)

Response: The revised manuscript incorporates complete regression tables and diagnostic metrics across the manuscript and Supplementary Appendix. In the main text, Table 1 reports descriptive statistics for the primary analytic sample ($N = 1,981$) against 2012 Census benchmarks across experimental conditions, and Table 2 provides a comprehensive summary of all hypotheses (H1–H6), theoretical mechanisms, empirical estimates, and statistical verdicts.

The Supplementary Appendix provides the full suite of econometric models: Table A.1 reports the complete ex-post Minimum Detectable Effect (MDE) power analysis; Table A.2 presents the full Difference-in-Differences regression models across the four educational tiers; Table A.3 details the complete 24-cell DiD matrix across status consistency quadrants; Table A.4 reports the education-alone and subjective-class-alone stepwise specifications estimated on the primary analytic sample; and Table A.5 provides a side-by-side comparison of unweighted and IPW-weighted pooled DiD models with random effect variance components. In addition, Appendix G provides the full unrestricted sample models ($N = 2,275$), including complete descriptive statistics (Table A.6), stepwise DiD models separating education and class (Table A.7), and the 24-cell status consistency matrix (Table A.8).

Point 13: Engagement with Recent Taste & Influence Debates

> *13. My only comment on the theoretical discussion is that the authors should engage more directly with the latest debates on taste and social influence, including more recent papers and book chapters by Lizardo and other scholars working in this area.*

Location in Revised Manuscript: Section 1 (Introduction), Section 1.1–1.2 (Theoretical Background), Section 4 (Discussion), References

Response: We have thoroughly updated our theoretical framing and bibliography to engage recent scholarship on aesthetic signaling and symbolic boundaries.

Response to Reviewer #2

Point 1: Research Question Rationale & Causal Hook

> *In section 1.1, the paper reviews past research that shows that people tend to share tastes with similar others and tend to have distaste for the tastes of dissimilar others. (I would agree that this is the state of the literature.) The authors then state the research question: "How does access to the evaluations of others influence our own aesthetic evaluations of cultural objects?" It's not clear why this research question needs to be asked, given the literature that just got reviewed and the authors' characterization of that literature. However, I think there is room to ask whether these relationships have been rigorously demonstrated with data that are designed to assess causal relationships. I think it might be the case that past work is observational and cross-sectional. Perhaps that should be the hook that highlights the novelty and contribution of the findings?*

Location in Revised Manuscript: Section 1 (Introduction), Section 1.1–1.2 (Theoretical Background)

Response: We have restructured the opening around this methodological and theoretical warrant. We now explicitly emphasize that while observational studies establish the possibility that taste is influenced by information about others' tastes, they cannot provide causal evidence for this effect. Our study addresses this gap by providing an experimental causal test of these micro-interactional dynamics in real time.

Point 2: Streamlining Section 1

> *The third paragraph of section 1.1 reviews work showing that high status groups' tastes tend to induce conformity in others' tastes and low status groups' tastes tend to induce distancing. The fourth paragraph brings the status of the self back in, which gets us back to the starting point, where taste conformity results from social similarity between groups and taste distancing results from social dissimilarity. It is not clear to me that the self was ever not in the analysis. The last paragraph of section 1.1 adds in a new consideration: the distinction between objective and subjective class status. For me, this distinction is itself distinct from the material that has come before, which only deals with objective social similarity and dissimilarity. I think section 1.1 has to be rewritten to be more focused and streamlined on the concepts central to the data analysis. Right now it is meandering.*

Location in Revised Manuscript: Section 1.1–1.3 (Theoretical Background)

Response: We agree with the reviewer that weaving together multiple theoretical threads produced circuitous conceptual loops in the original draft. We have completely rewritten Section 1 to follow a linear, cumulative progression across three central concepts.

We begin with the distinction between informational and status-based social influence, explaining how cues from others act as heuristics under aesthetic ambiguity while remaining fundamentally

structured by valence asymmetry and alter status. Next, we address relational status matching, examining how influence is governed by the structural congruence between ego's social location and alter's status (same-status alignment versus cross-status reactance). Finally, we introduce status consistency, showing how structural alignment between objective credentials and subjective class anchors cognitive certainty and evaluative inertia ("staying"), whereas status inconsistency generates evaluative fluidity.

Point 3: Demand Characteristics and Respondent Reactivity

> *My understanding of the data being analyzed is that they are responses to a question on liking of the painting, compared to a repeated question on liking that followed the provision of information about who else liked/disliked the painting, with the occupation and education of those others being varied. I'm wondering whether this design made respondents aware that a shift in their preference was being measured, and how that might influence their responses. Where any steps taken to minimize demand awareness or respondent reactivity?*

Location in Revised Manuscript: Section 2.3 (Mitigation of Experimenter Demand Characteristics and Respondent Reactivity)

Response: We thank Reviewer 2 for raising this important methodological consideration. In response, we have added a dedicated subsection—Section 2.3—detailing four procedural and design safeguards that minimized demand awareness and respondent reactivity.

First, pre- and post-treatment evaluations were separated by an extensive 15-to-20 minute intervening survey buffer comprising detailed employment questions, parental demographics, and a comprehensive 20-genre musical taste battery, preventing mechanical recall of initial ratings. Second, aggregate peer feedback was presented in a naturalistic manner as observational survey summaries, keeping the apparent purpose focused on social perception rather than persuasion. Third, benchmarking all treatment conditions against the unexposed Control Condition, as we do in the current analysis, helps to separate treatment effects from any potential repeated-testing reactivity. Fourth, the empirical behavioral distributions provide evidence against passive compliance with experimenter demand: approximately 74% of respondents exhibited complete evaluative stability ("staying"), and significant oppositional reactance emerged among discordant status groups.

Point 4: Framing of Valence Findings: Valence Asymmetry

> *In section 4, the authors write, "While previous research suggests that information about others' choices acts as a heuristic to resolve uncertainty (Strang and Macy, 2001; Salganik et al., 2006), we find that valence matters." Isn't it true that valence mattered in past research as well, insofar as others' likes vs dislikes mattered? Instead, isn't the novel finding here is that there is an asymmetry to the valence?*

Location in Revised Manuscript: Abstract, Section 1.2, Section 3.1, Section 4.1, Section 4.3

Response: We thank Reviewer 2 for this sharp and constructive correction. We have revised our framing throughout the manuscript to focus specifically on *valence asymmetry* rather than suggesting that prior literature ignored valence altogether. We now explicitly highlight that while positive peer feedback produces selective, status-contingent shifts, negative peer feedback acts as an authoritative asymmetric veto that broadly halts positive exposure drift across groups. Furthermore, in both the Results (Section 3.1) and Discussion (Sections 4.1 and 4.3), we interpret this asymmetry through the lens of the psychological literature on the *negativity bias* (Rozin and Royzman, 2001; Baumeister et al., 2001), showing how negative consensus functions as a diagnostic social alarm that interrupts automatic perceptual fluency and triggers deliberative evaluative downgrading.

Point 5: Theoretical Mechanism for Status Inconsistency as a Driver of Evaluative Fluidity

> *Although the discussion very briefly describes status inconsistency as a driver of evaluative fluidity, that finding is not sufficiently integrated with the hypotheses and there is no clear mechanism attributed to it. The authors should formally test whether status inconsistency moderates the experimental conditions. Is the paper also meant to demonstrate that there is a general propensity to change that is distinct from conformity to particular status groups? If so, there needs to be more explanation in the discussion to describe that and to explain the significance of that.*

Location in Revised Manuscript: Section 1.2, Section 1.4 (Hypothesis 6), Section 3.3, Section 4.3

Response: We have fully articulated the theoretical mechanism grounded in Lenski's (1954) status crystallization theory and contemporary cultural sociology. Actors occupying congruent structural locations (*high-high* or *low-low*) are embedded within coherent social milieus with unambiguous normative expectations, generating high subjective certainty and behavioral stability ("staying"). In contrast, when actors possess college degrees while identifying as working-class, or claim a middle-class identity without institutional credentials, they experience cross-pressured reference norms. This structural ambiguity lowers evaluative inertia and heightens susceptibility to external cues, leading to elevated conformity and reactivity.

Point 6: Clarification of "Stay" Response Across Consistent Groups

This sentence needs clarification: "For instance, we see that status consistent individuals, and particularly high-status consistent individuals have the highest propensity to stay with their initial judgment of taste (Middle Class/College: 79.0%; Working Class/No College: 74.4%), a classic behavioral indicator of high status (Correll and Ridgeway, 2003)." The lowest status individuals have a much higher propensity to stay than the two inconsistent groups who both have higher status. What does this imply about the link between high status and staying with a taste judgment? And it amplifies the need for a mechanism to explain the link between status inconsistency and evaluative fluidity.

Location in Revised Manuscript: Section 3.3, Section 4.3 (Discussion)

Response: We have corrected the narrative to clarify that evaluative stability ("staying") is not a monotonic function of high status alone, but rather an outcome of *status consistency.*

Point 7: Clarity and Utility for Researchers

For me, one of the virtues of a research note is that it presents clear findings that other researchers can rely on for having established an empirical set of facts. If find that this research note is lacking that clarity, so I'm not sure how readily it will be useful to other researchers.

Location in Revised Manuscript: Abstract, Section 4.1 (Table 2), Section 4 (Discussion)

Response: To ensure maximum clarity and utility for other researchers, we have synthesized the study's core empirical findings into *Table 2 (Summary of Hypotheses, Expectations, and Empirical Verdicts)* in Section 4.1. This table presents a clear summary of the theoretical expectations and empirical verdicts across all six hypotheses.

References

Baumeister, R. F., Bratslavsky, E., Finkenauer, C., and Vohs, K. D. (2001). Bad is stronger than good. *Review of General Psychology*, 5(4):323–370.

Rozin, P. and Royzman, E. B. (2001). Negativity bias, negativity dominance, and contagion. *Personality and Social Psychology Review*, 5(4):296–320.